Sparsity in Convolutional Neural Networks

Yuan YAO

HKUST

Based on Qingyun Sun, Vardan Papyan talks etc.

Acknowledgement

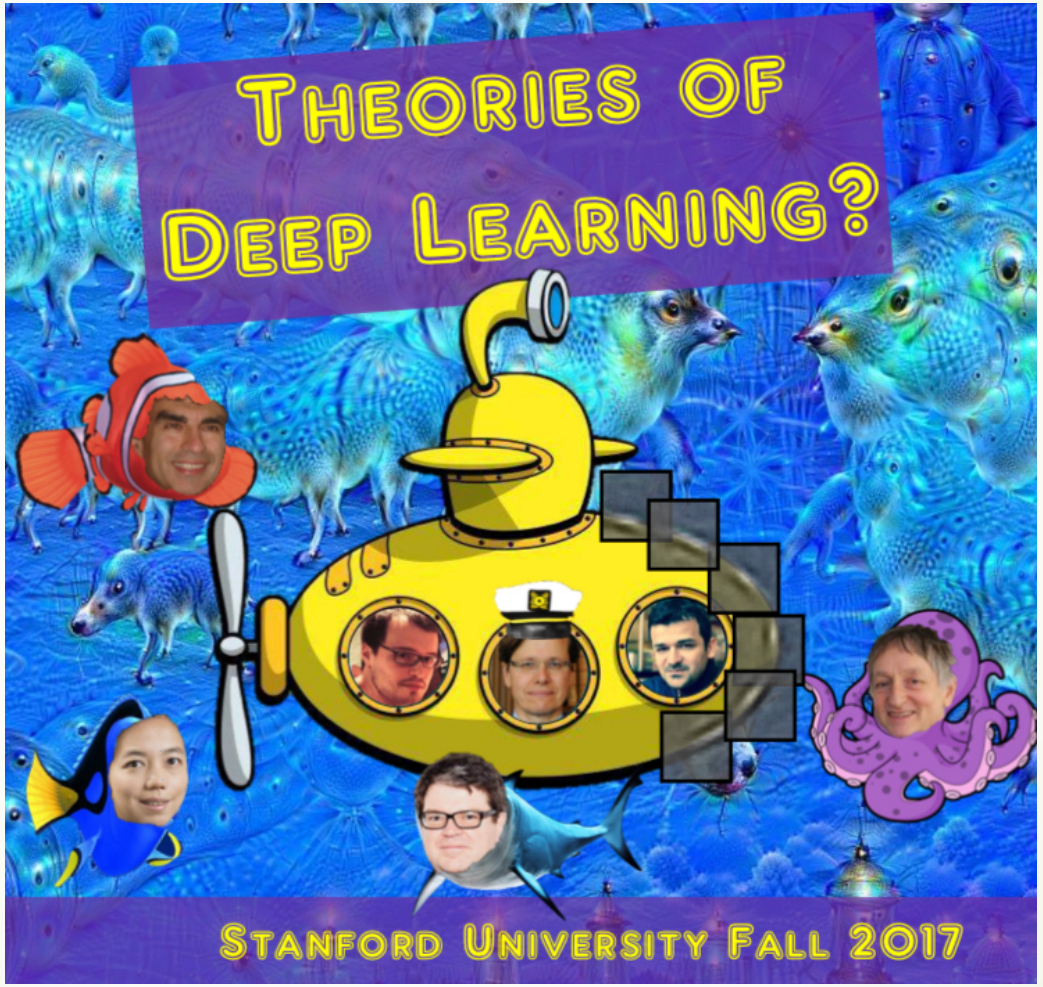

A following-up course at HKUST: https://deeplearning-math.github.io/

Sparsity: Central idea in Stats

Compressive Sensing:

$$\mathbf{X} = \mathbf{D}\boldsymbol{\Gamma}$$

$$\hat{\boldsymbol{\Gamma}} = \arg\min_{\boldsymbol{\Gamma}} \|\boldsymbol{\Gamma}\|_1 \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\boldsymbol{\Gamma}$$

Sparsity: Central idea in Stats

Lasso:

$$\mathbf{Y} = \mathbf{D}\boldsymbol{\Gamma} + \mathbf{E}$$

$$\hat{\boldsymbol{\Gamma}} = \arg\min_{\boldsymbol{\Gamma}} \frac{1}{2}\|\mathbf{Y} - \mathbf{D}\boldsymbol{\Gamma}\|_2^2 + \lambda\|\boldsymbol{\Gamma}\|_1$$

After Deep Revolution, is sparsity still important?

Sparsity observed in CNN

Sparsity in Practice

The activation of RELU layer is sparse.

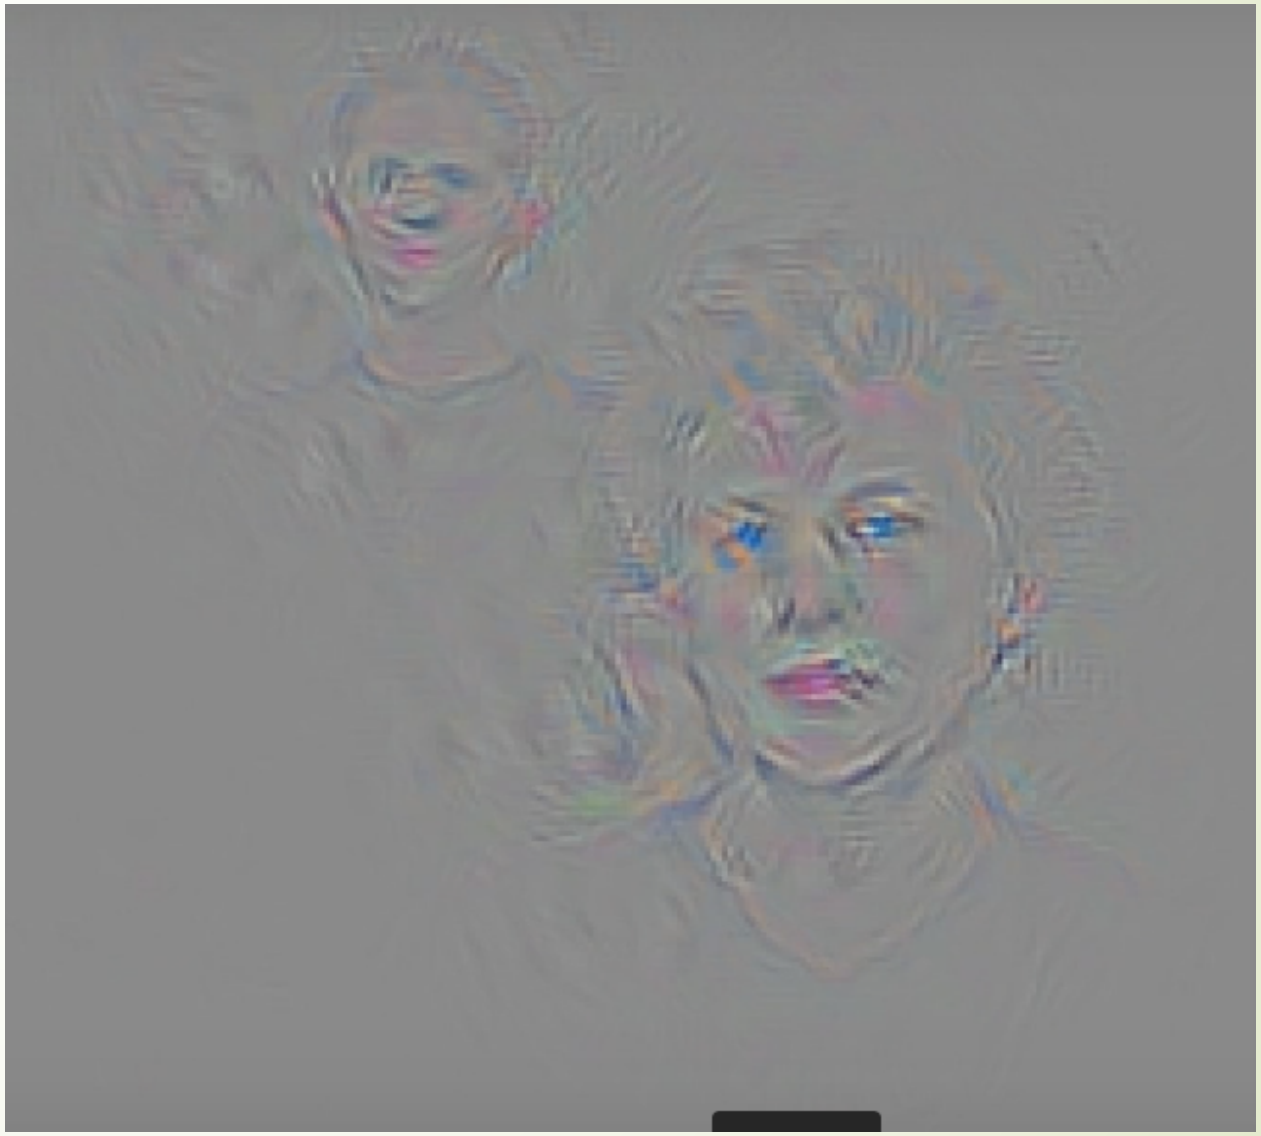

Sparsity in Practice

The activation of RELU layer is sparse.

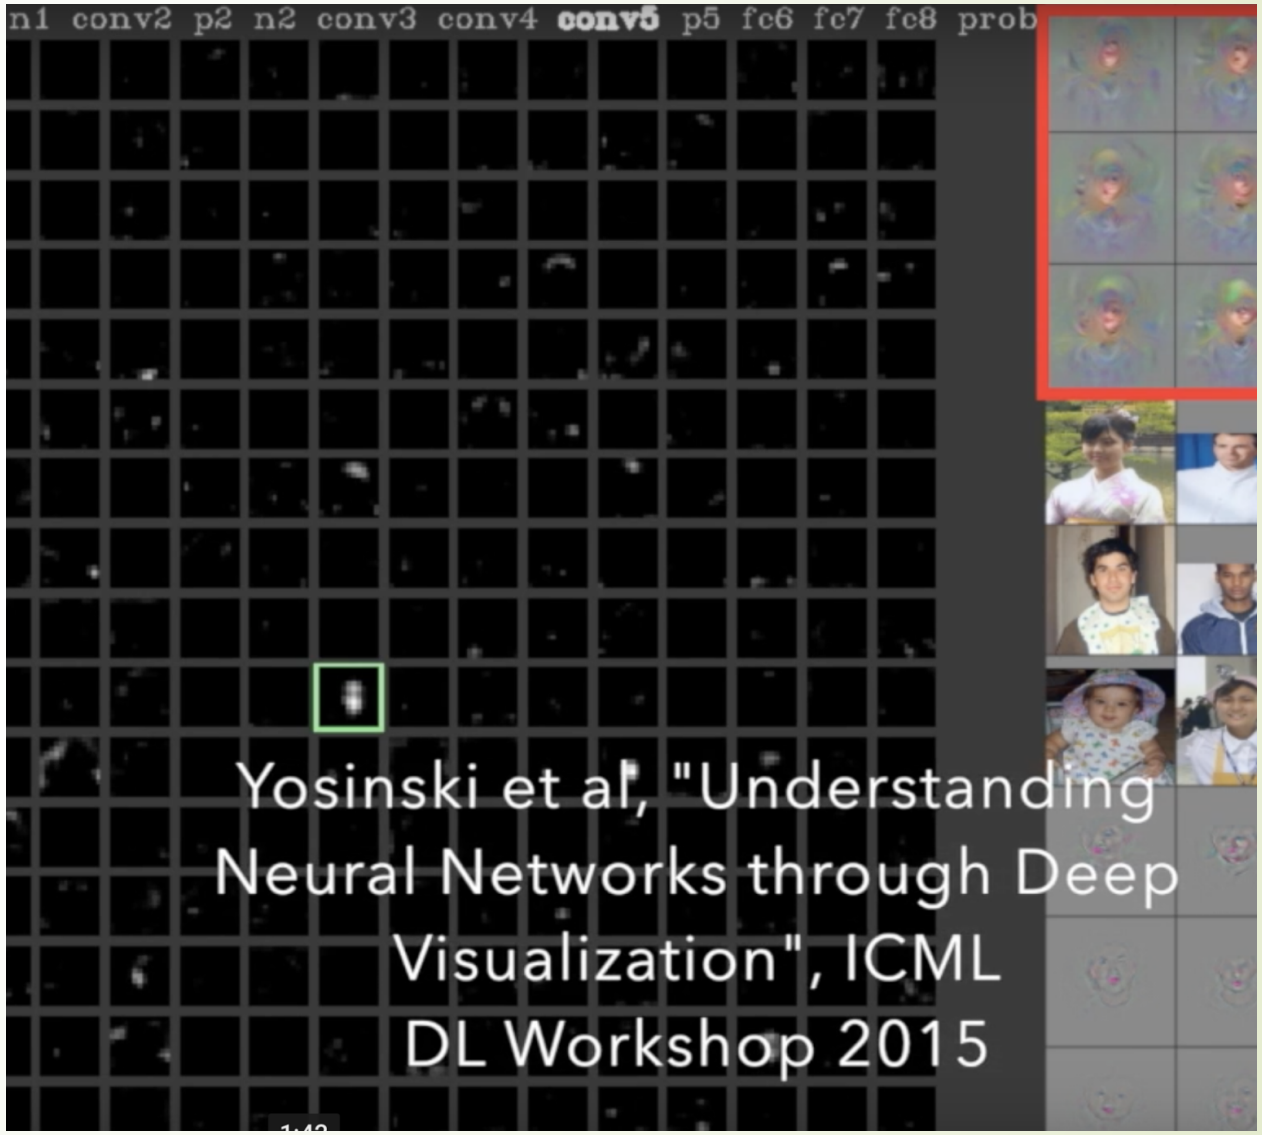

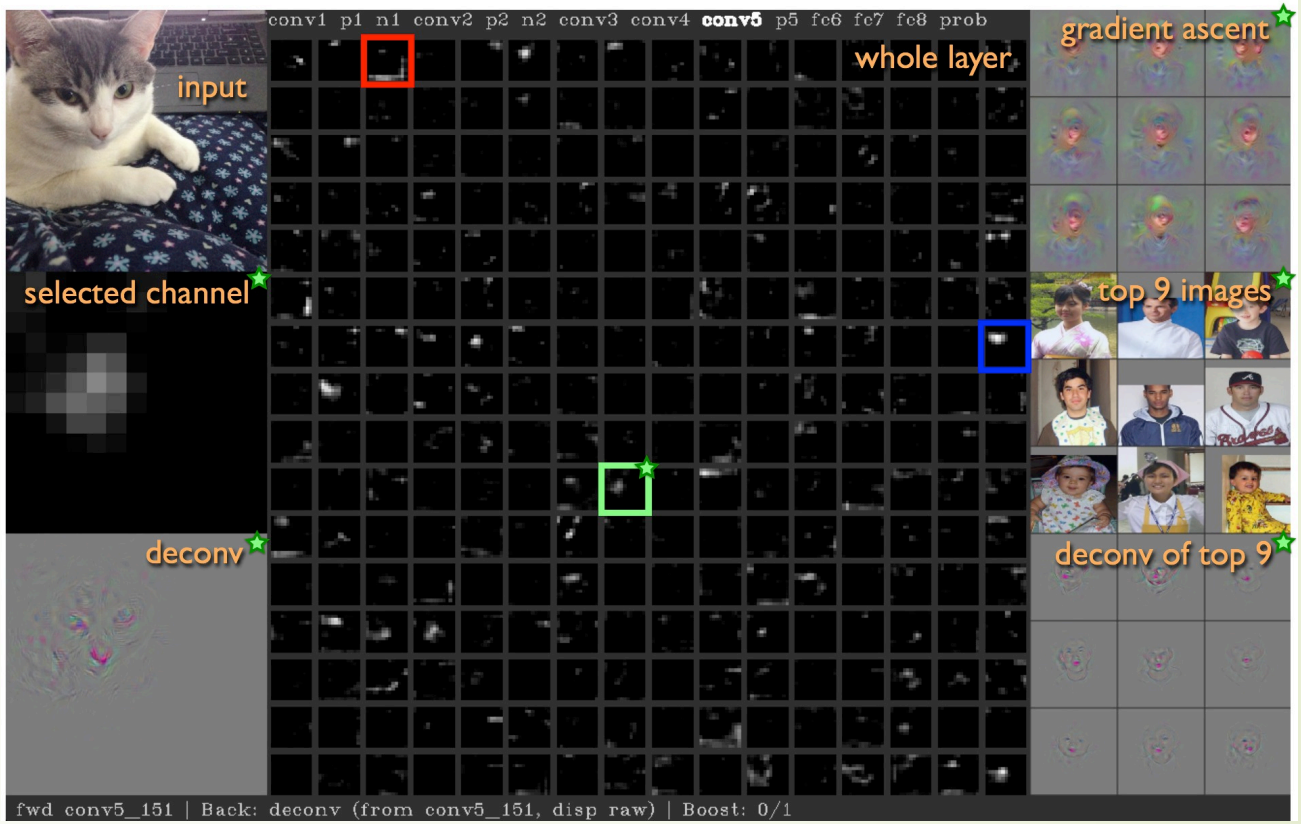

Yosinski et al. Understanding Neural Networks Through Deep Visualization. ICML workshop 2015.

Sparsity in Practice

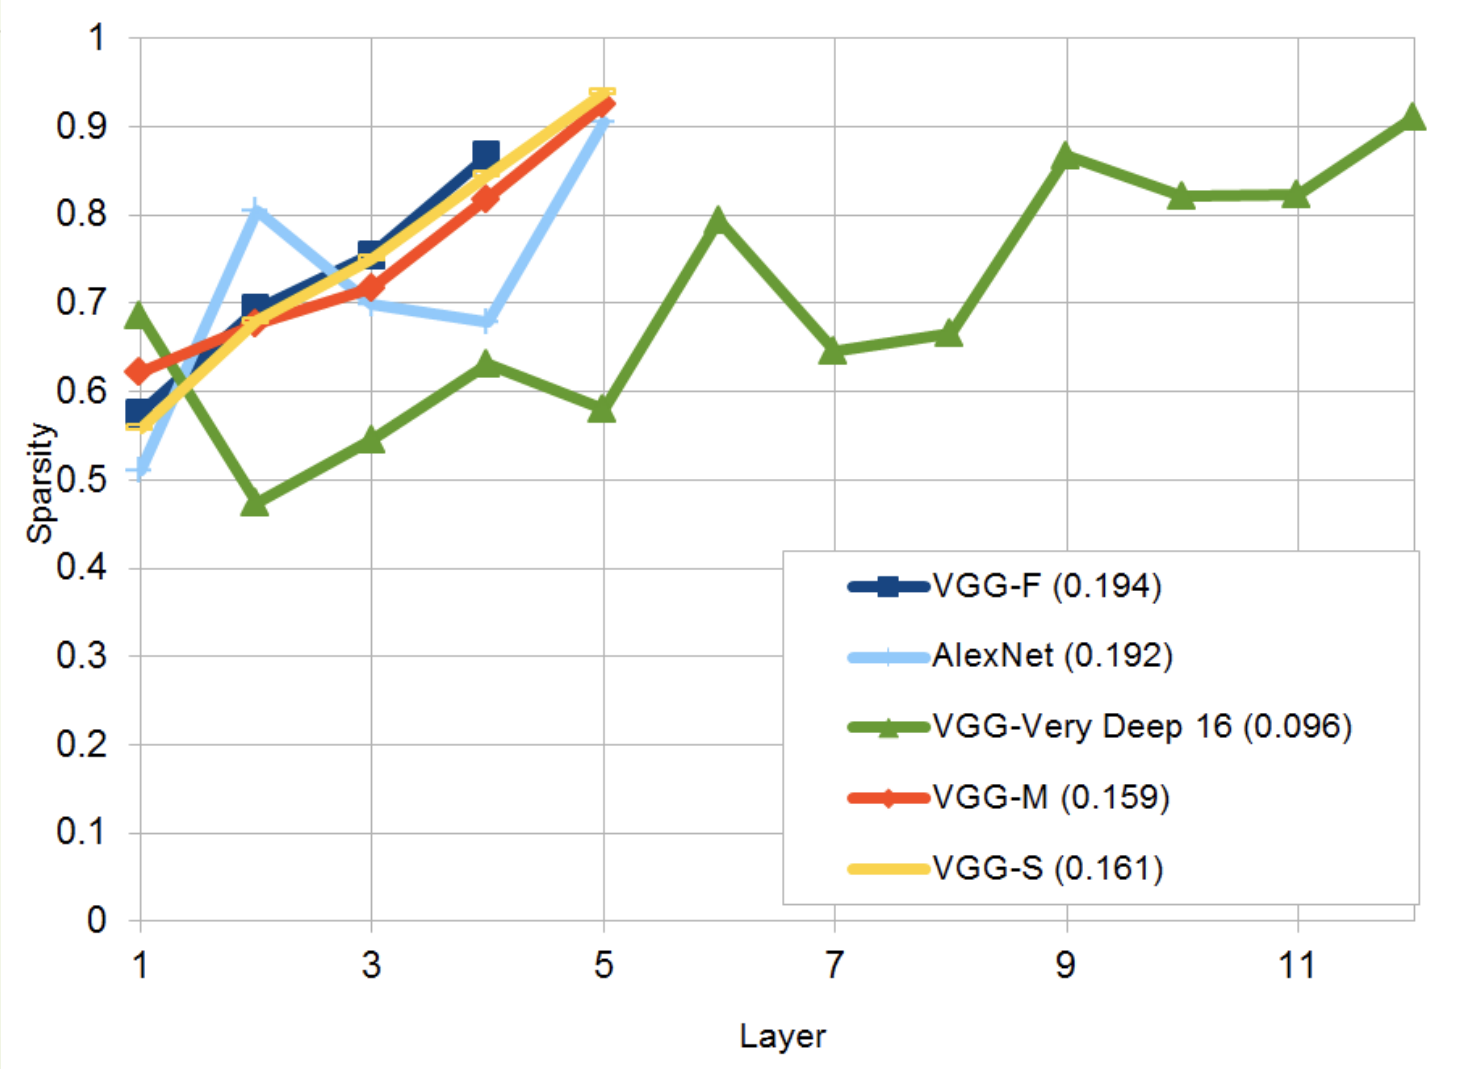

Olshausen & Field and AlexNet

Olshausen & Field

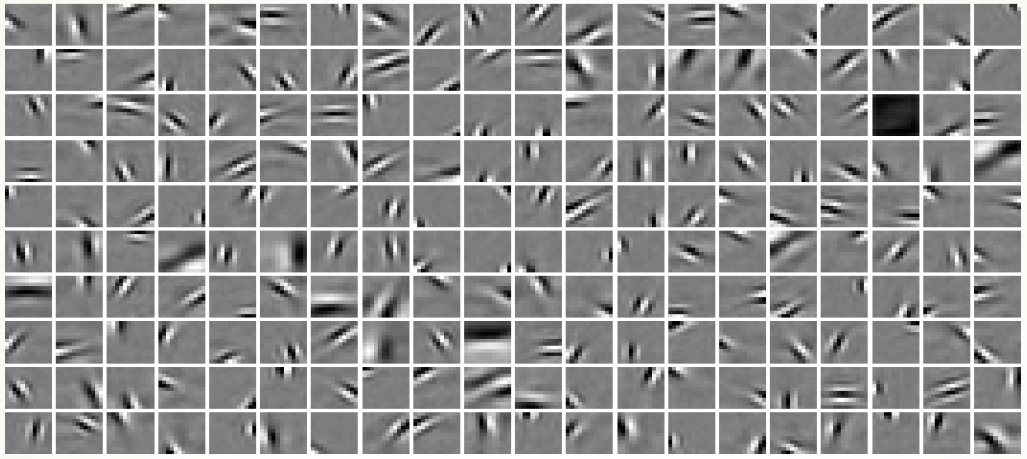

explicit sparsity

AlexNet

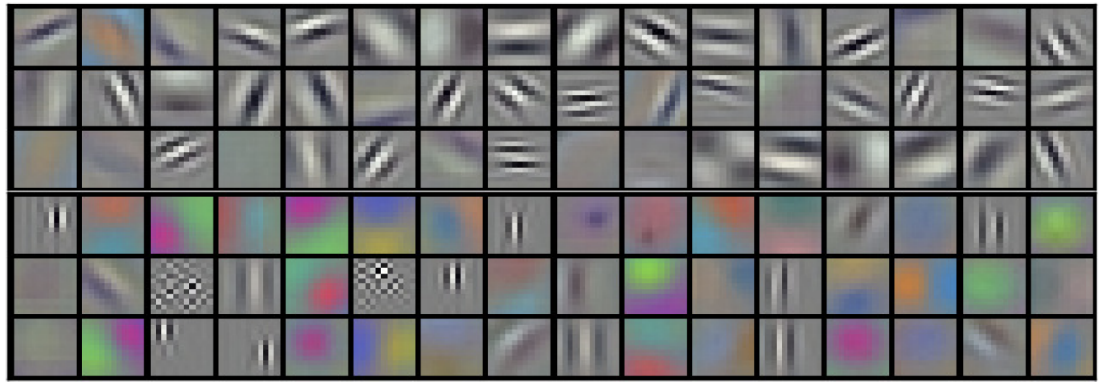

implicit sparsity

Credit to: Vardan, Stats385@Stanford

[Fu, Li, Zhang, Sun, Yao, 2018]

- Pruning in Training, submitted to ICLR, 2018.

Ratio	100%	25%	12.5%	6.25%	3.13%	1.57%
Plain	99.12±0.02	80.46±5.46	62.61±9.05	45.49±0.97	32.34±1.53	21.30±4.83
Rand	99.19±0.01	62.23±10.12	37.71±4.34	23.58±6.96	18.58±4.70	14.36±3.27
R-P	99.16±0.06	75.47±8.11	60.31±5.19	37.97±2.99	26.11±2.13	18.11±1.05
L-P	98.95±0.04	98.95±0.04	90.29±1.30	60.37±5.45	32.91±3.35	20.31±1.10
S-P	98.97±0.07	98.96±0.08	98.67±0.15	68.27±11.22	42.10±7.83	24.95±8.92

(a) Pruning the conv.c5 layer

Ratio	100%	25%	12.5%	6.25%	3.13%	1.57%
Plain	99.10±0.02	96.73±1.05	95.65±1.76	89.60±3.49	78.40±5.48	64.17±6.93
Rand	99.09±0.01	91.56±3.89	71.05±6.08	51.92±8.87	33.65±6.17	29.91±8.49
R-P	99.13±0.05	96.39±0.48	95.31±0.79	91.29±3.46	82.75±5.59	68.35±6.29
L-P	99.10±0.03	98.89±0.05	98.89±0.05	98.89±0.05	98.89±0.05	98.89±0.05
S-P	99.13±0.03	98.73±0.10	98.61±0.15	98.23±0.40	96.75±0.88	92.53±3.17

(a) Pruning the fc.c6 layer

(%)	100	50	25	12.5	6.25	3.13	1.57
Plain	83.92	13.53	8.12	5.32	5.29	5.92	6.31
Rand	82.36	17.90	6.52	6.38	7.90	9.58	8.67
R-P	82.79	13.72	7.10	6.38	6.29	6.52	5.76
L-P / GL-P	81.09	81.09	76.43	75.06	68.42	55.25	33.49
S-P / GS-P	78.95	77.81	73.92	70.65	68.67	67.58	*65.17*
Com-Rat(%)	100	63.00	44.51	35.26	30.64	28.33	27.17

Table 5: Top 5 accuracy on *mini*Imagenet by pruning ResNet-18, the fully connected layer, Block#4.0 and #4.1 layers.

Also Andrea Madotto et al. **Low-Rank Matrix Factorization of LSTM as Effective Model Compression**

Theory of Sparsity in CNN?

Breiman's "Two Cultures"

Generative modeling

Predictive modeling

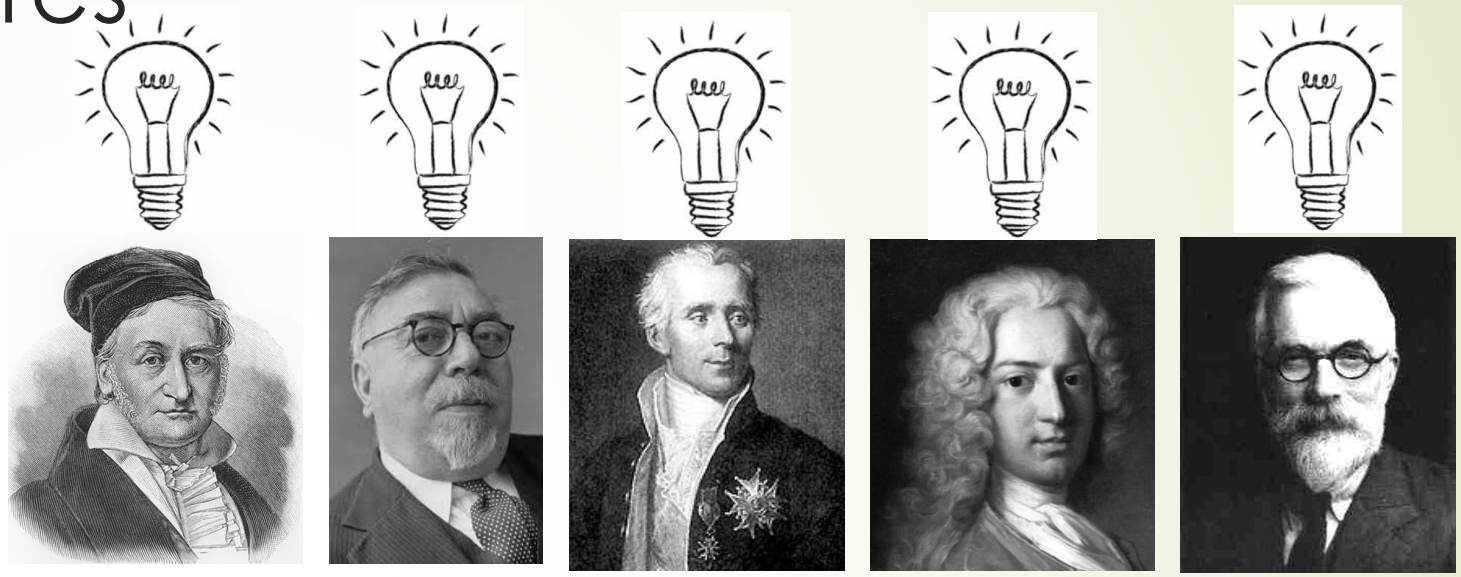

Gauss Wiener Laplace Bernoulli Fisher

Generative modeling

Seeks to develop stochastic models which fit the data, and then make inferences about the data-generating mechanism based on the structure of those models.

Predictive modeling

Predictive modeling is effectively silent about the underlying mechanism generating the data, and allows for many different predictive algorithms, preferring to discuss only accuracy of prediction made by different algorithm on various datasets.

Generative Modeling

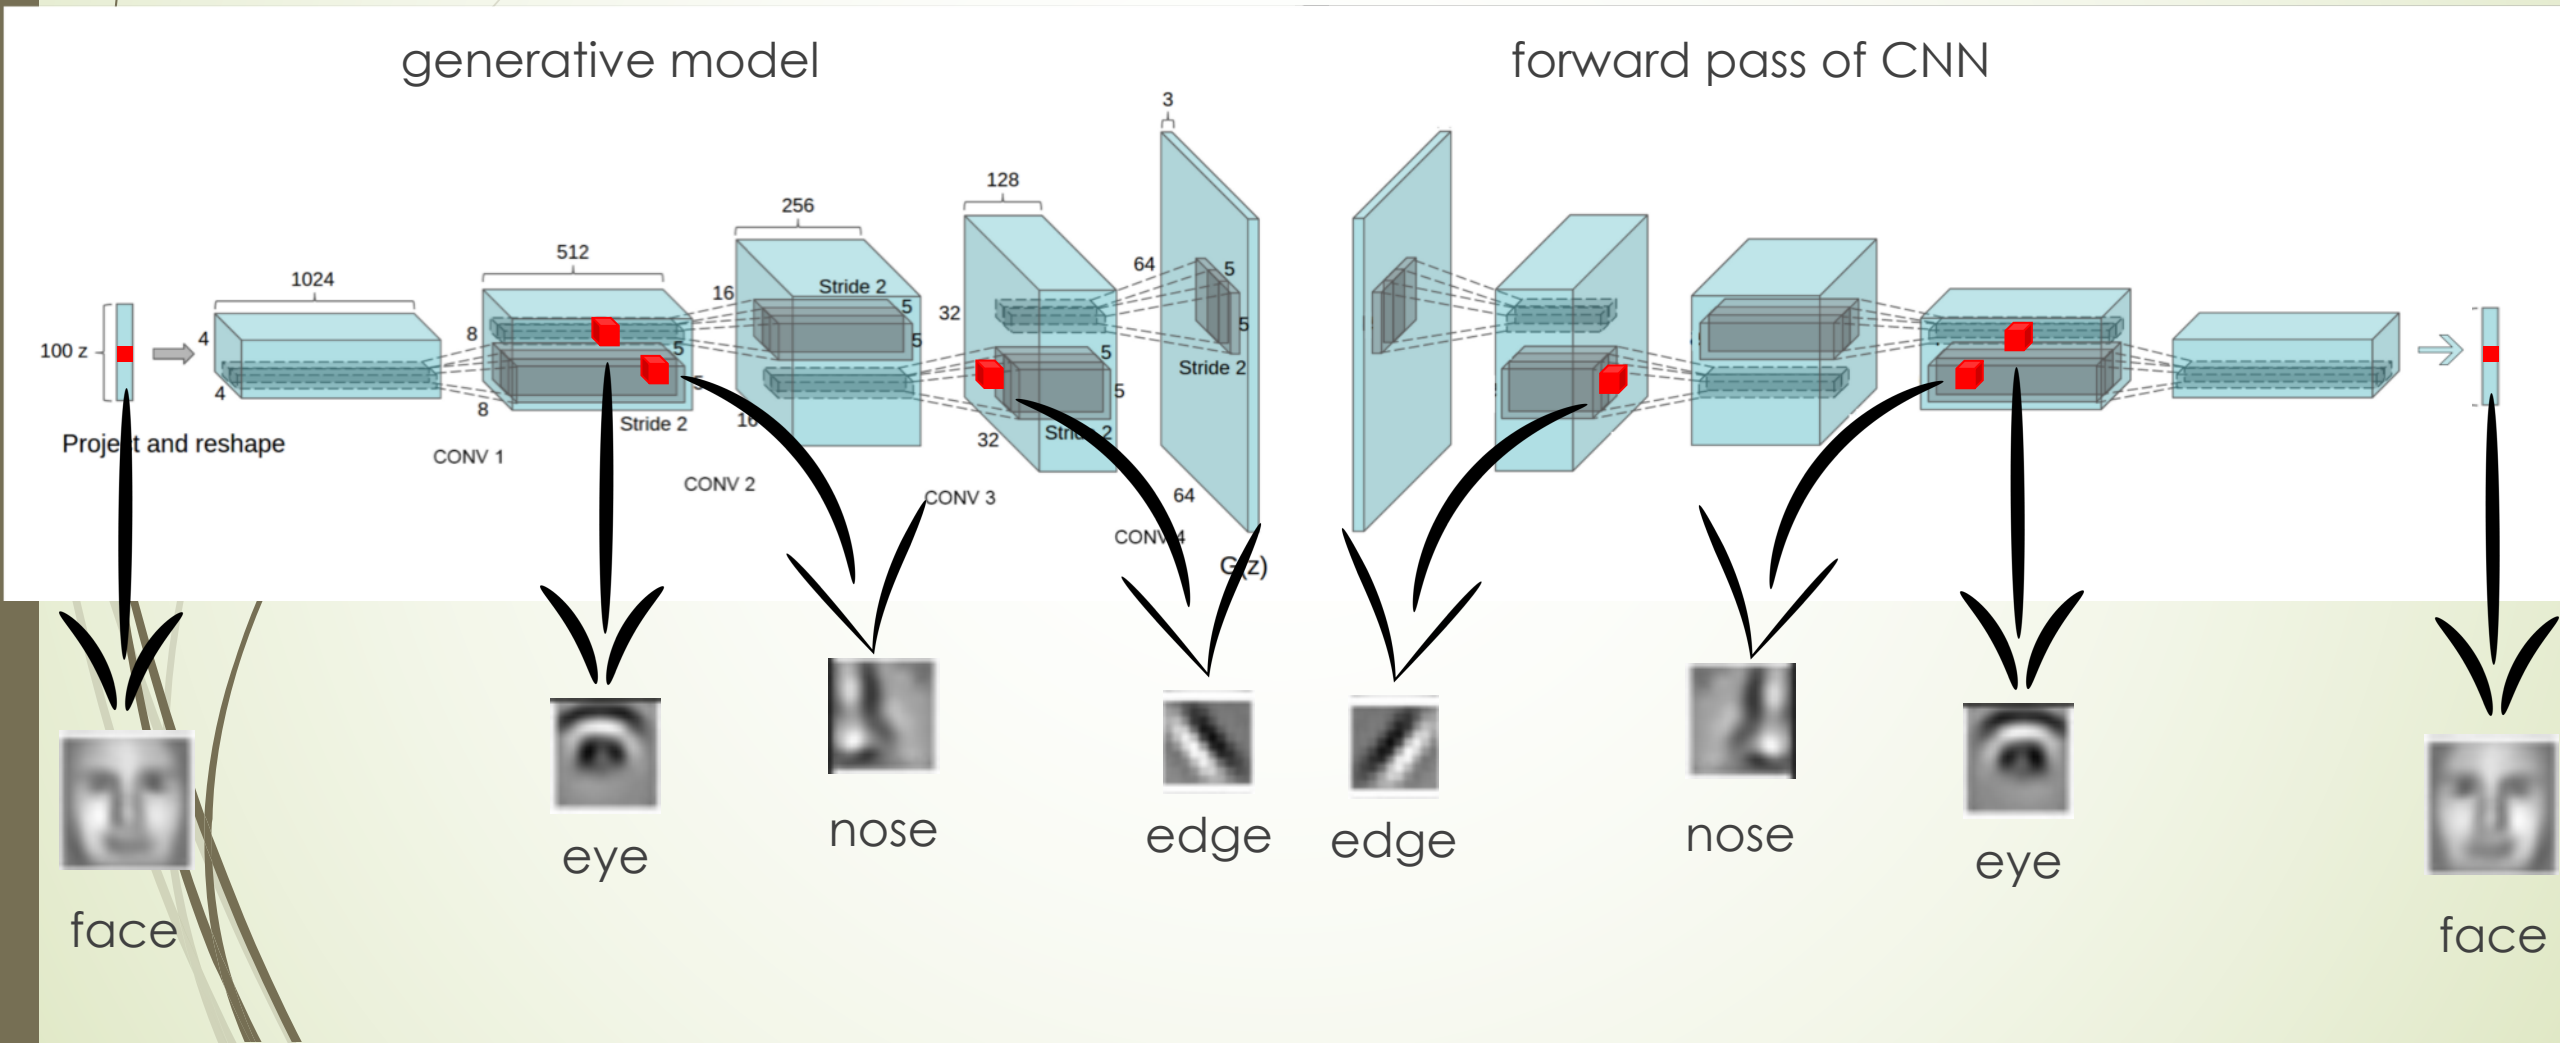

Sparse Representation Generative Model

- Receptive fields in visual cortex are spatially localized, oriented and bandpass
- Coding natural images while promoting sparse solutions results in a set of filters satisfying these properties [Olshausen and Field 1996]
- Two decades later…
  - vast theoretical study
  - different inference algorithms
  - different ways to train the model

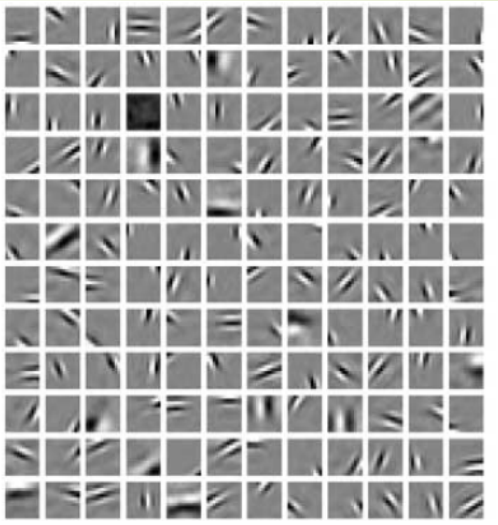

Evolution of Models

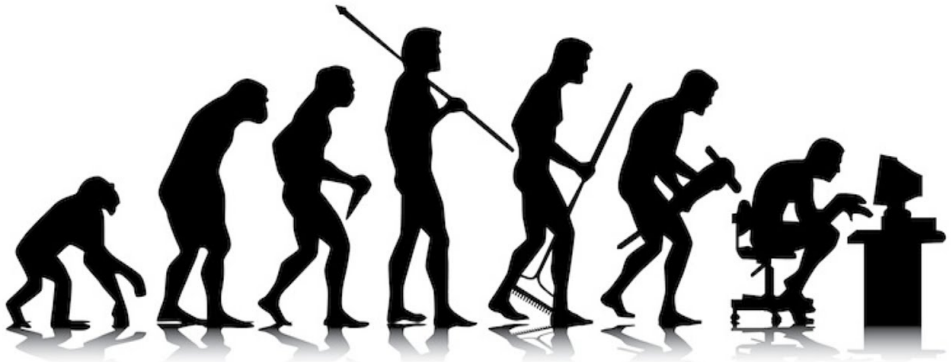

Sparse representations

Convolutional sparse representation

Multi-Layered Convolutional Sparse Representation

First Layer of a Convolutional Neural Network

First Layer of a Neural Network

Multi-Layered Convolutional Neural Network

First Layer of a Neural Network

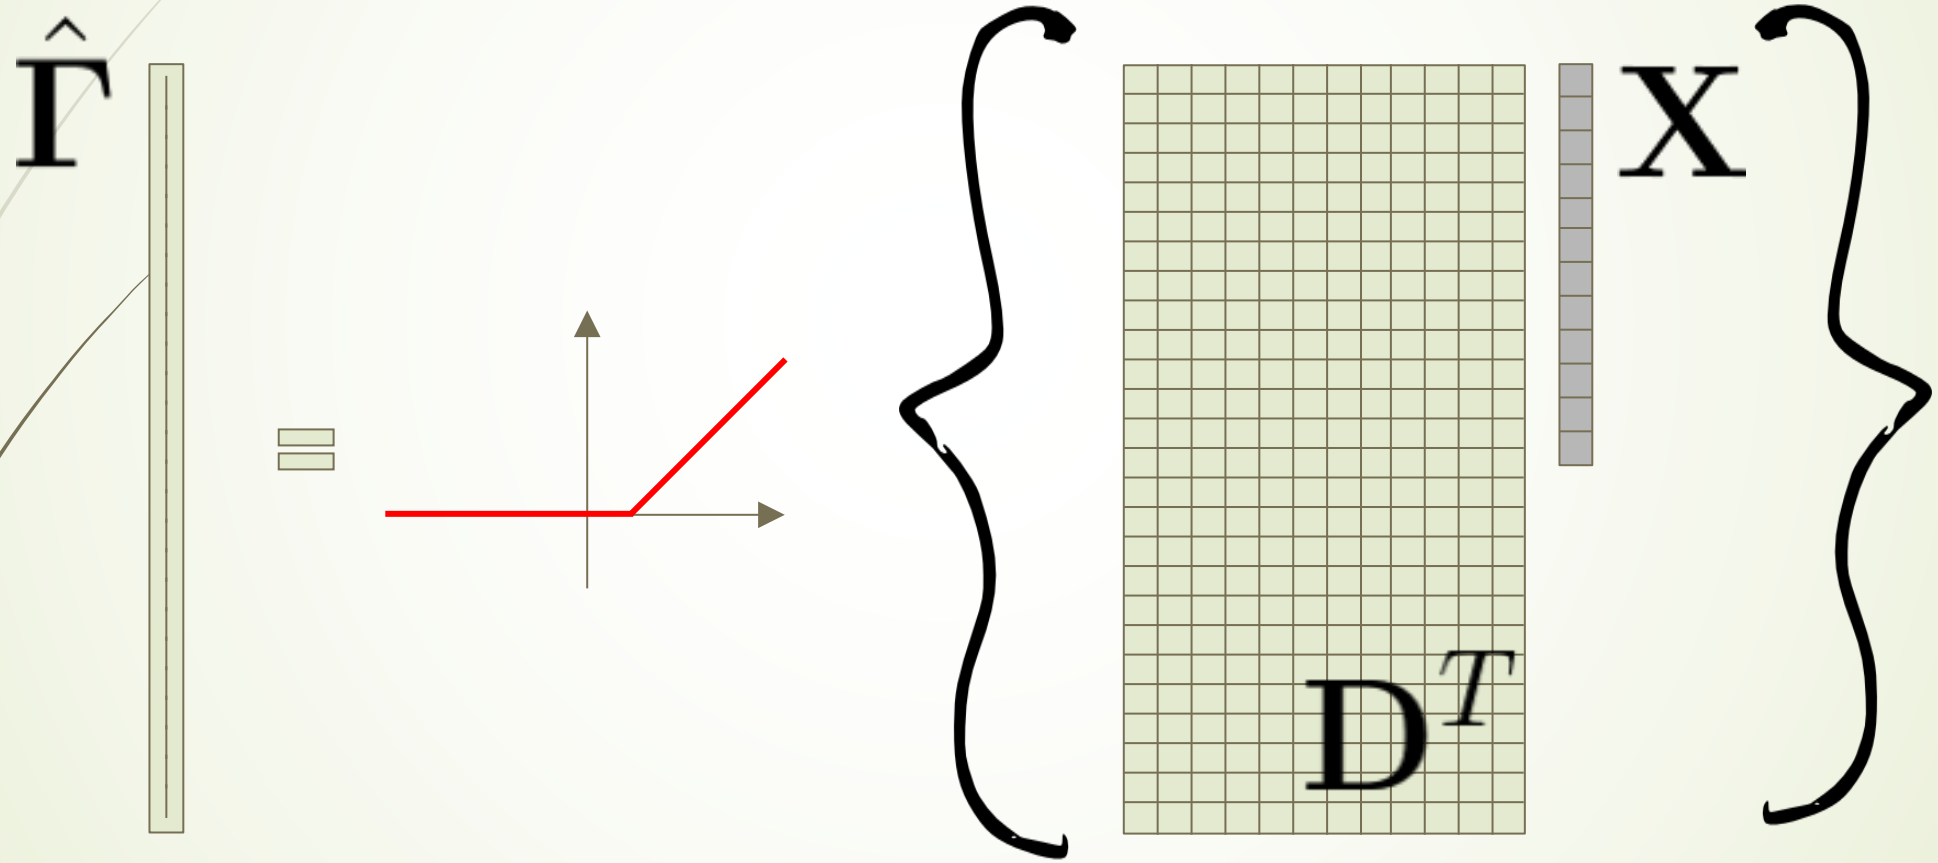

Sparse Modeling, e.g.

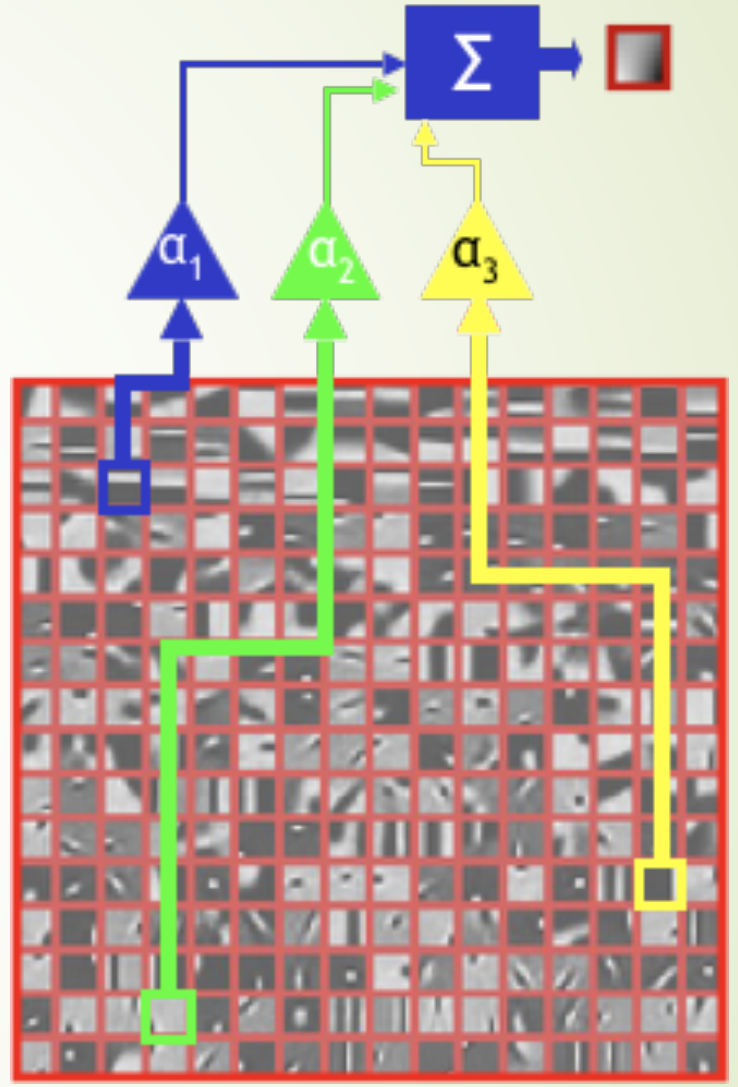

Task: model image patches of size 8x8 pixels

We assume a dictionary of such image patches is given, containing 256 atoms (overcomplete)

Assumption: every patch can be described as a linear combination of a few atoms

Key properties: sparsity and redundancy

Matrix Notation

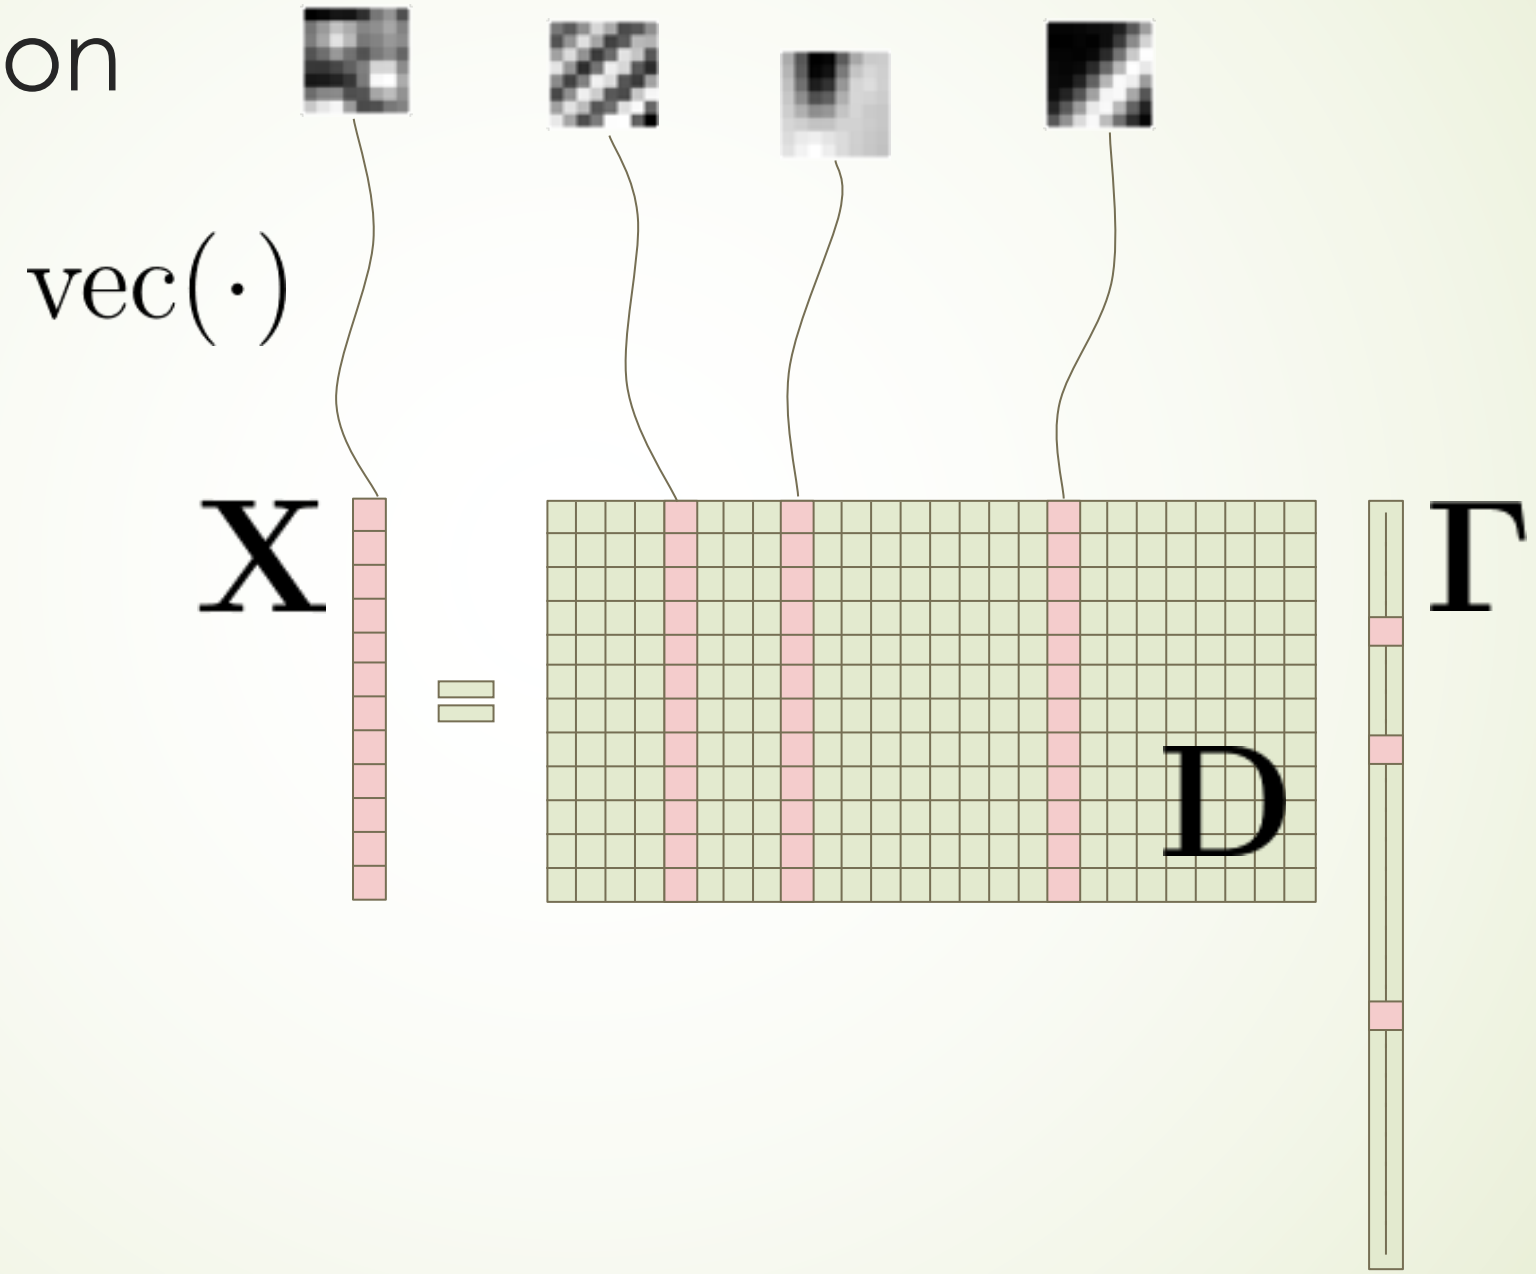

Sparse Coding

Given a signal, we would like to find its sparse representation

Convexify

$$\min_{\mathbf{\Gamma}} \|\mathbf{\Gamma}\|_0 \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\mathbf{\Gamma}$$

$$\min_{\mathbf{\Gamma}} \|\mathbf{\Gamma}\|_1 \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\mathbf{\Gamma}$$

Sparse Coding

Given a signal, we would like to find its sparse representation

$$\min_{\mathbf{\Gamma}} \|\mathbf{\Gamma}\|_0 \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\mathbf{\Gamma}$$

Convexify

$$\min_{\mathbf{\Gamma}} \|\mathbf{\Gamma}\|_1 \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\mathbf{\Gamma}$$

Crude approximation

$$\mathcal{S}_\beta\{\mathbf{D}^T\mathbf{X}\}$$

Soft-Thresholding Algorithm

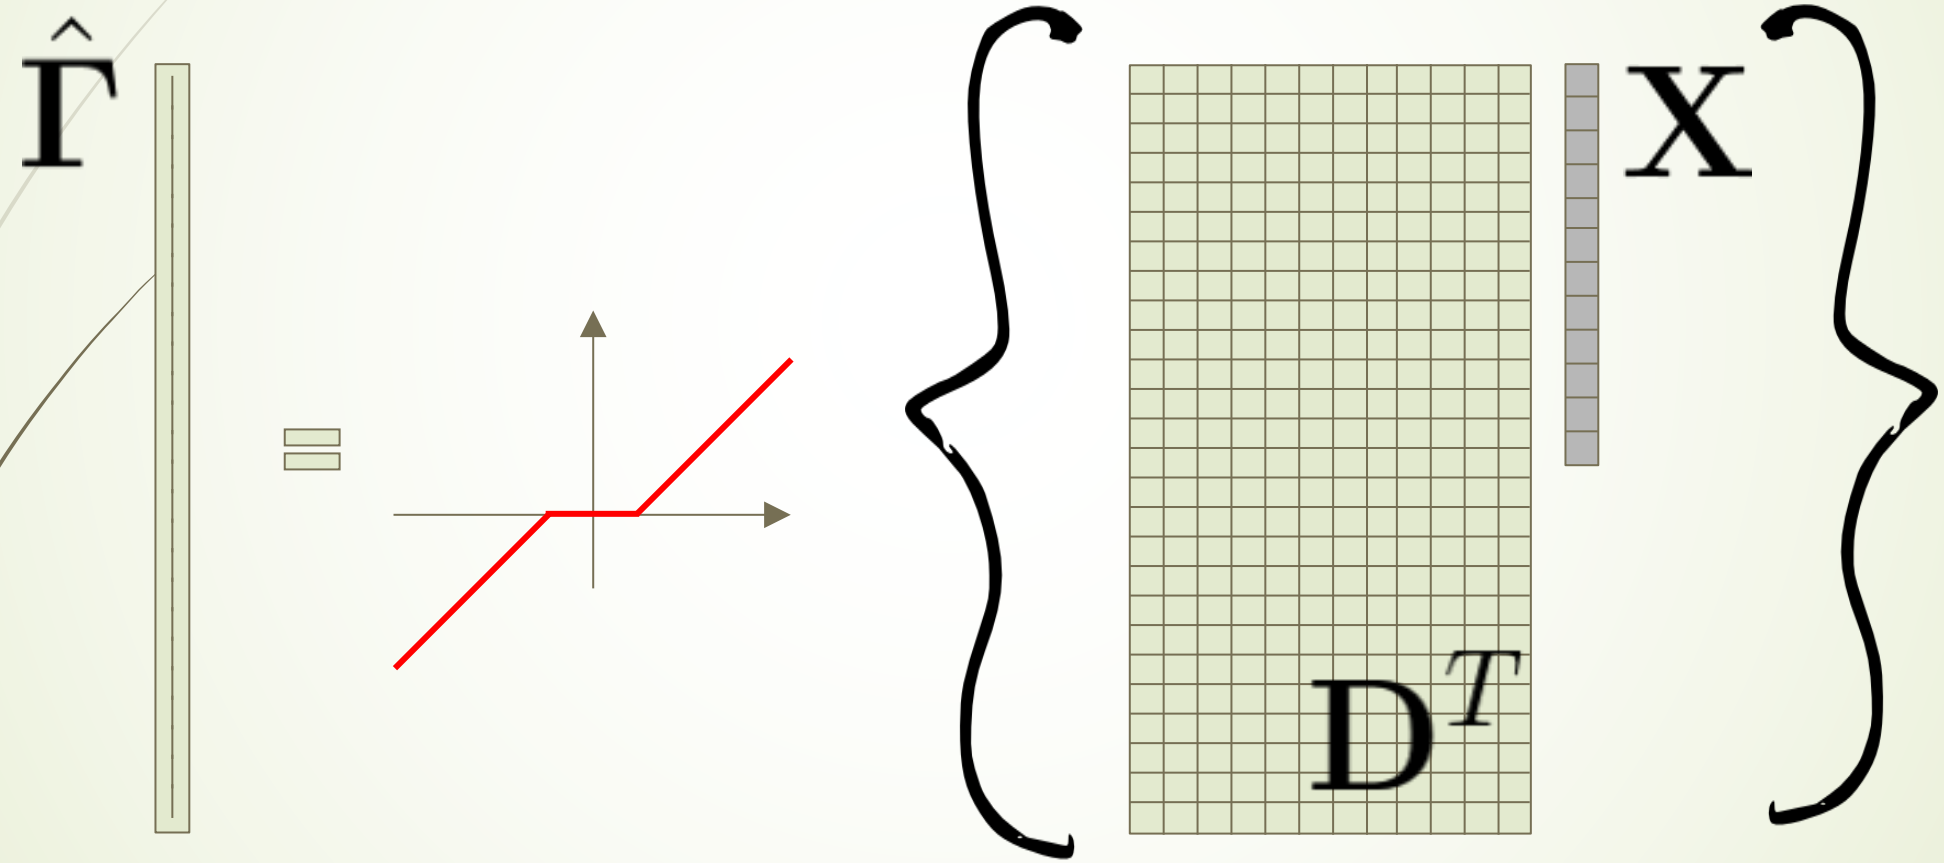

First Layer of a Neural Network

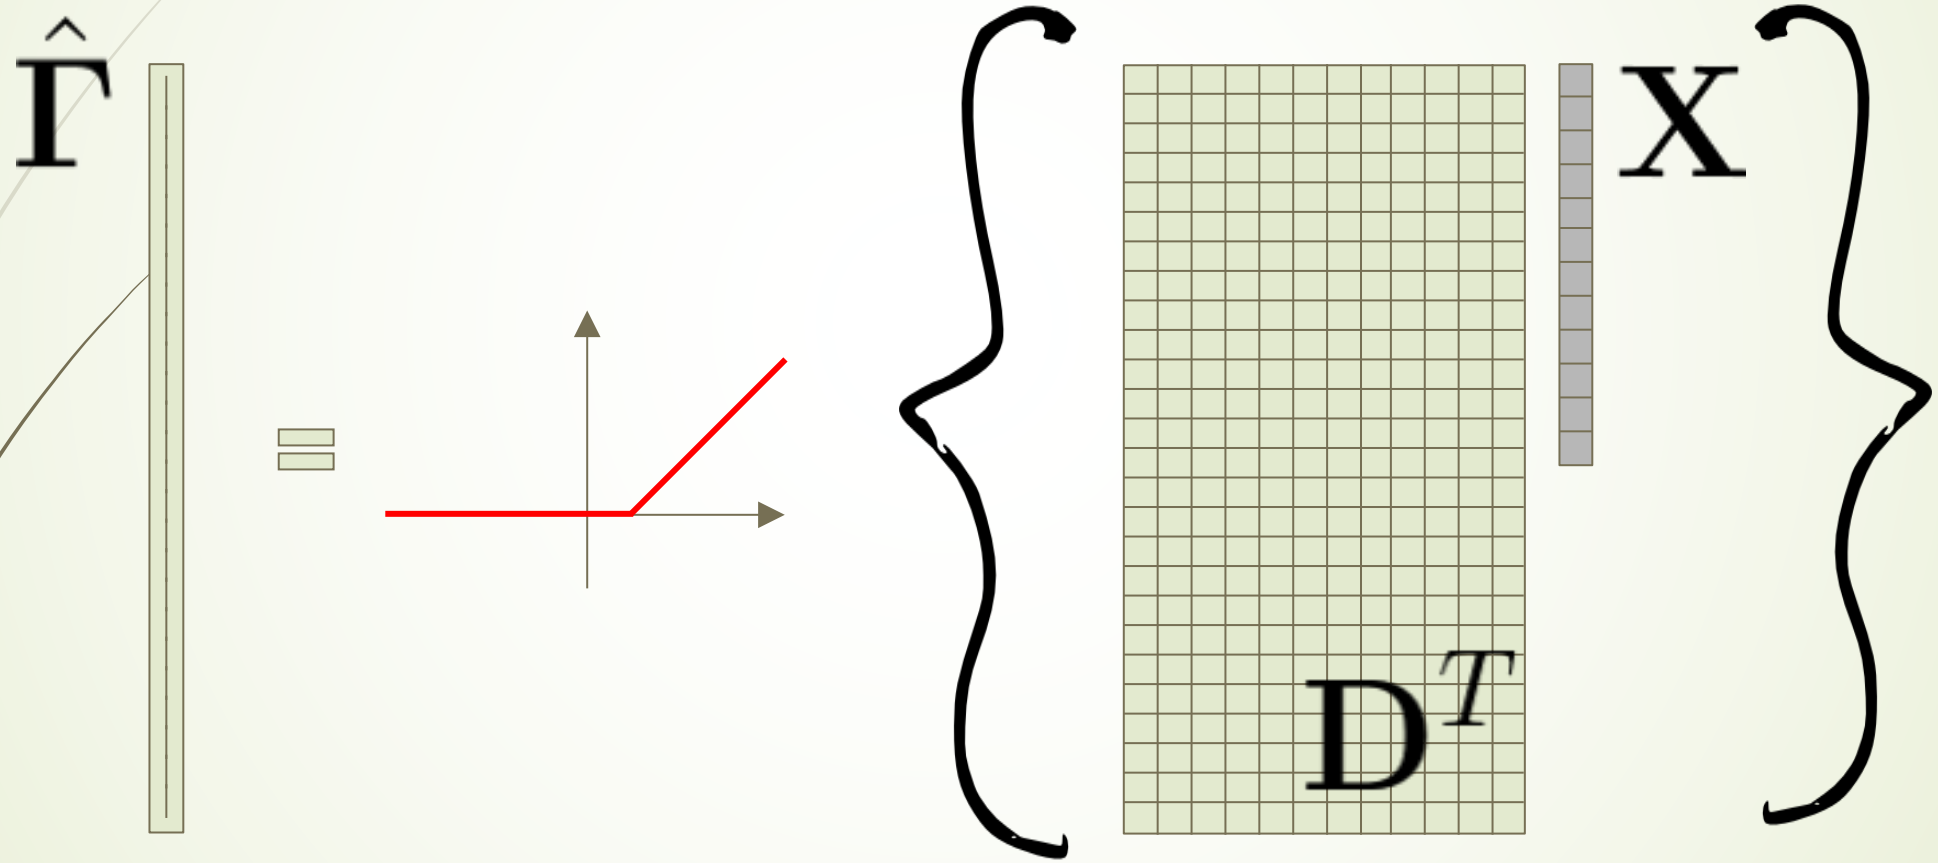

ReLU = Soft Nonnegative Thresholding

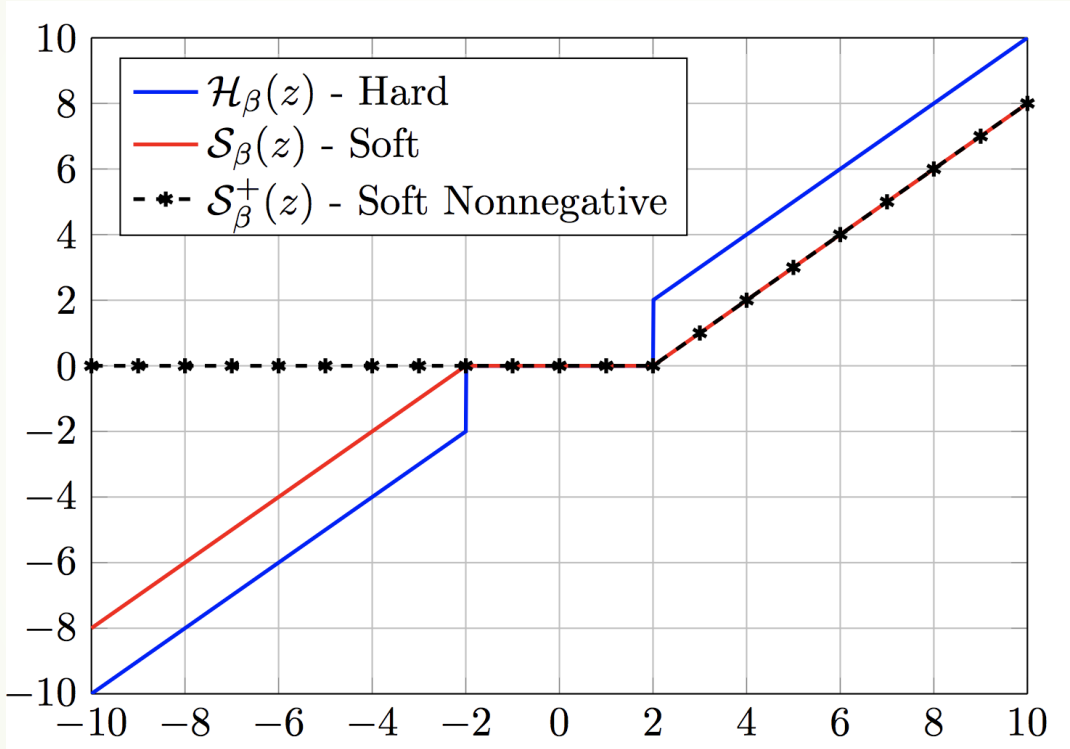

ReLU is equivalent to soft nonnegative thresholding

First layer of a **Convolutional** Neural Network

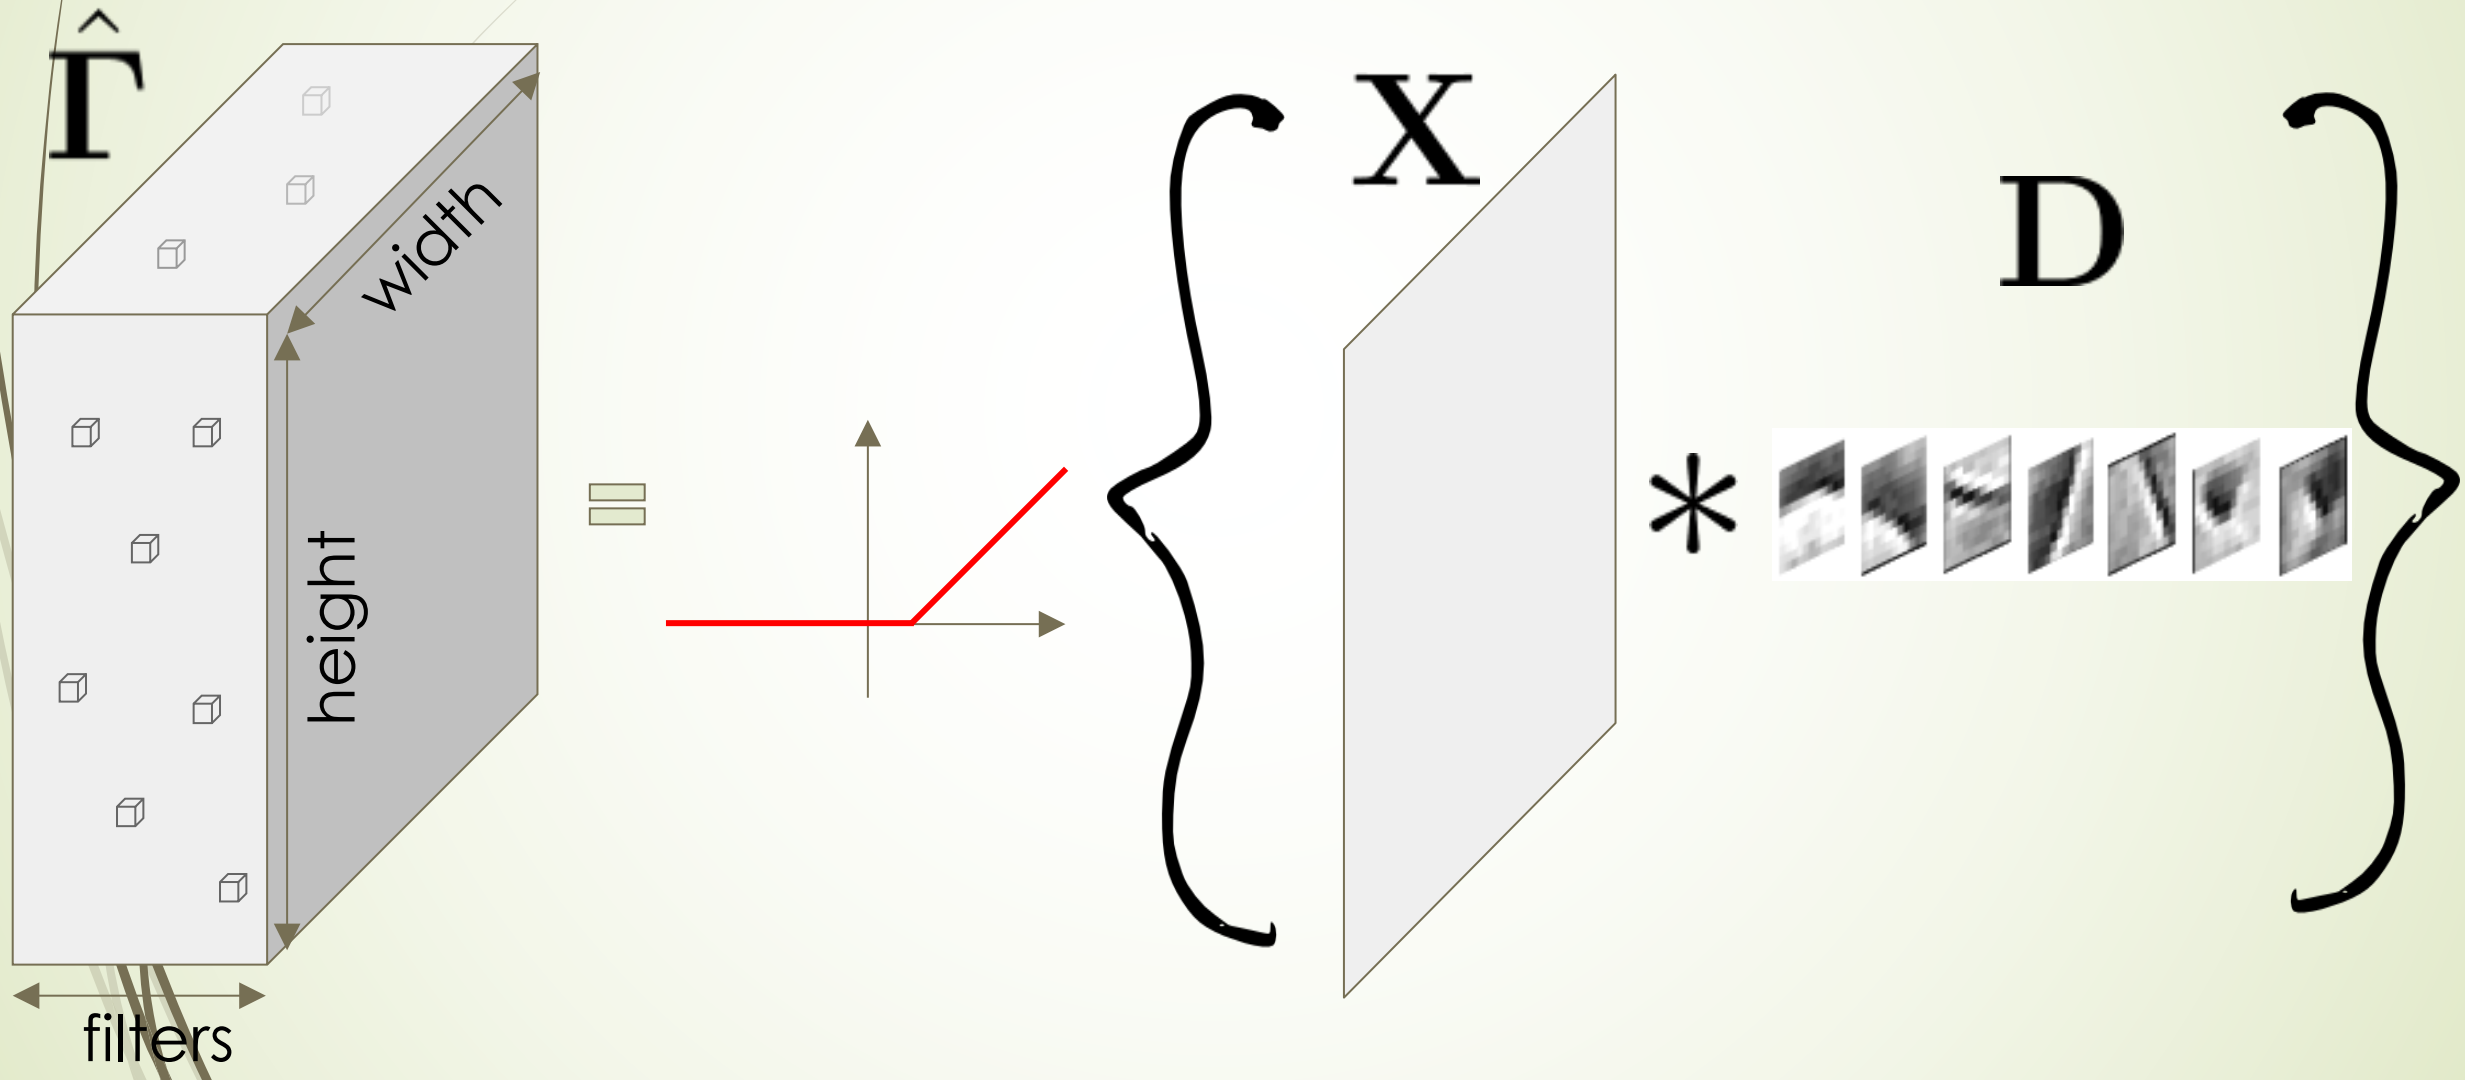

Convolutional Sparse Modeling

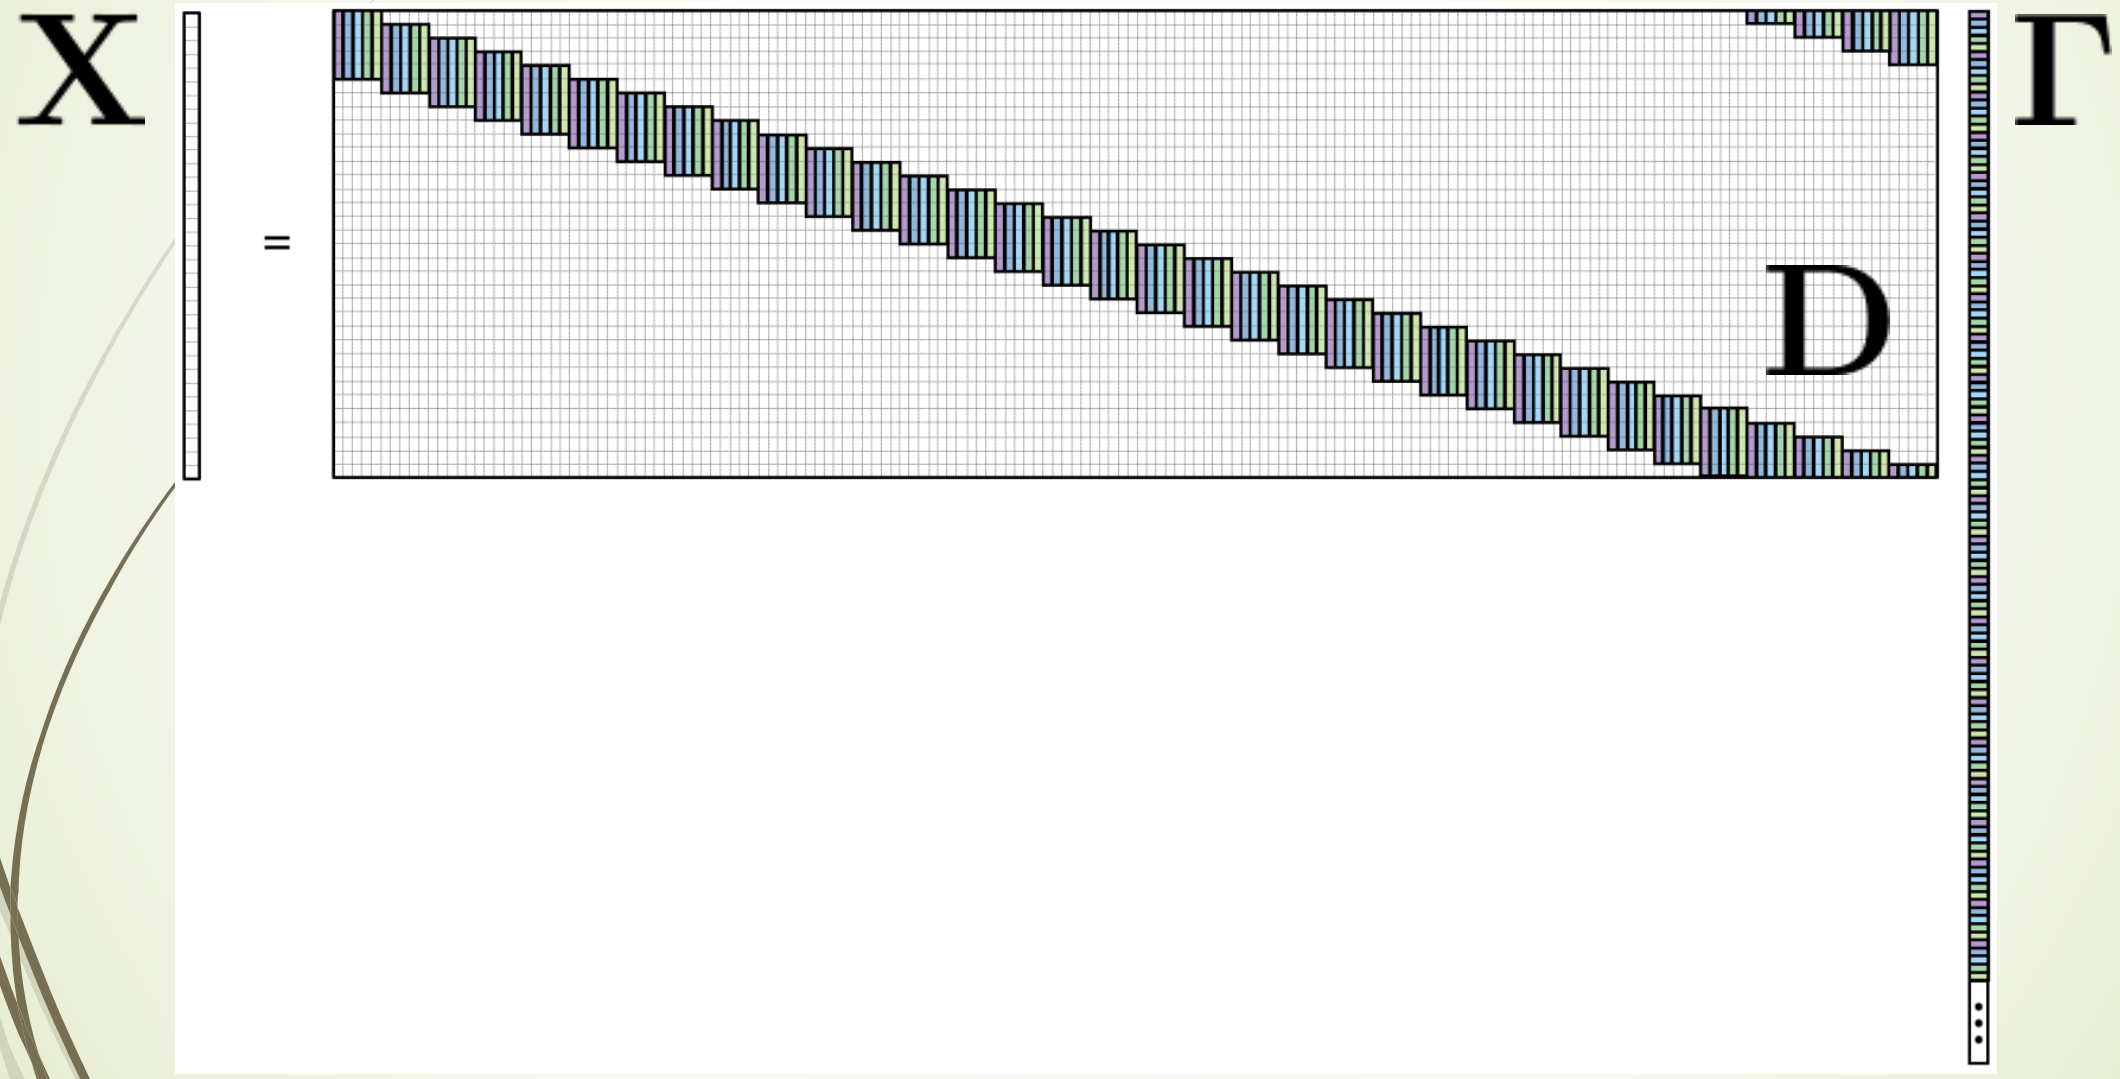

Convolutional Sparse Modeling

$$\mathbf{X} = \sum_{k=1}^{K} \mathbf{d}_k * \mathbf{z}_k$$

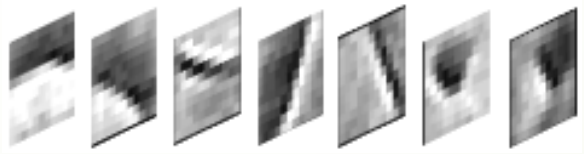

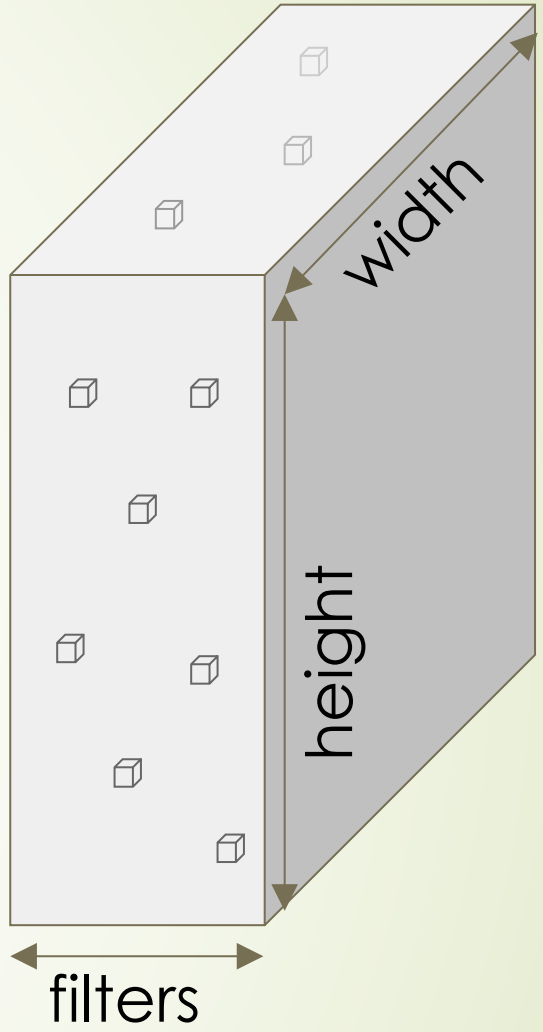

Convolutional Sparse Modeling

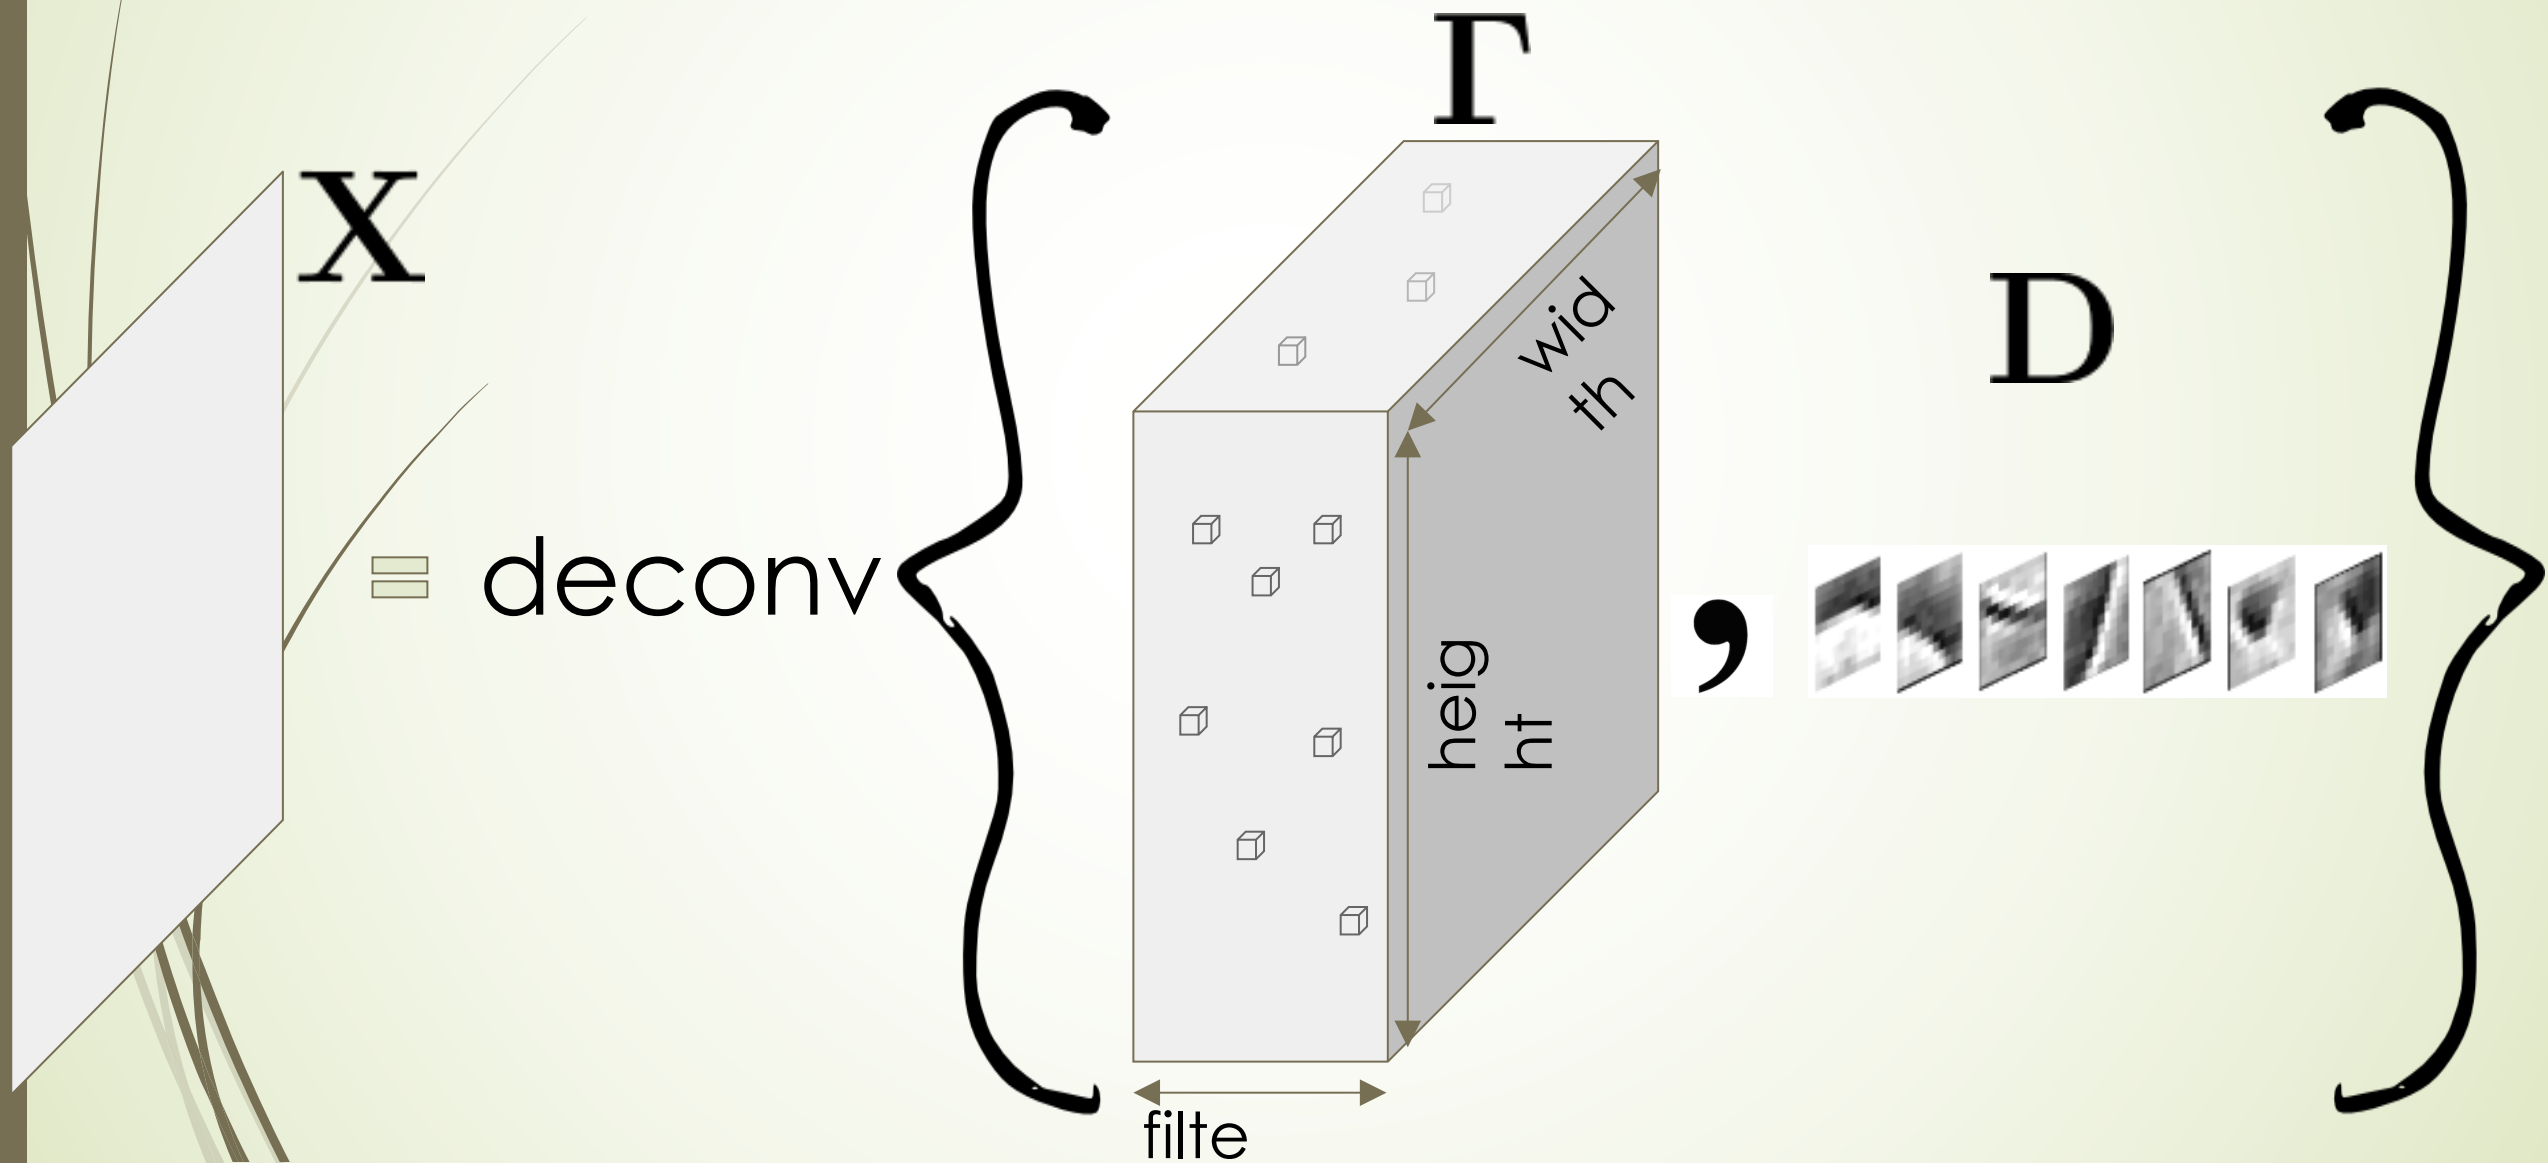

Thresholding Algorithm

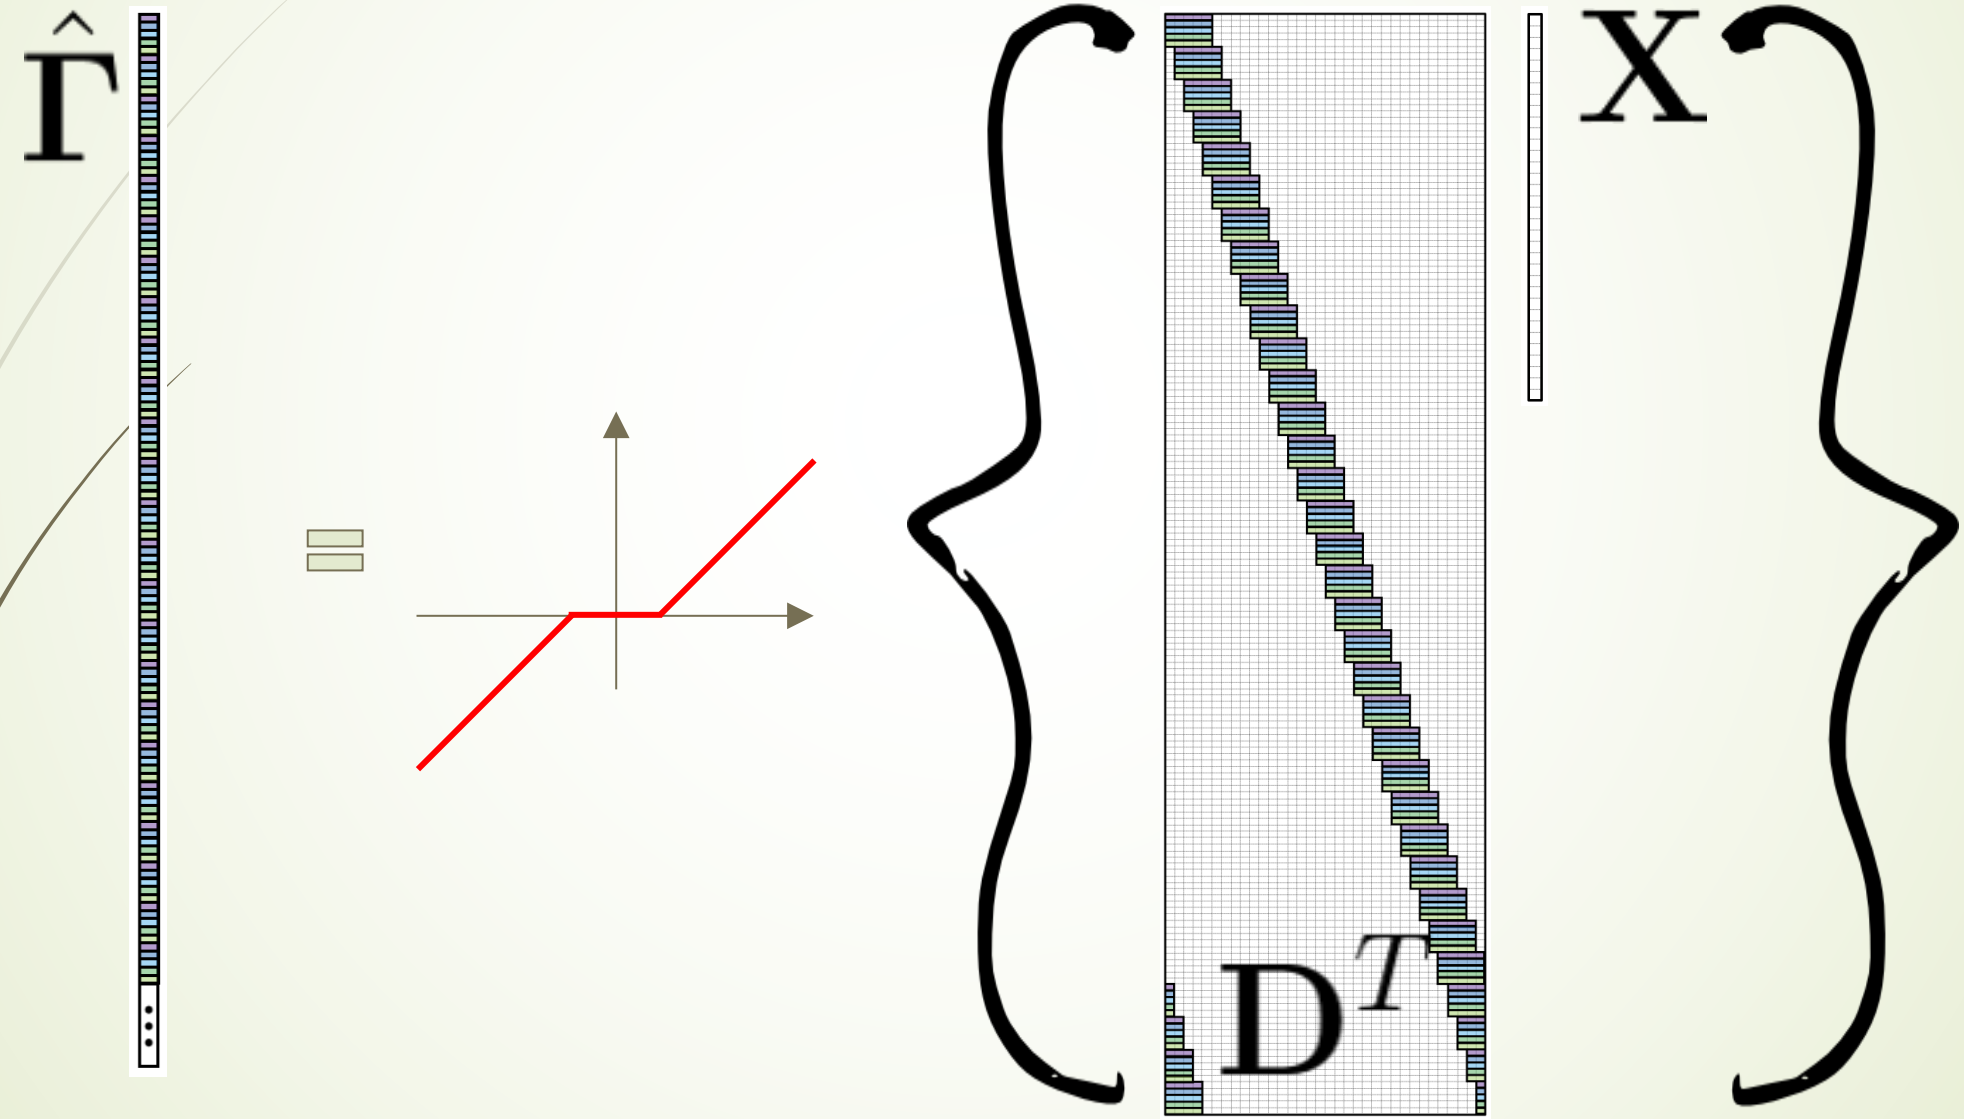

First layer of a Convolutional Neural Network

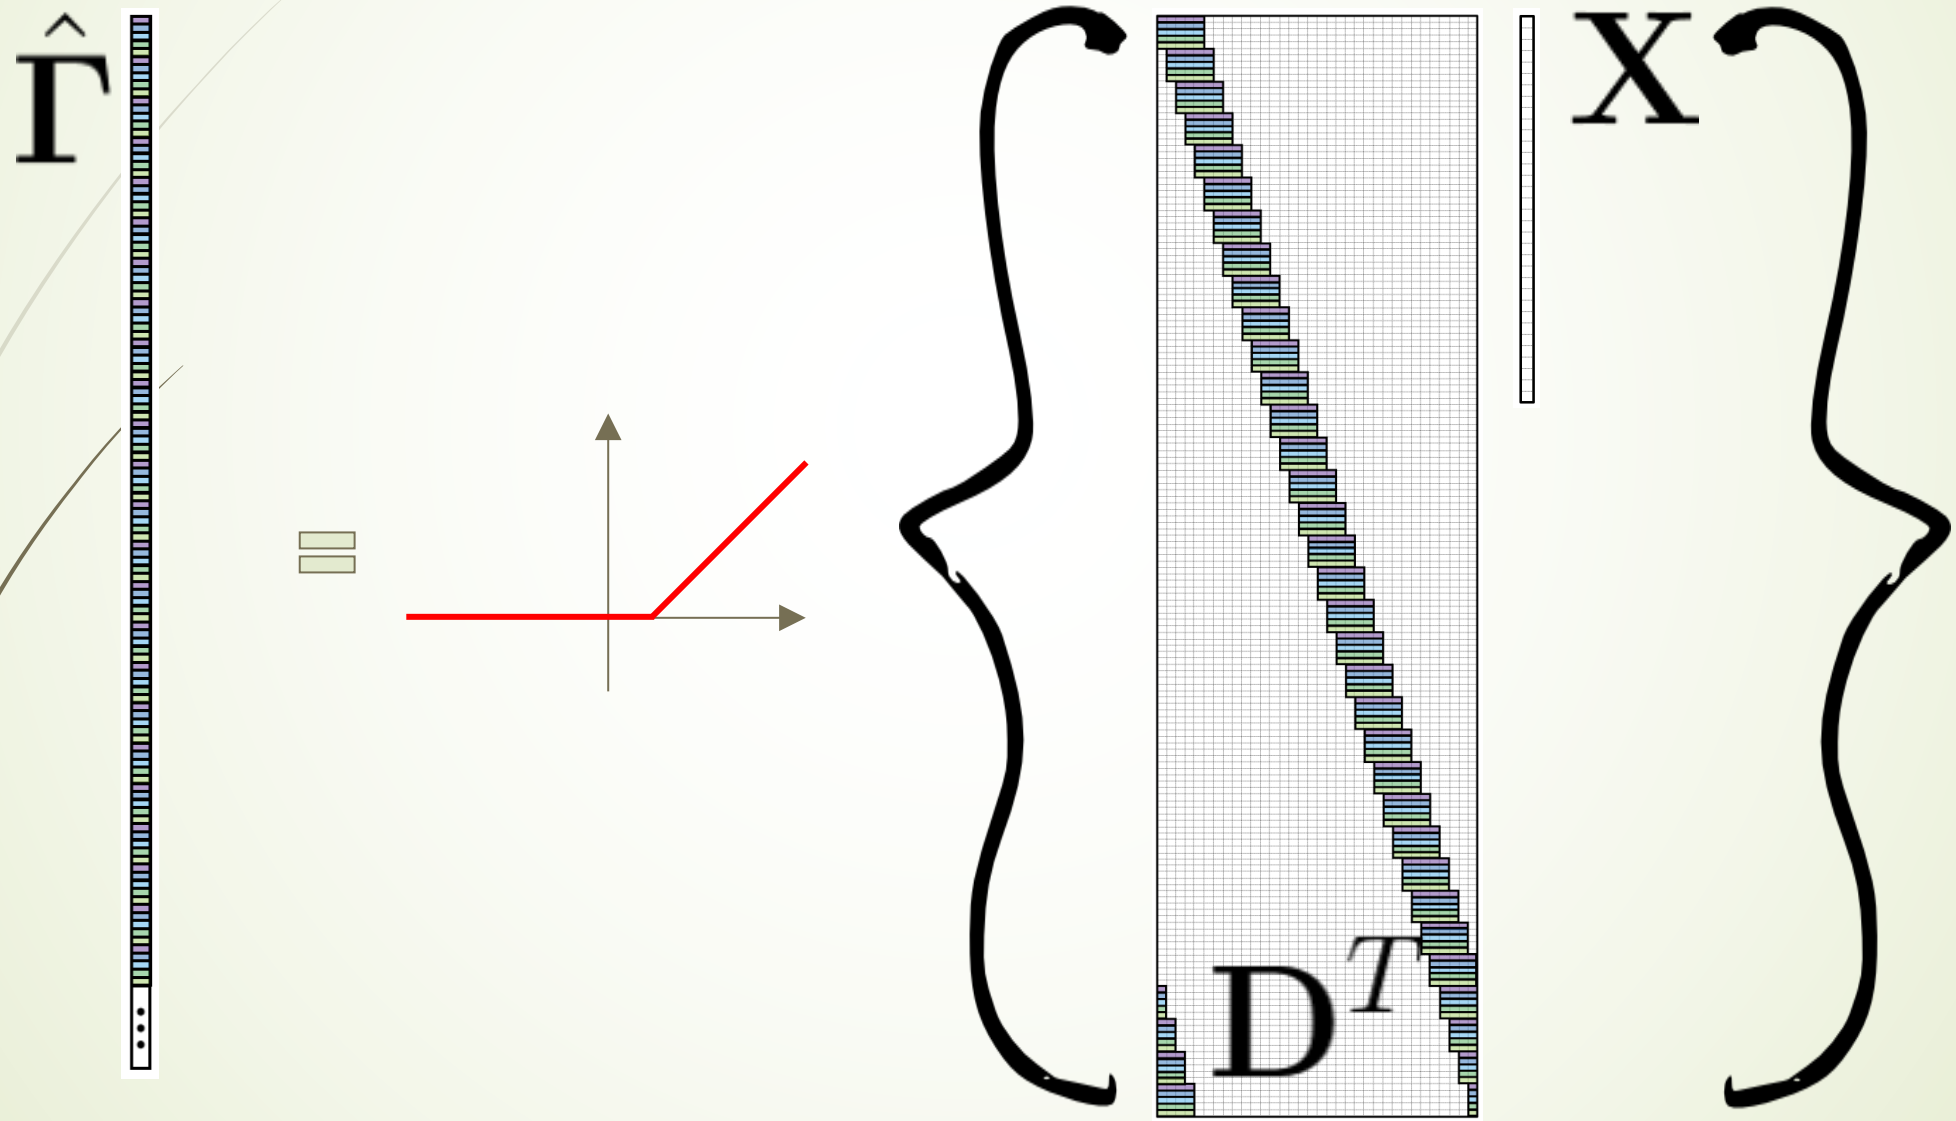

First layer of a Convolutional Neural Network

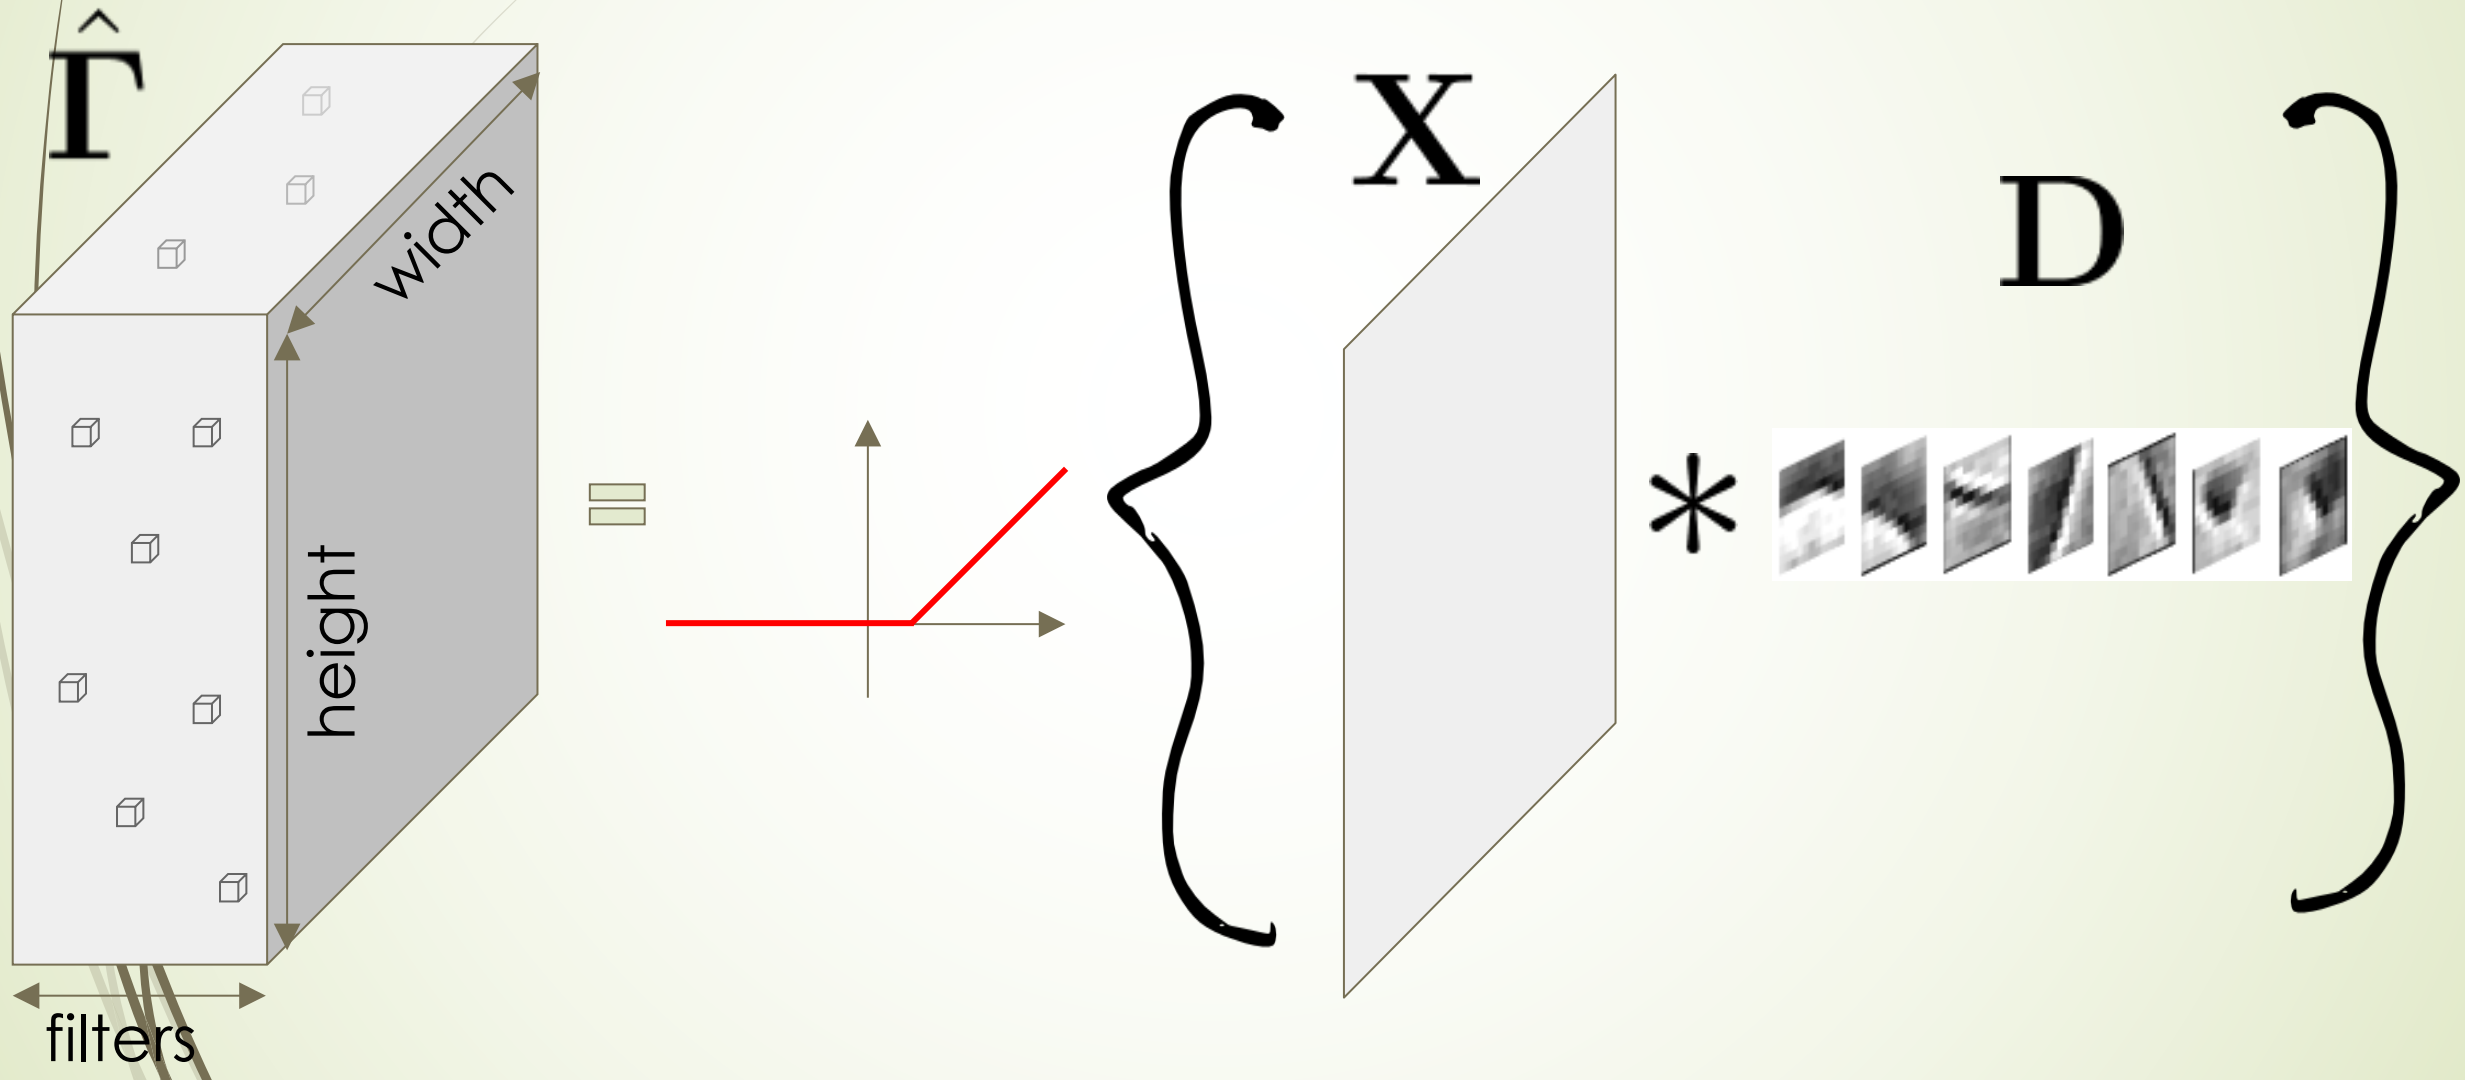

Convolutional Neural Network

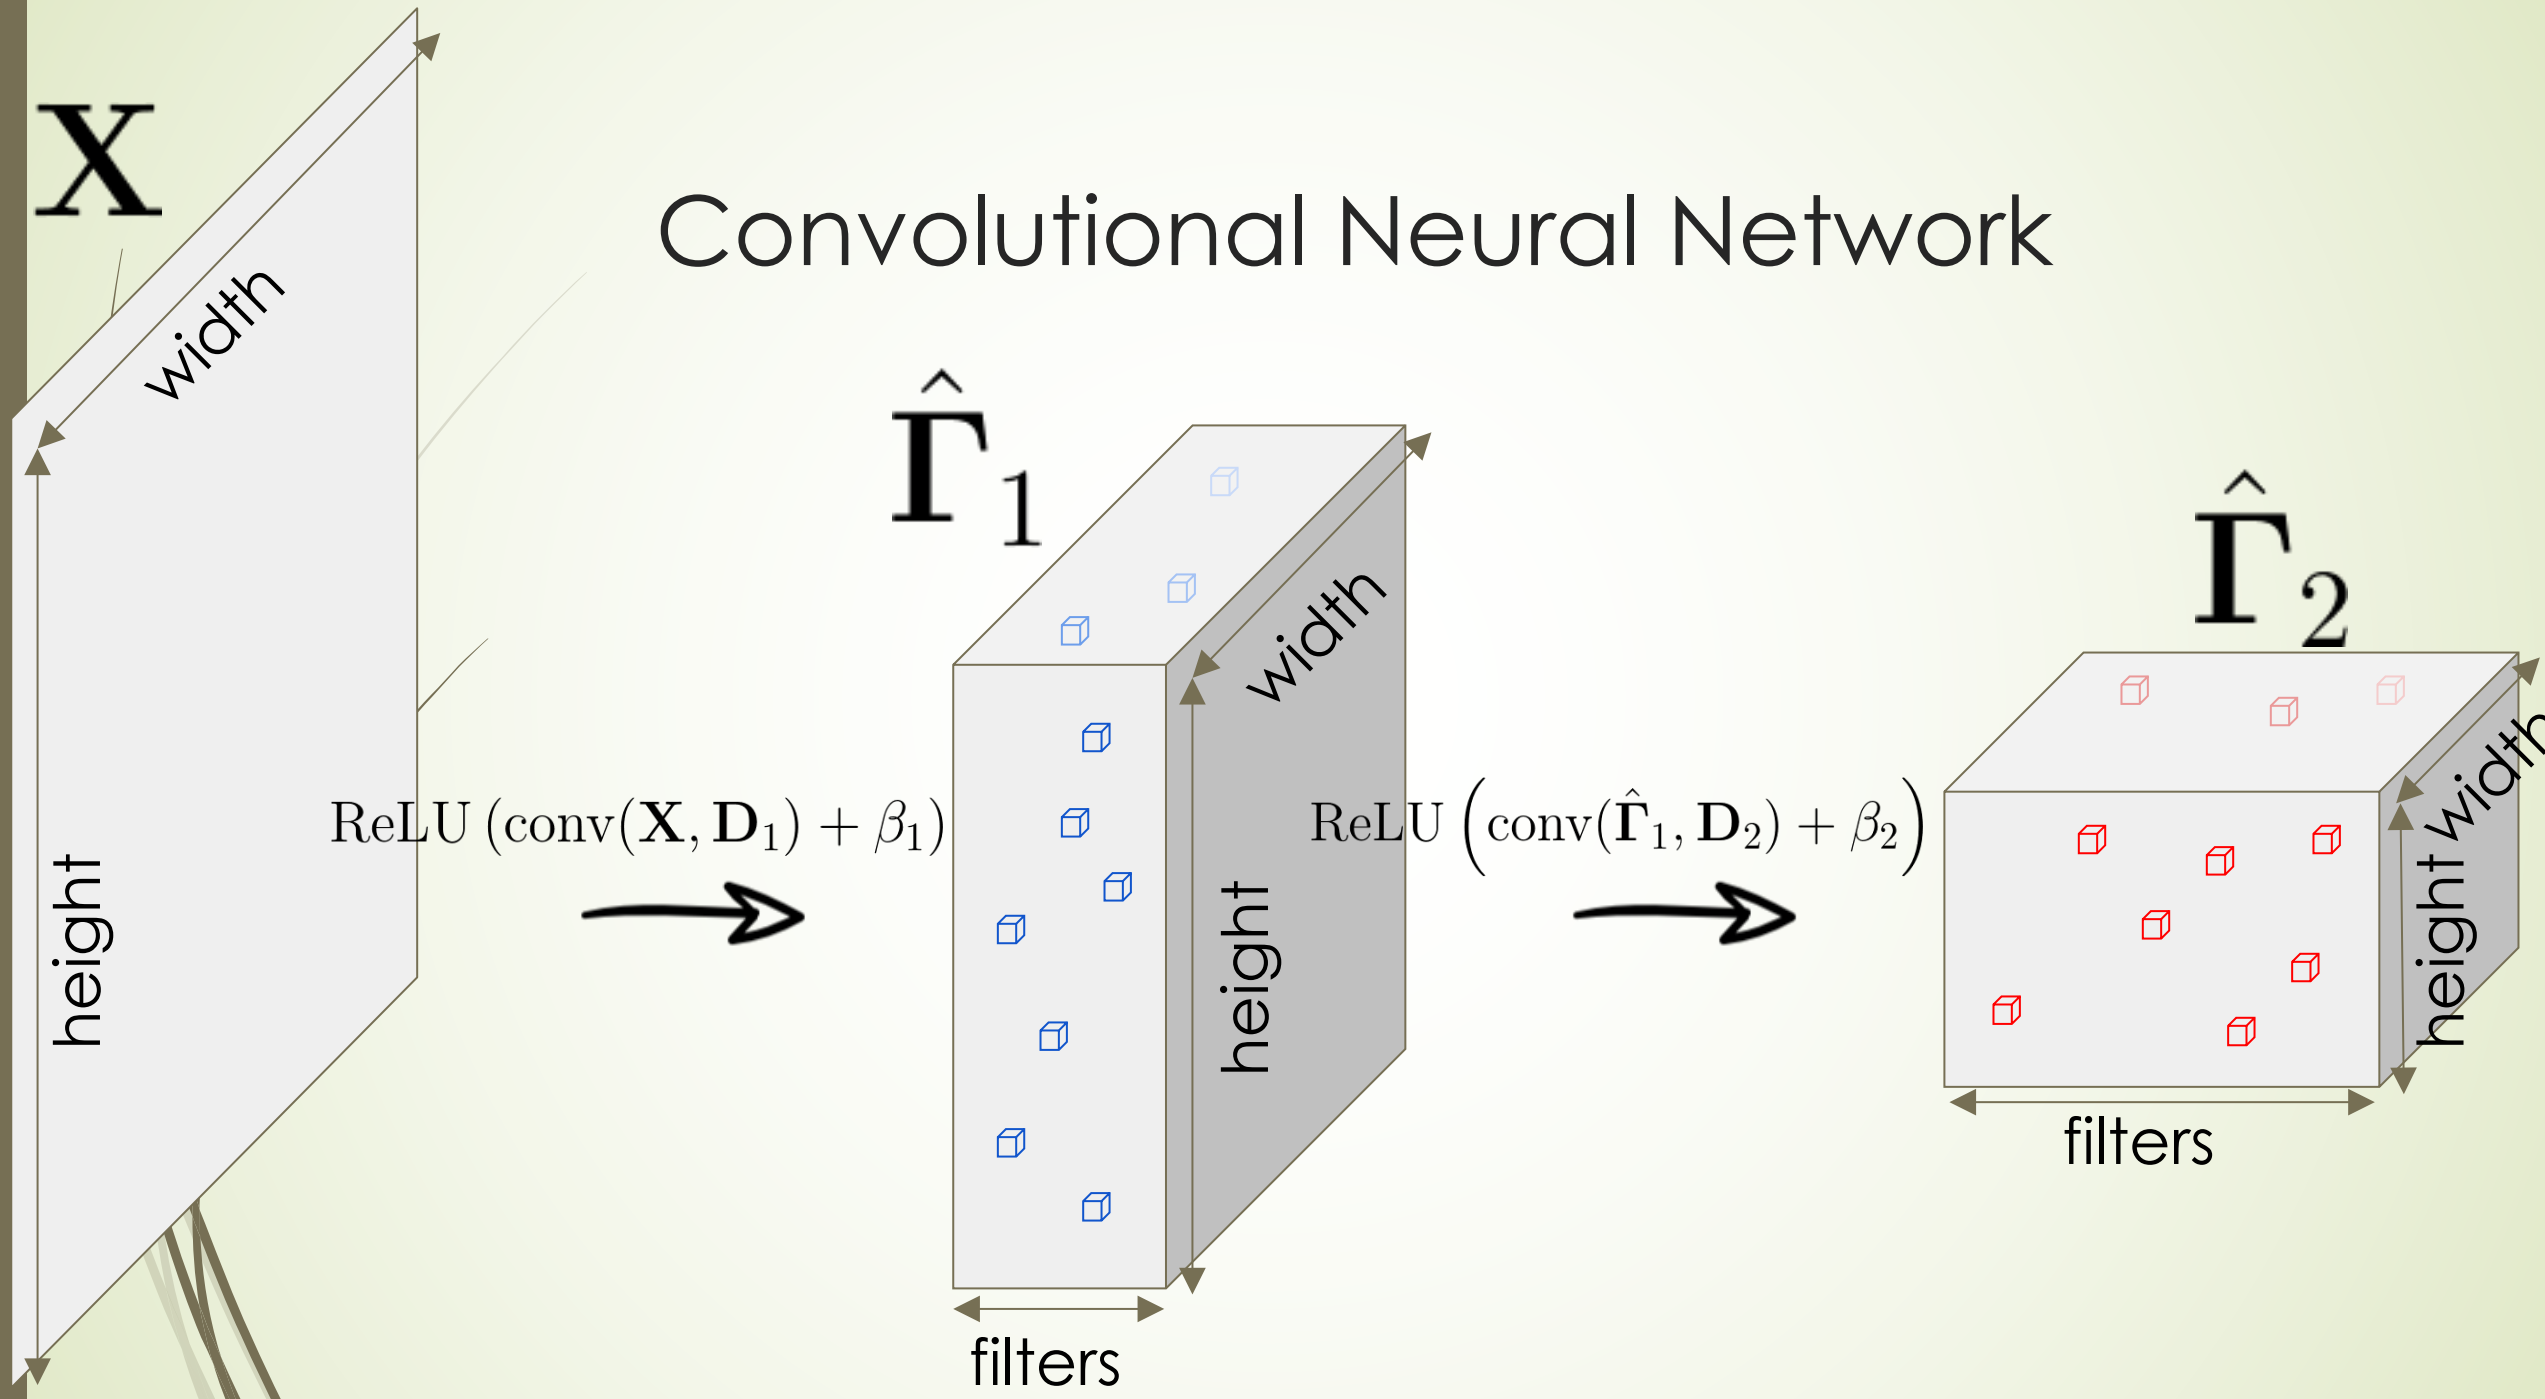

Multi-layered Convolutional Sparse Modeling

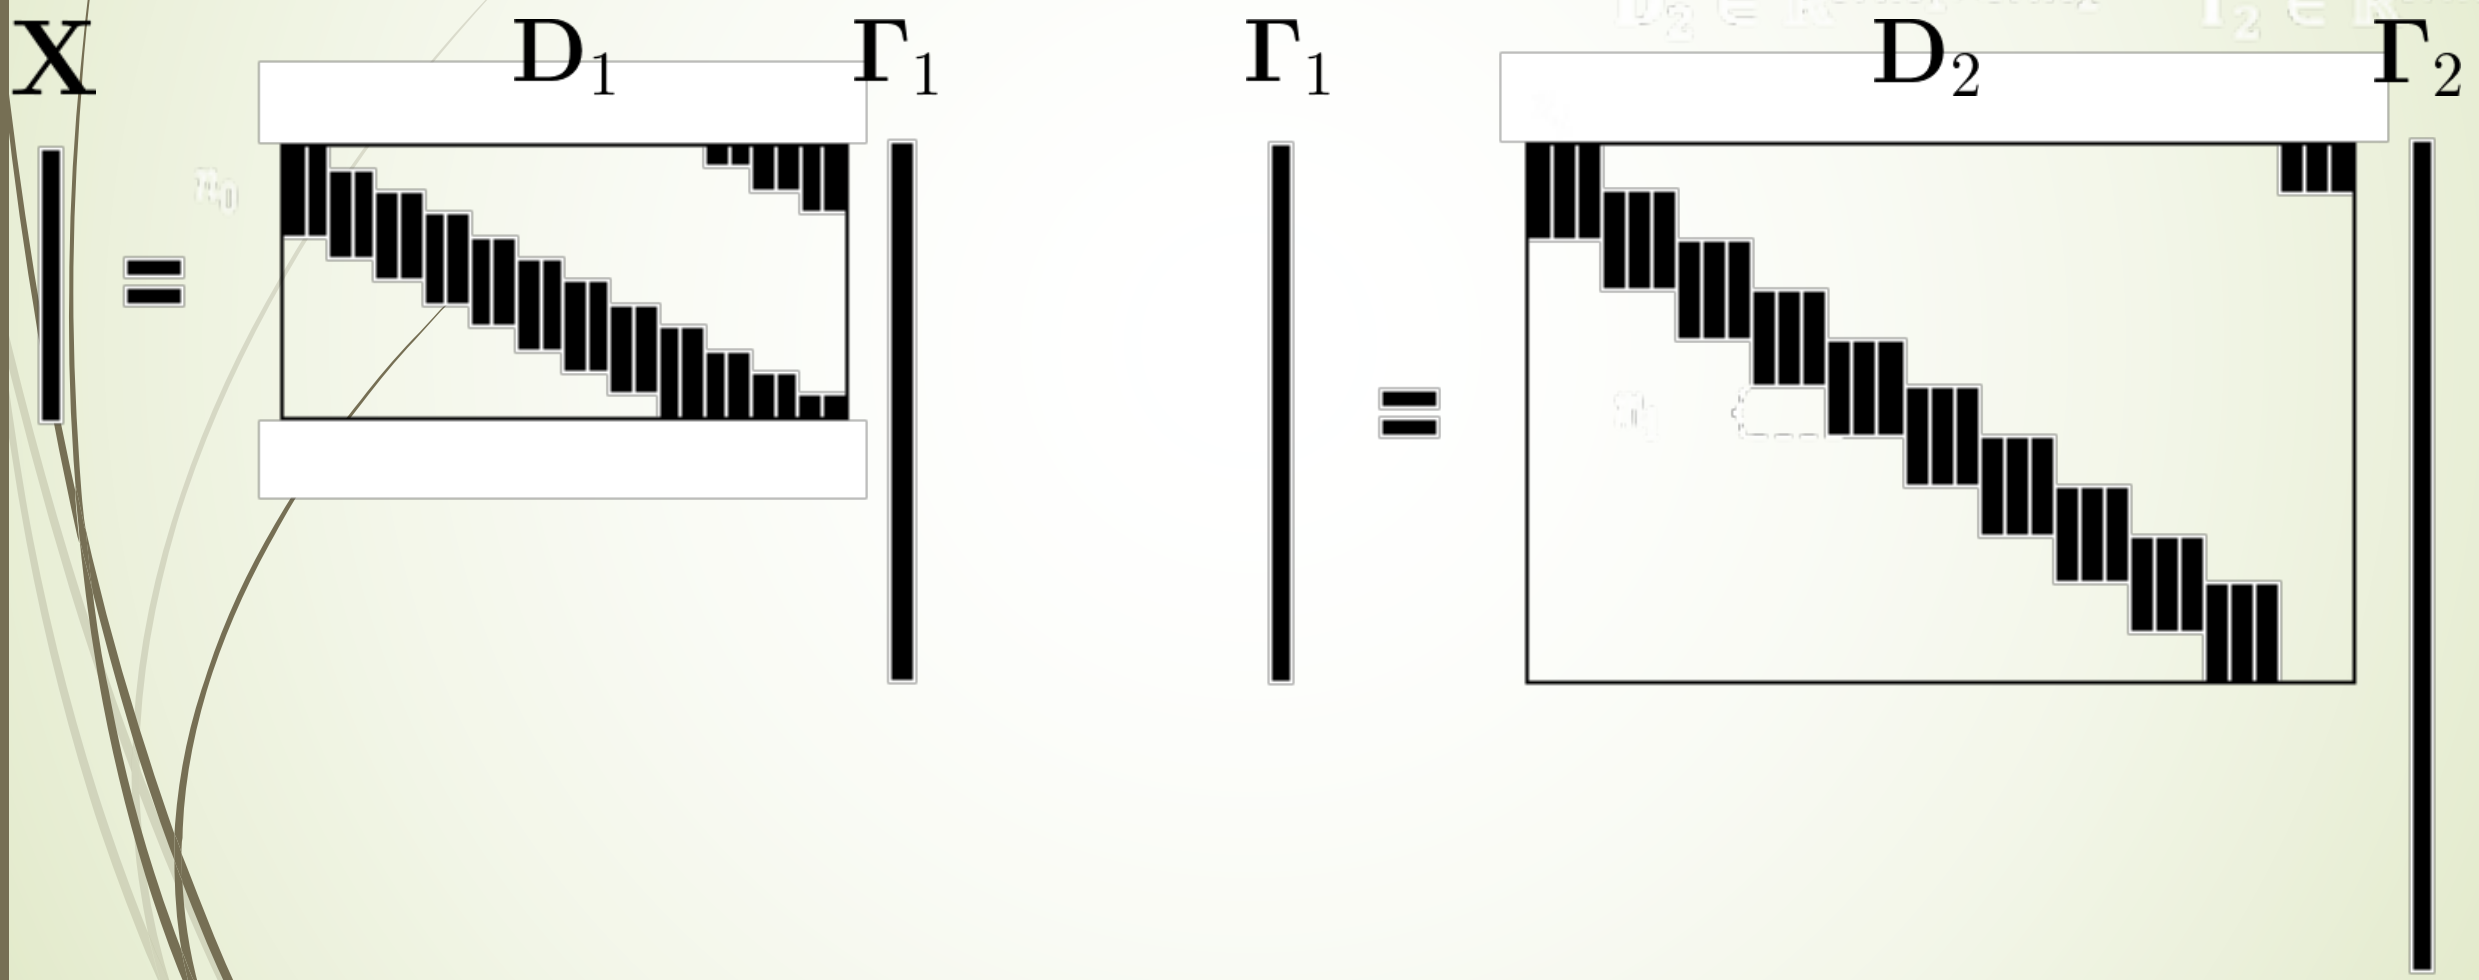

Multi-layered Convolutional Sparse Modeling

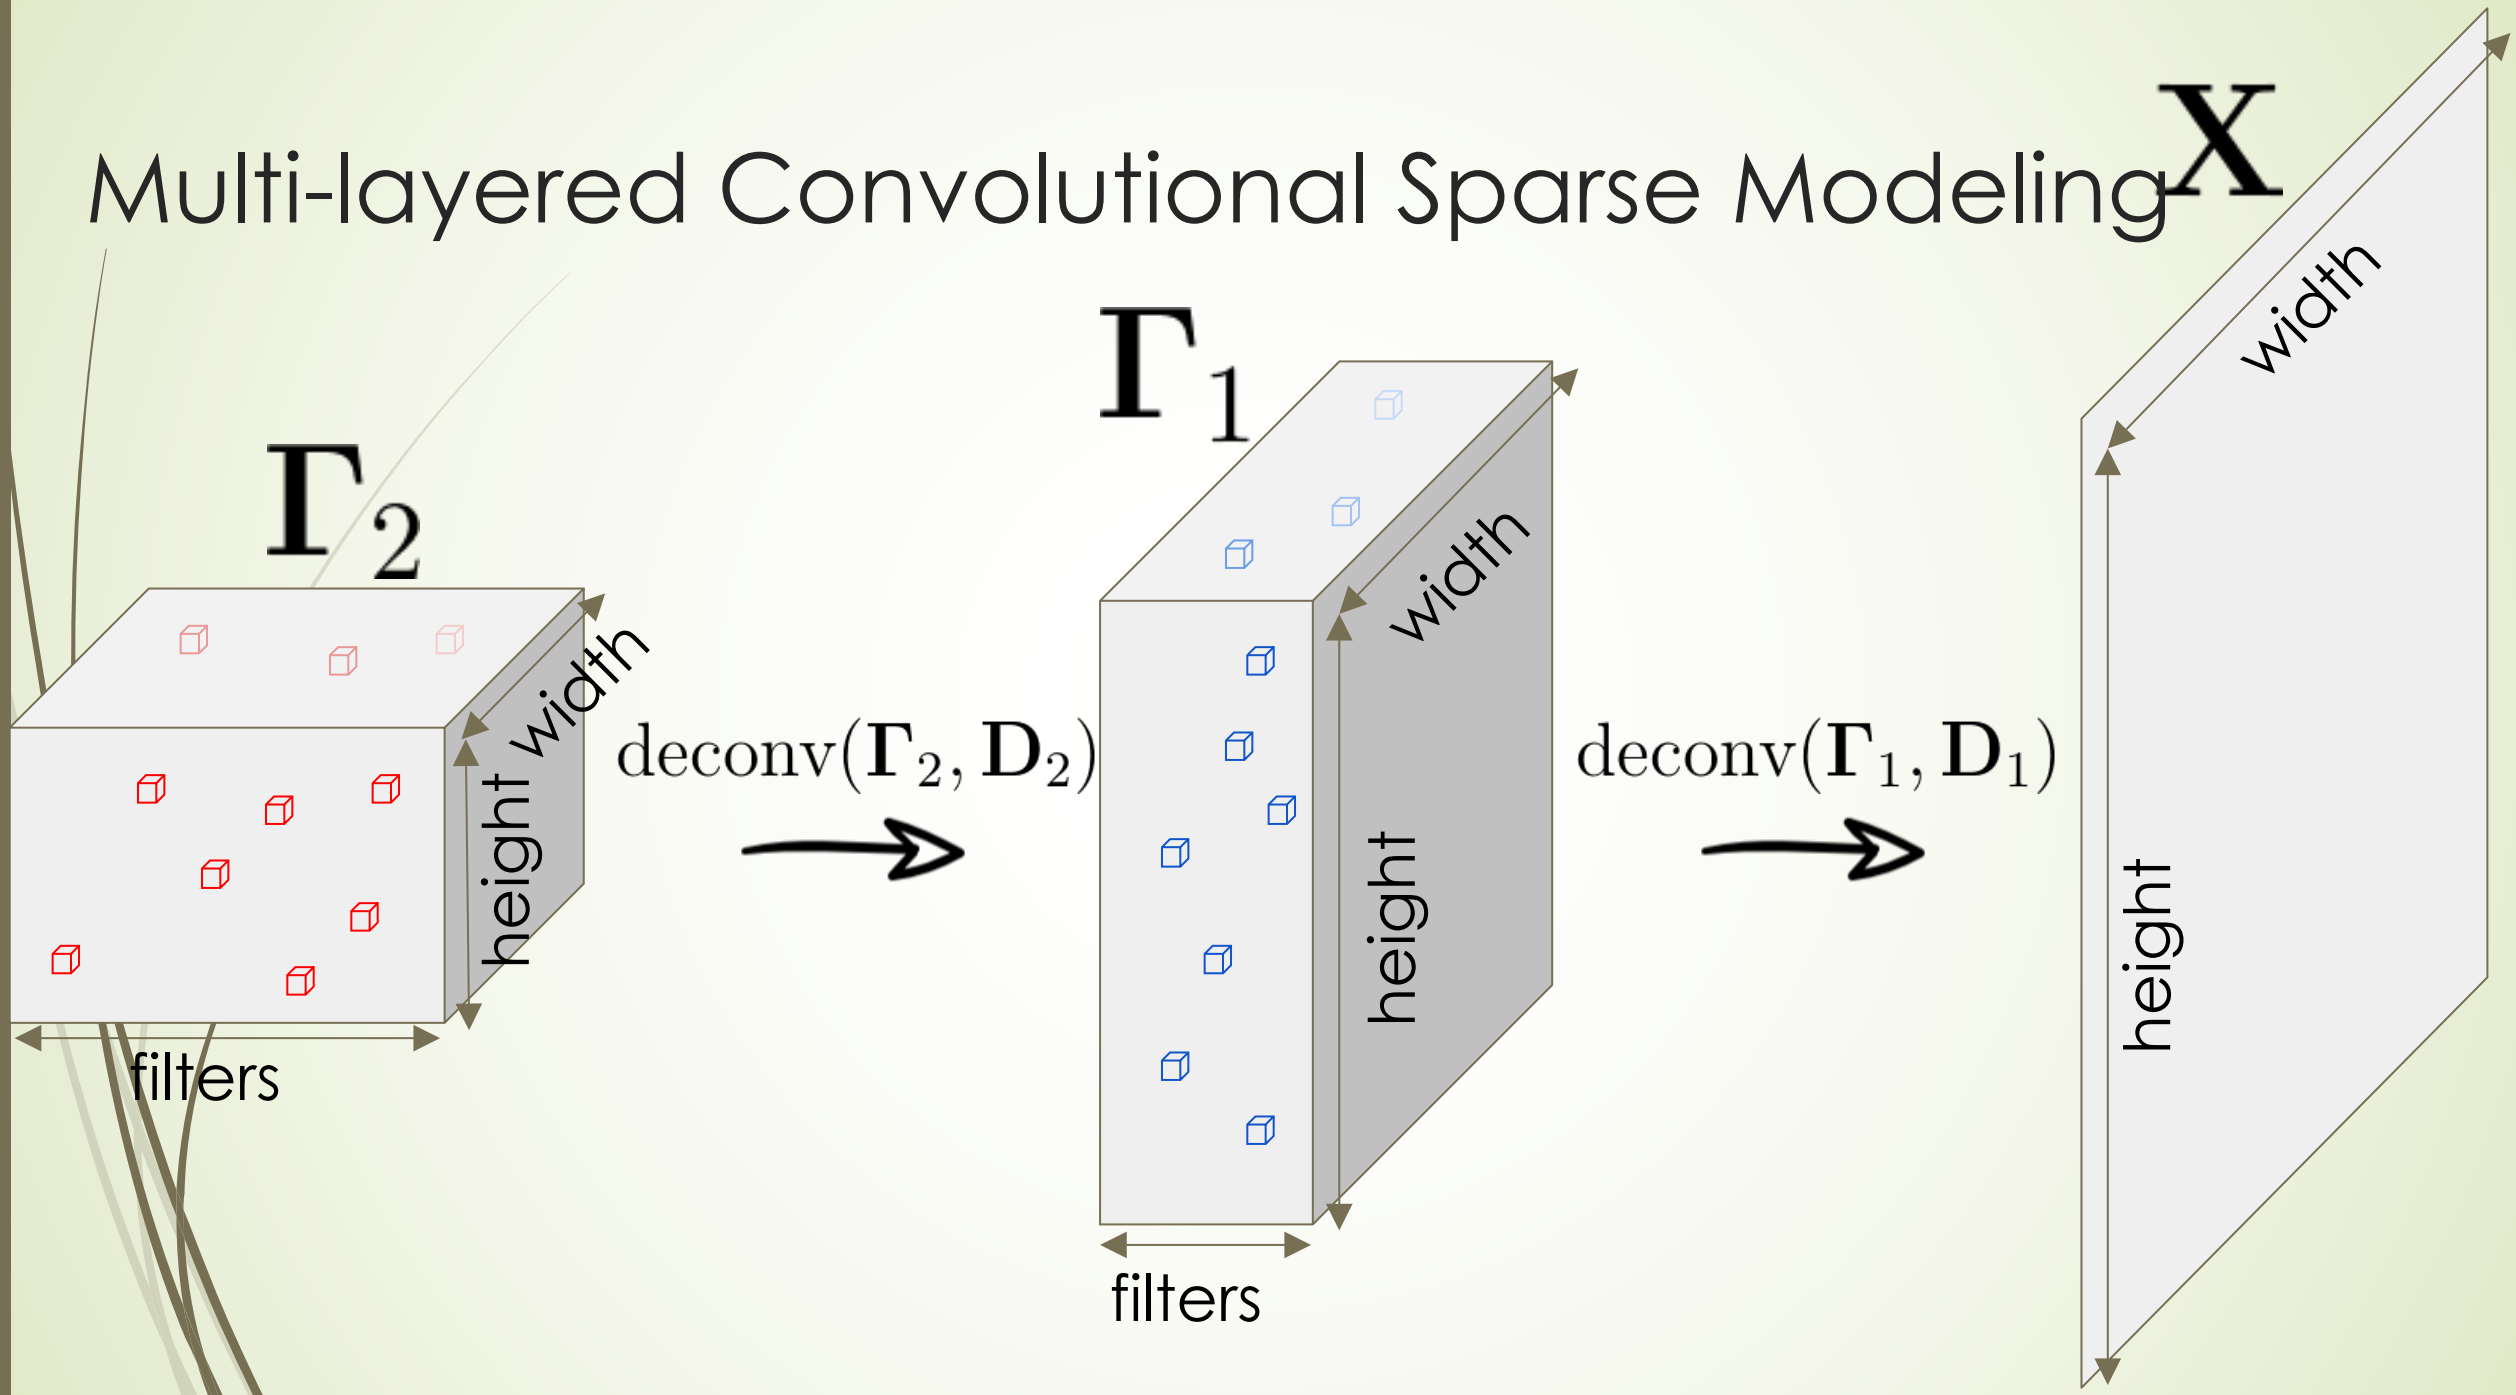

Layered Thresholding

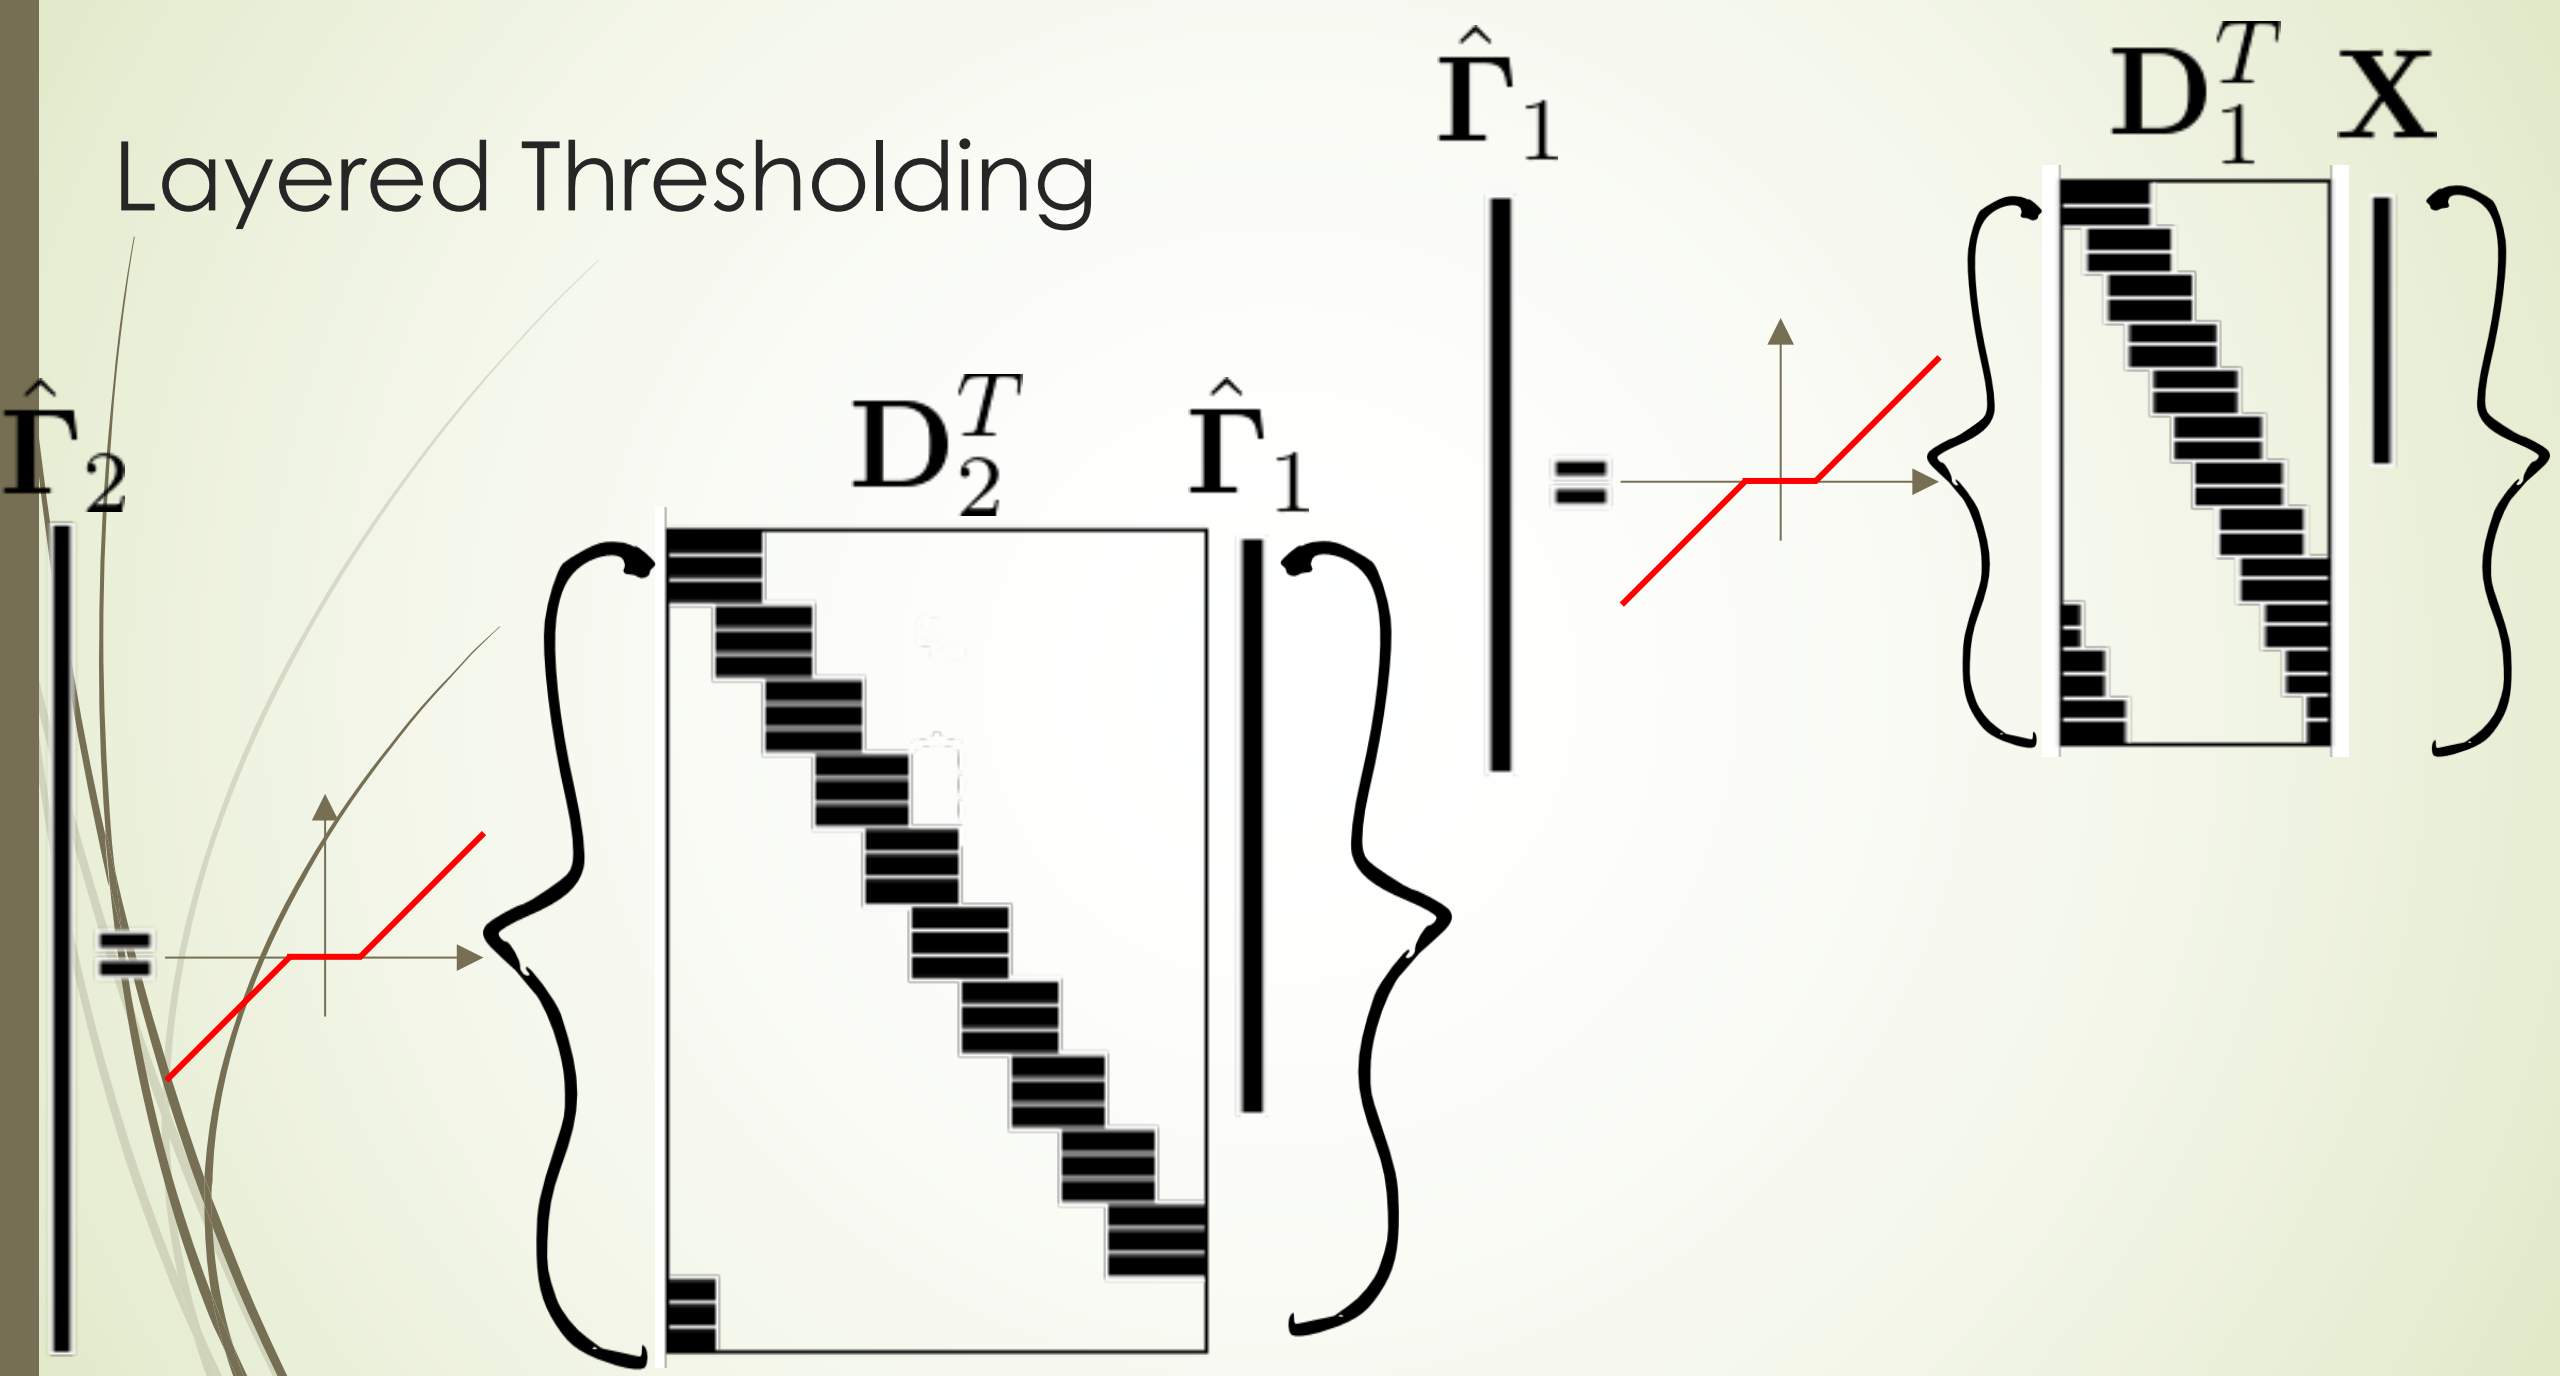

Convolutional Neural Network

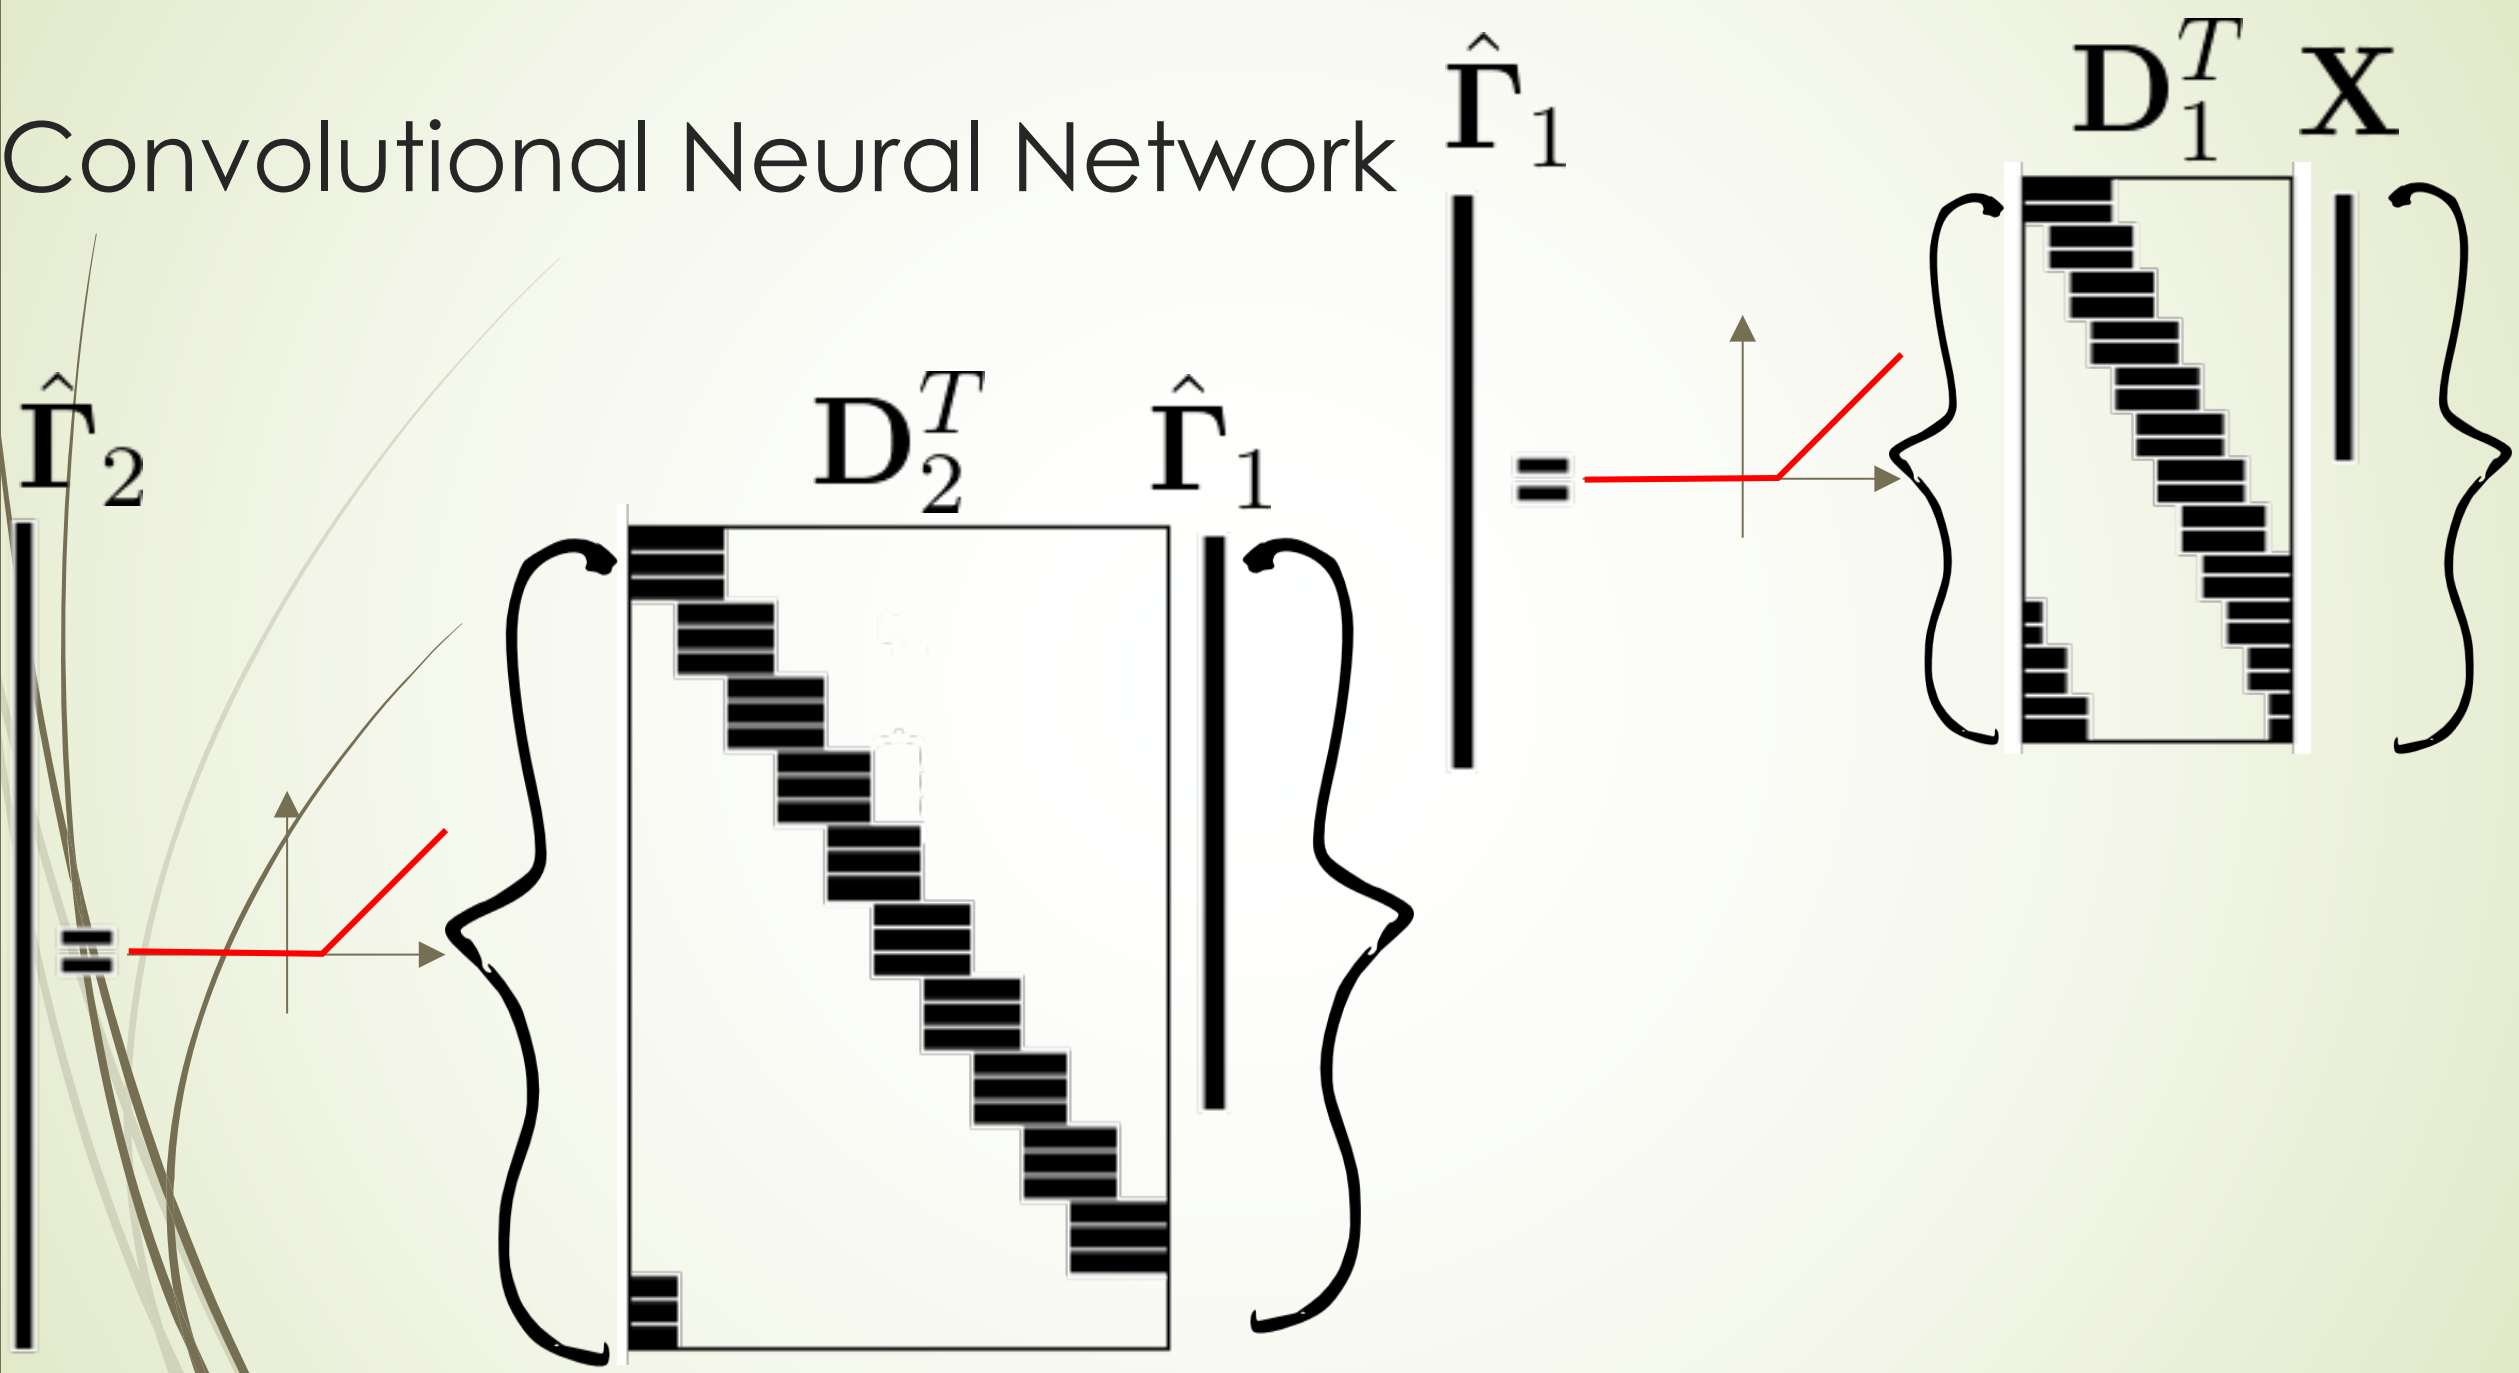

Convolutional Neural Network

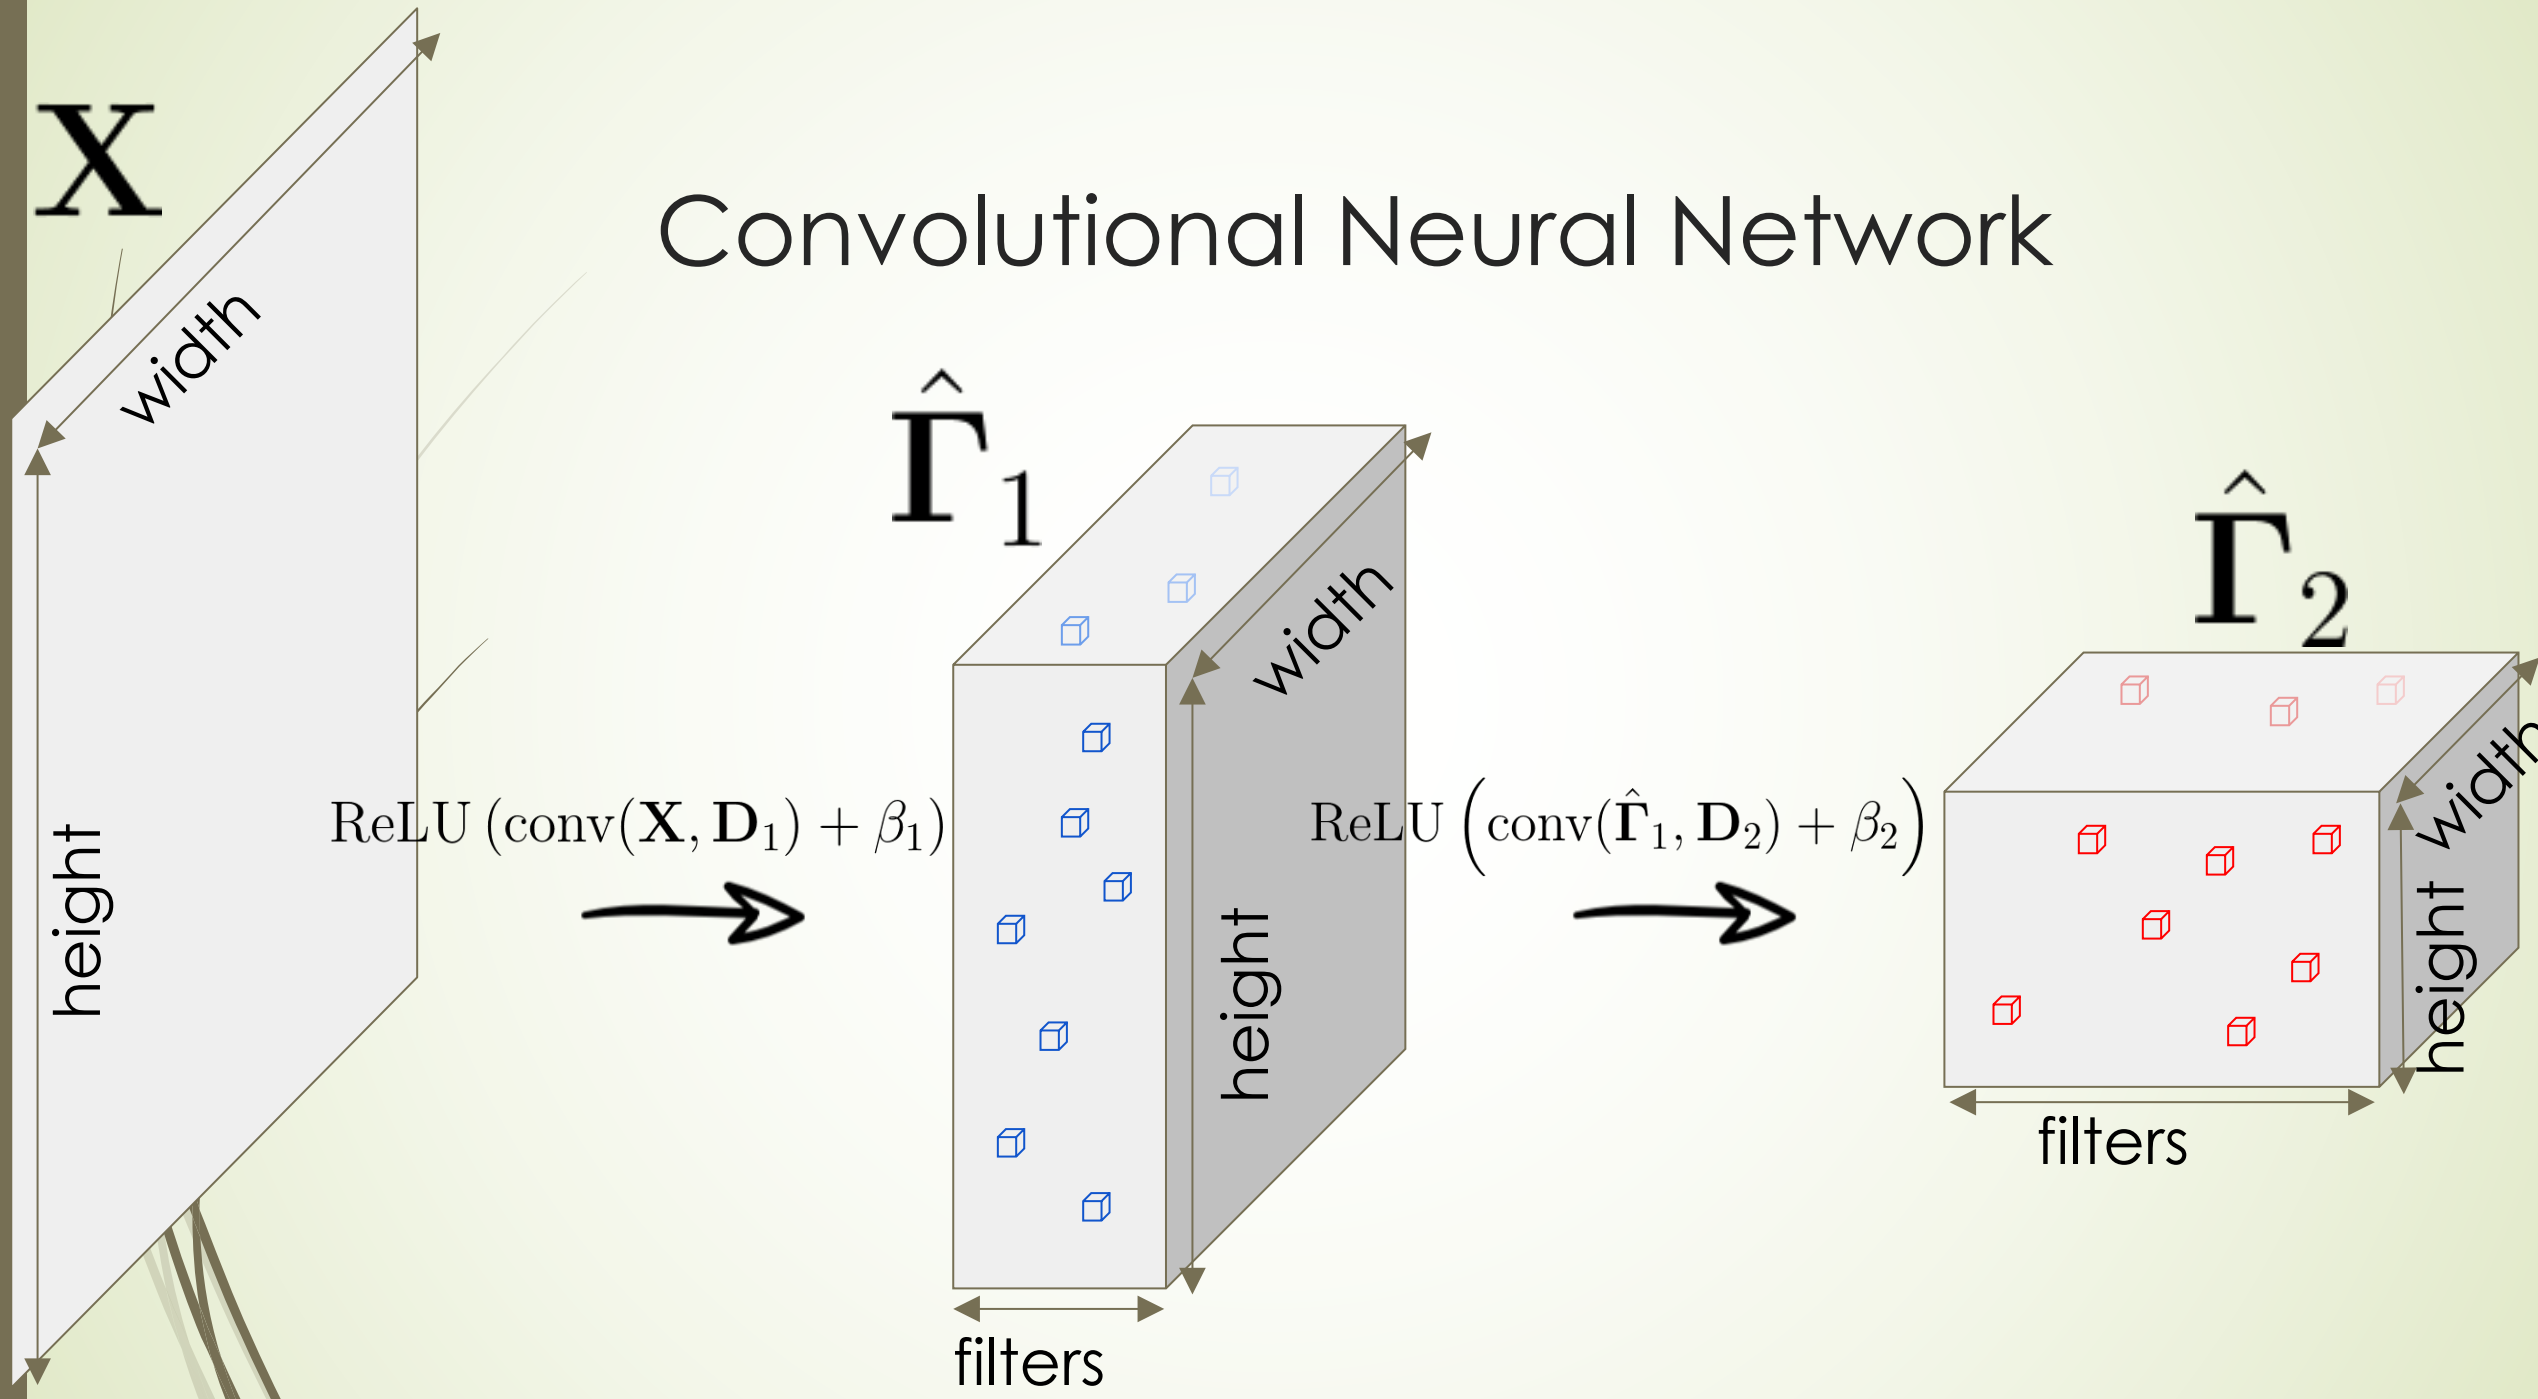

Theories of Sparse Coding

Sparse Modeling

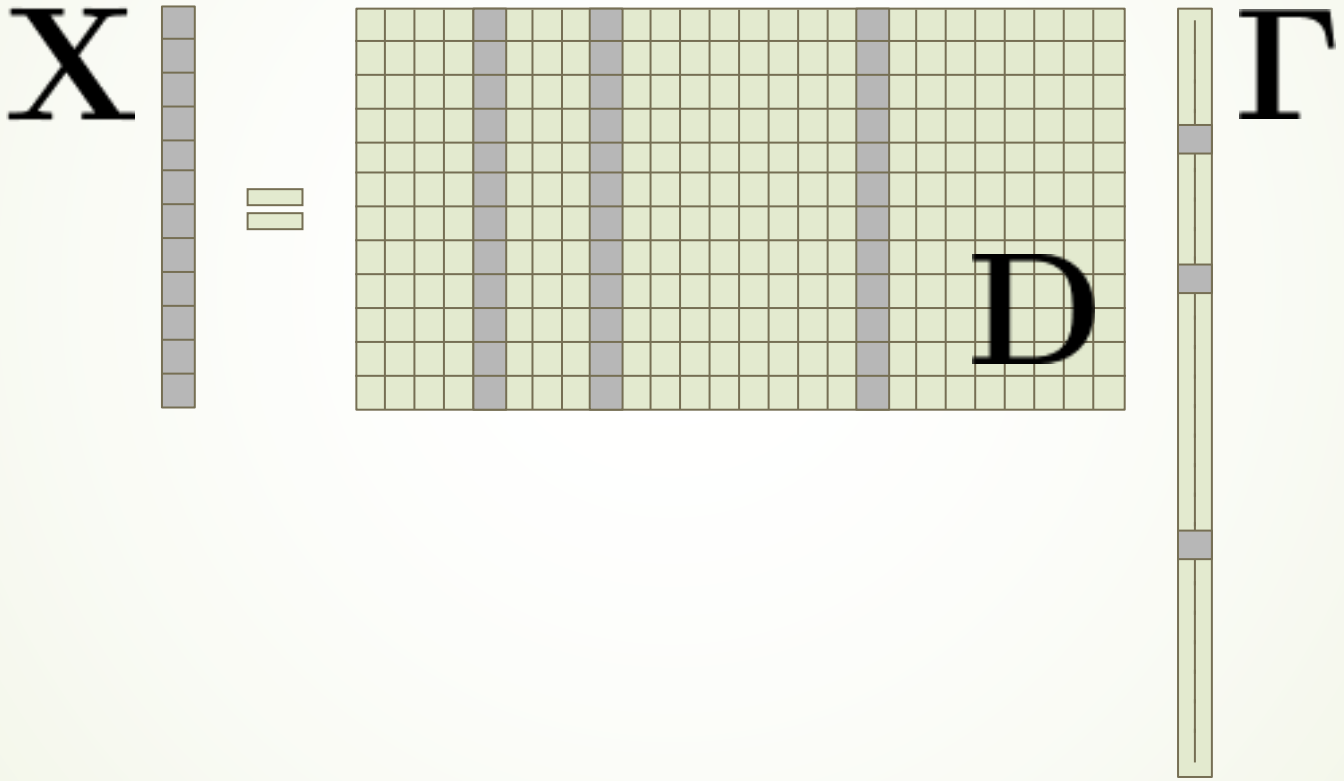

Classic Sparse Theory: Basis Pursuit

$$\mathbf{X} = \mathbf{D}\boldsymbol{\Gamma}$$

$$\hat{\boldsymbol{\Gamma}} = \arg\min_{\boldsymbol{\Gamma}} \|\boldsymbol{\Gamma}\|_1 \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\boldsymbol{\Gamma}$$

Theorem: [Donoho and Elad, 2003]
Basis pursuit is guaranteed to recover the true sparse vector assuming that

$$\|\boldsymbol{\Gamma}\|_0 < \frac{1}{2}\left(1 + \frac{1}{\mu(\mathbf{D})}\right)$$

Mutual Coherence: $\mu(\mathbf{D}) = \max_{i \neq j} \left| (\mathbf{D}^T \mathbf{D})_{i,j} \right|$

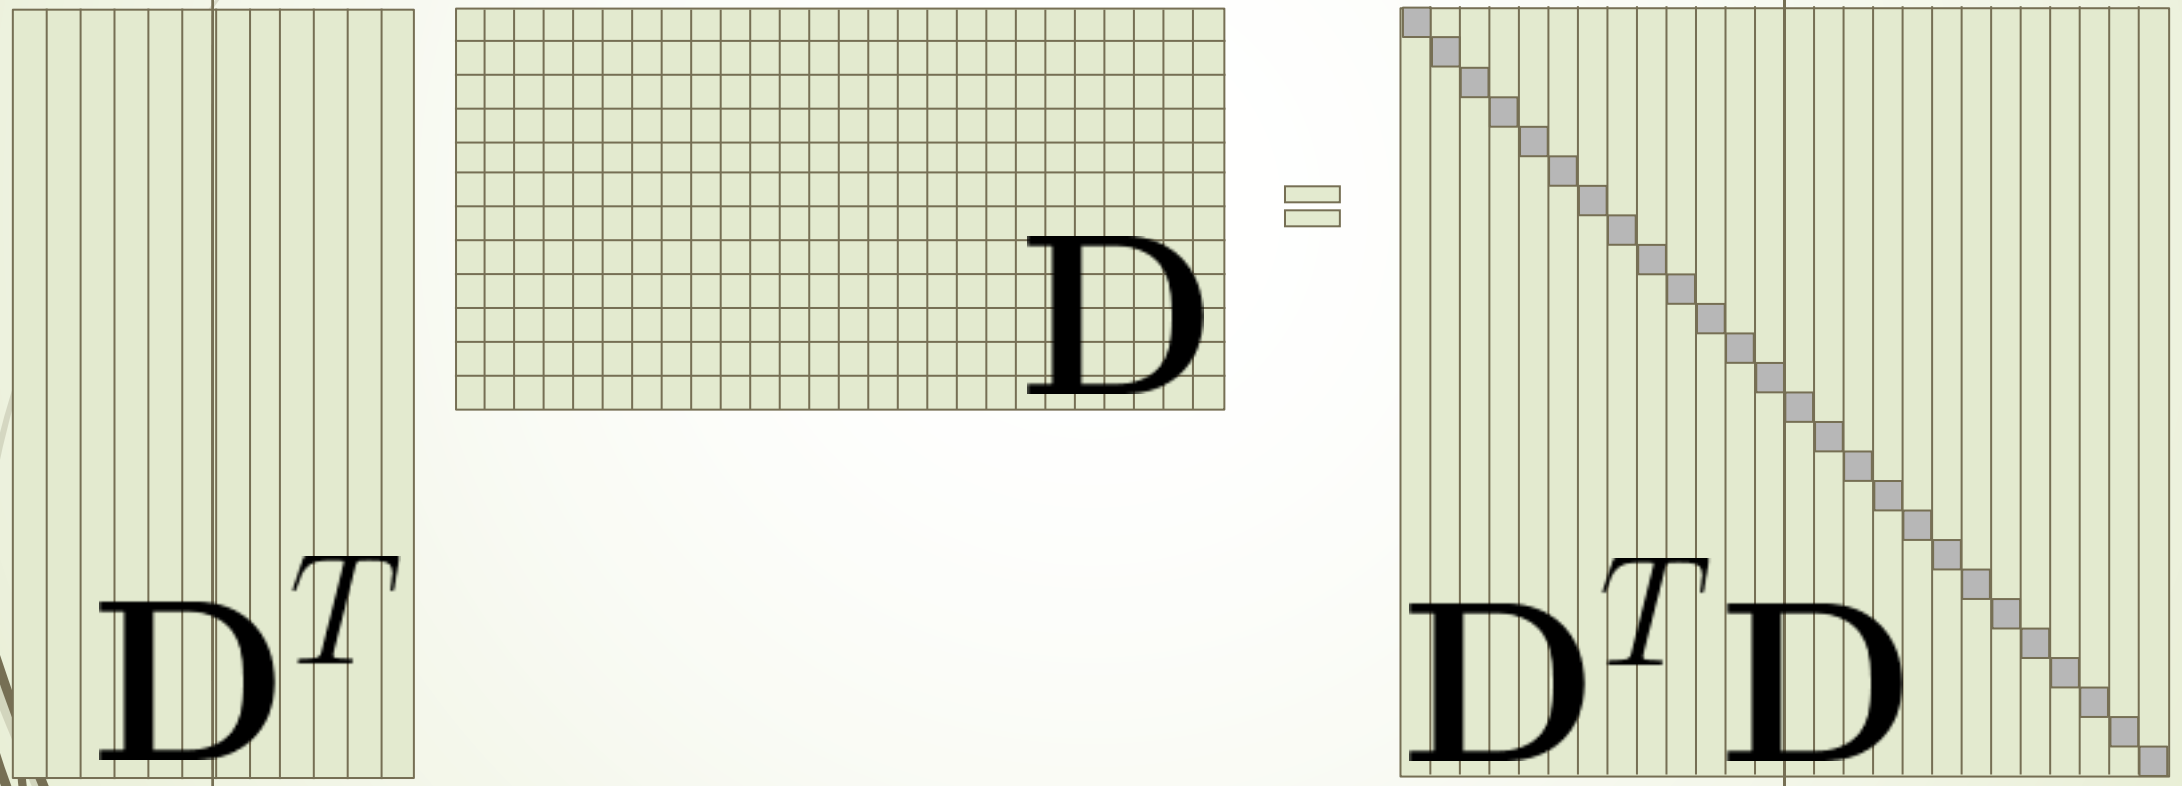

Convolutional Sparse Modeling

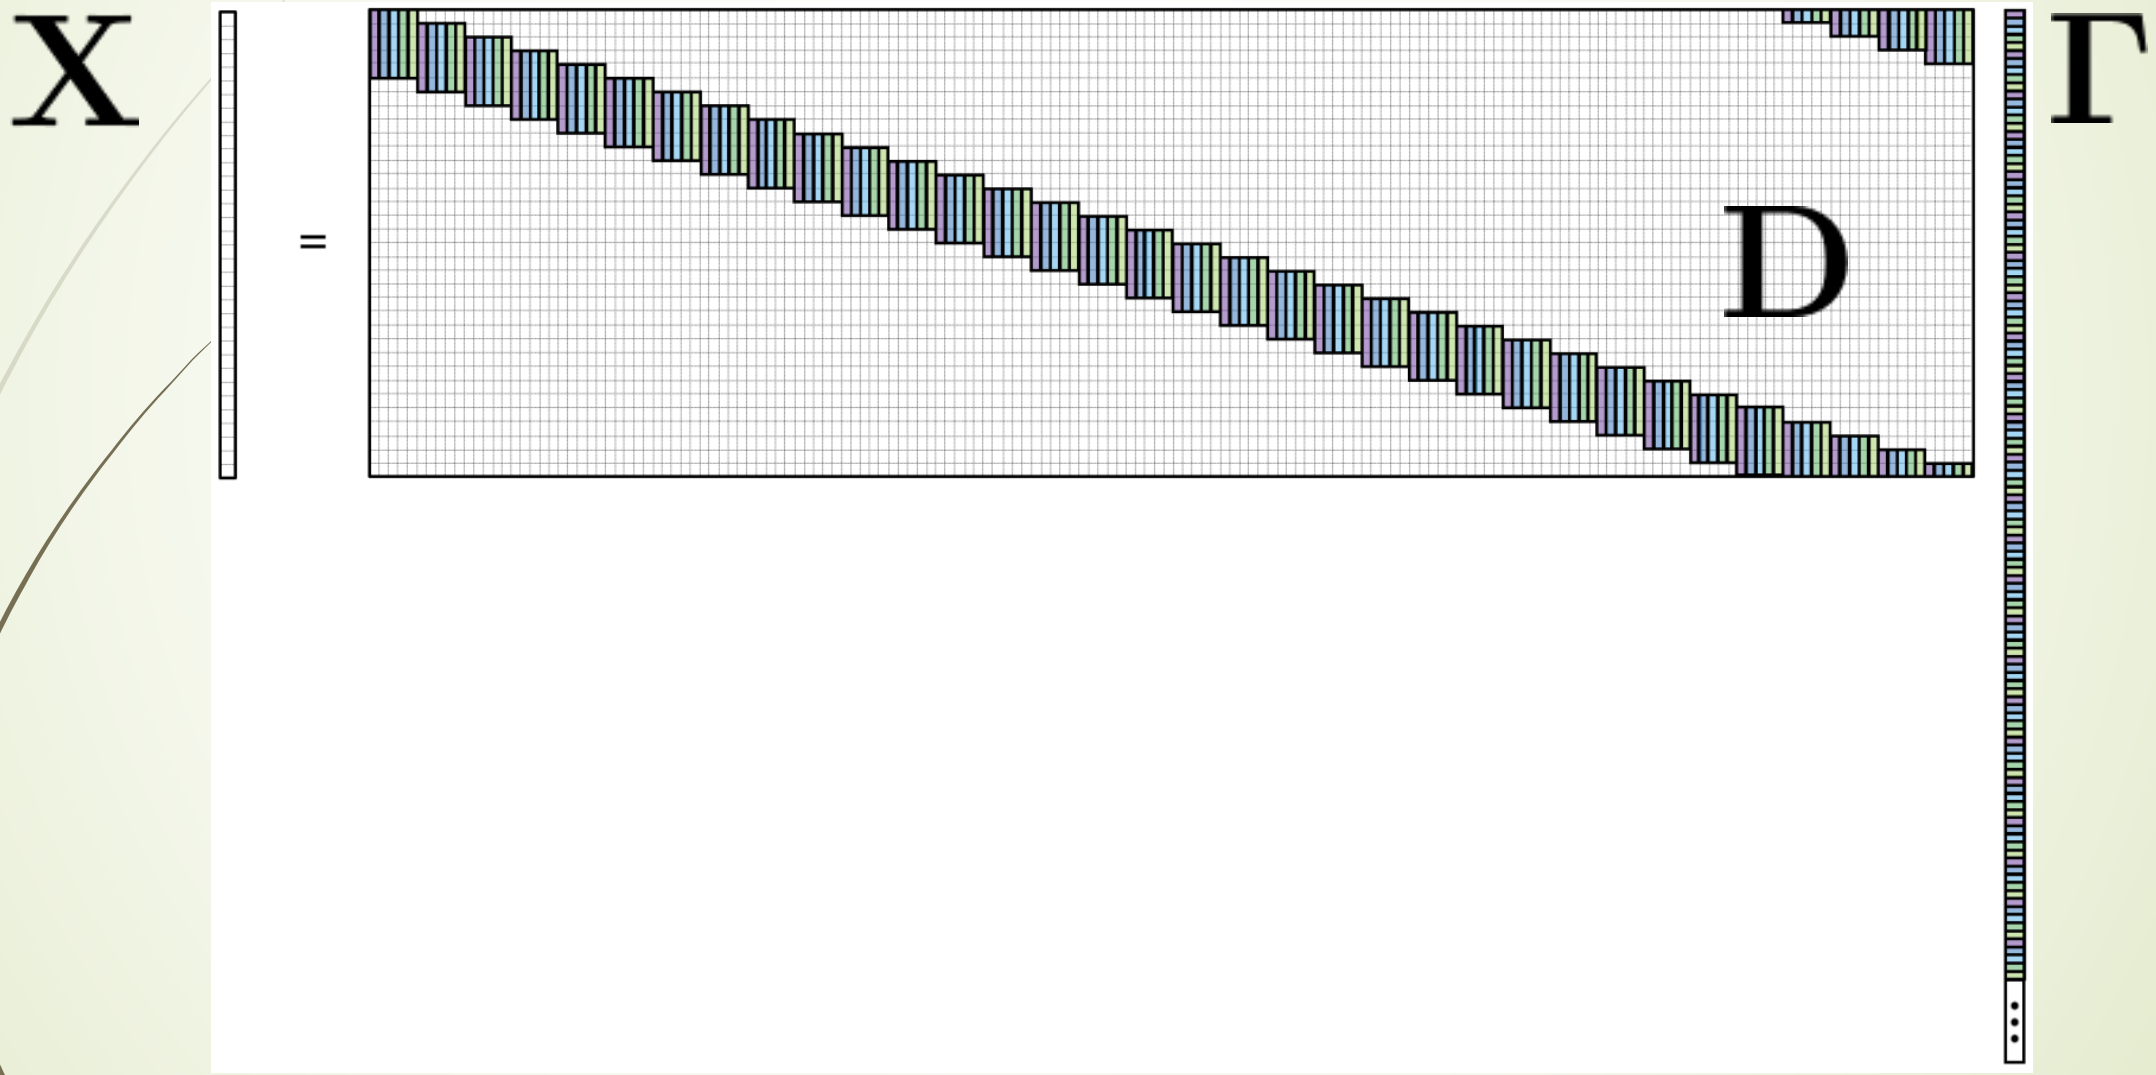

Classic Sparse Theory for Convolutional Case

Theorem: [Donoho and Elad, 2003]
Basis pursuit is guaranteed to recover the true sparse vector assuming that

$$\|\mathbf{\Gamma}\|_0 < \frac{1}{2}\left(1 + \frac{1}{\mu(\mathbf{D})}\right)$$

Assuming 2 atoms of length 64 $\mu(\mathbf{D}) \geq 0.063$ [Welch, 1974]

Success guaranteed when

$$\|\mathbf{\Gamma}\|_0 < 8.43$$

Local Sparsity

$$\|\mathbf{\Gamma}\|_{0,\infty}$$

maximal number of non-zeroes
in a local neighborhood

$$\min_{\mathbf{\Gamma}} \|\mathbf{\Gamma}\|_{0,\infty} \quad \text{s.t.} \quad \mathbf{X} = \mathbf{D}\mathbf{\Gamma}$$

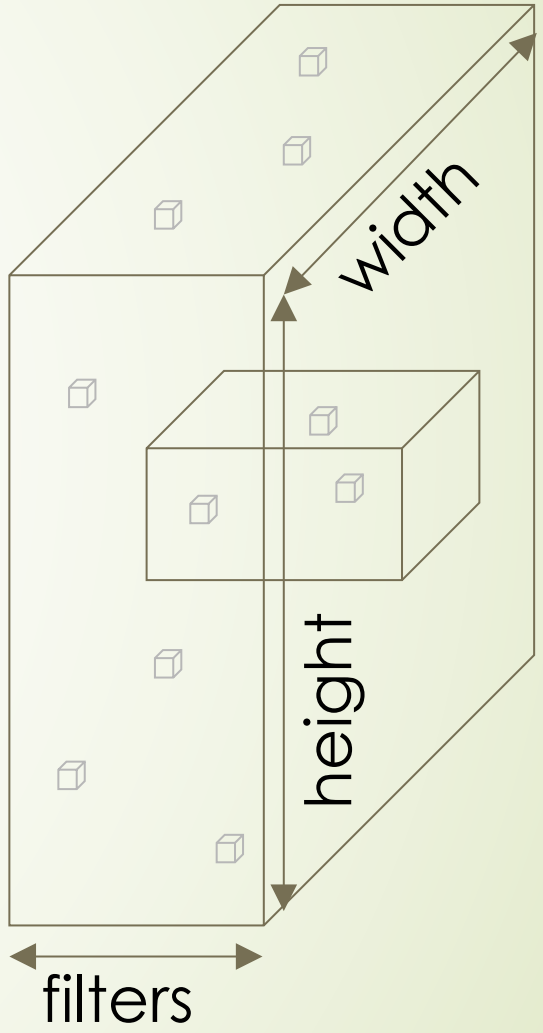

Success of Basis Pursuit DeNoising

$$\mathbf{Y} = \mathbf{D}\boldsymbol{\Gamma} + \mathbf{E}$$

$$\hat{\boldsymbol{\Gamma}} = \arg\min_{\boldsymbol{\Gamma}} \frac{1}{2}\|\mathbf{Y} - \mathbf{D}\boldsymbol{\Gamma}\|_2^2 + \lambda\|\boldsymbol{\Gamma}\|_1$$

Theorem: [Papyan, Sulam and Elad, 2016]

Assume: $\|\boldsymbol{\Gamma}\|_{0,\infty} < \frac{1}{3}\left(1 + \frac{1}{\mu(\mathbf{D})}\right)$

Then: $\|\hat{\boldsymbol{\Gamma}} - \boldsymbol{\Gamma}\|_\infty \leq 7.5\|\mathbf{E}\|_{2,\infty}$

Theoretical guarantee for:

- [Zeiler et. al 2010]
- [Wohlberg 2013]
- [Bristow et. al 2013]
- [Fowlkes and Kong 2014]
- [Zhou et. al 2014]
- [Kong and Fowlkes 2014]
- [Zhu and Lucey 2015]
- [Heide et. al 2015]
- [Gu et. al 2015]
- [Wohlberg 2016]
- [Šorel and Šroubek 2016]
- [Serrano et. al 2016]
- [Papyan et. al 2017]
- [Garcia-Cardona and Wohlberg 2017]
- [Wohlberg and Rodriguez 2017]
- ...

Multi-layered Convolutional Sparse Modeling

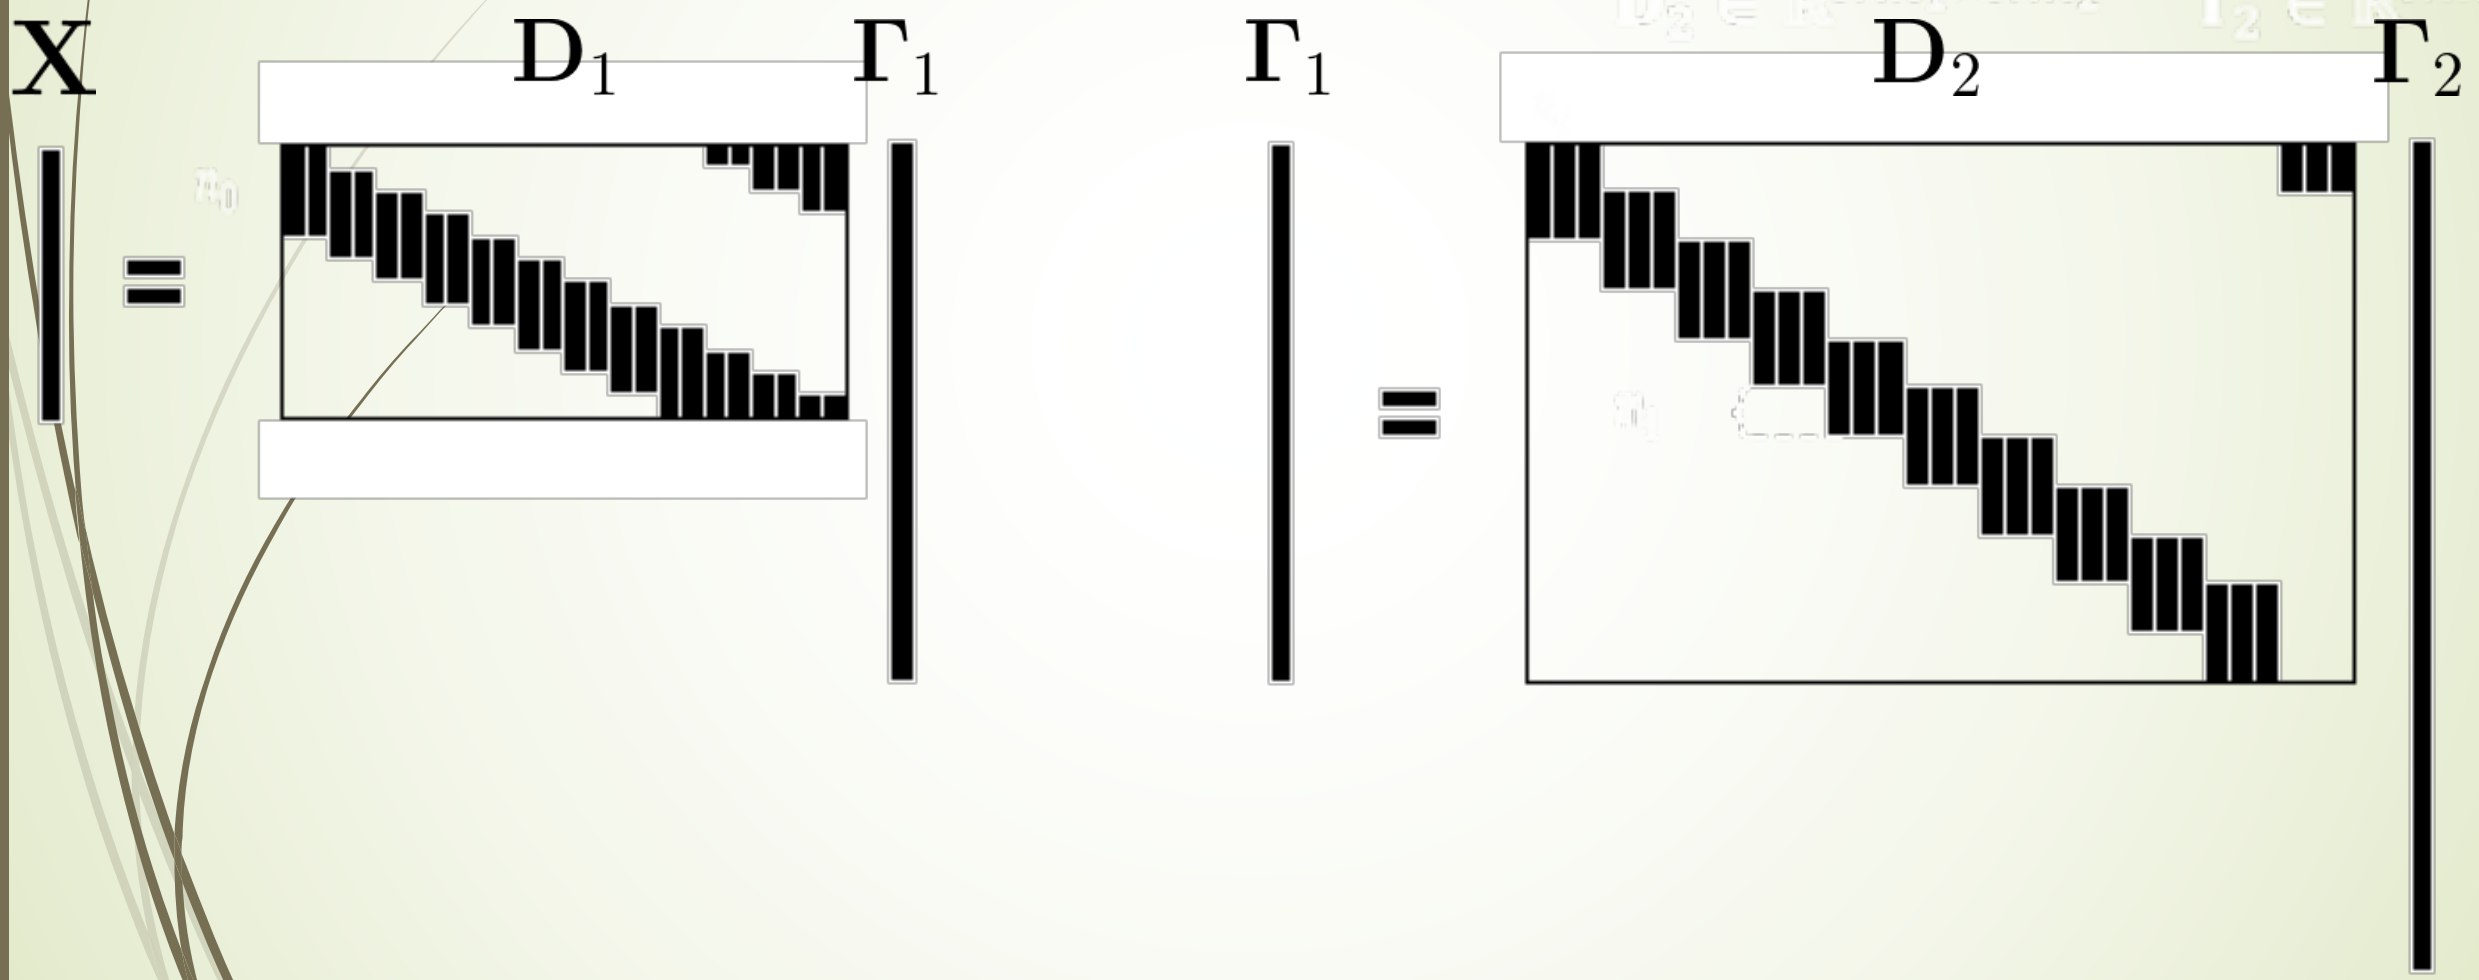

Deep Coding Problem

Given $\mathbf{X}$, find a set of representations satisfying:

$$\mathbf{X} = \mathbf{D}_1\boldsymbol{\Gamma}_1, \qquad \|\boldsymbol{\Gamma}_1\|_{0,\infty} \leq \lambda_1$$

$$\boldsymbol{\Gamma}_1 = \mathbf{D}_2\boldsymbol{\Gamma}_2, \qquad \|\boldsymbol{\Gamma}_2\|_{0,\infty} \leq \lambda_2$$

$$\vdots$$

$$\boldsymbol{\Gamma}_{L-1} = \mathbf{D}_L\boldsymbol{\Gamma}_L, \qquad \|\boldsymbol{\Gamma}_L\|_{0,\infty} \leq \lambda_L$$

Deep Coding Problem

Given $\mathbf{Y}$, find a set of representations satisfying:

$$\|\mathbf{Y} - \mathbf{D}_1\boldsymbol{\Gamma}_1\|_2 \leq \epsilon, \qquad \|\boldsymbol{\Gamma}_1\|_{0,\infty} \leq \lambda_1$$

$$\boldsymbol{\Gamma}_1 = \mathbf{D}_2\boldsymbol{\Gamma}_2, \qquad \|\boldsymbol{\Gamma}_2\|_{0,\infty} \leq \lambda_2$$

$$\vdots$$

$$\boldsymbol{\Gamma}_{L-1} = \mathbf{D}_L\boldsymbol{\Gamma}_L, \qquad \|\boldsymbol{\Gamma}_L\|_{0,\infty} \leq \lambda_L$$

Uniqueness

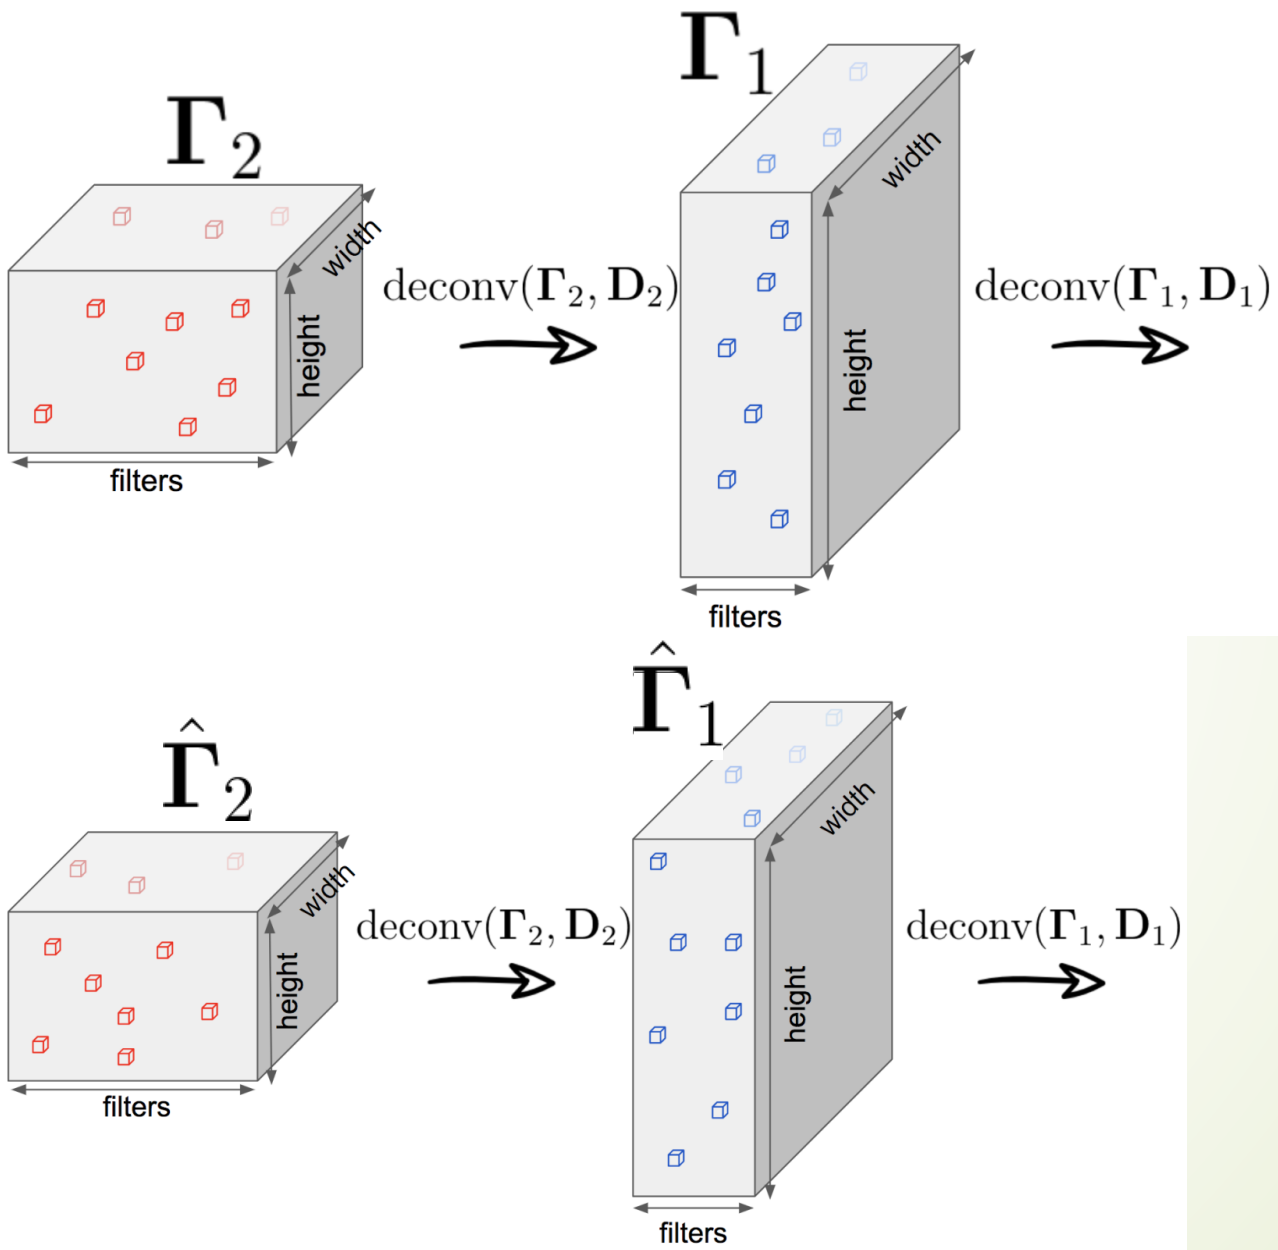
$\boldsymbol{\Gamma}_2$
width
height
filters
$\text{deconv}(\boldsymbol{\Gamma}_2, \mathbf{D}_2)$
$\boldsymbol{\Gamma}_1$
width
height
filters
$\text{deconv}(\boldsymbol{\Gamma}_1, \mathbf{D}_1)$
$\mathbf{X}$
width
height
$\hat{\boldsymbol{\Gamma}}_2$
width
height
filters
$\text{deconv}(\boldsymbol{\Gamma}_2, \mathbf{D}_2)$
$\hat{\boldsymbol{\Gamma}}_1$
width
height
filters
$\text{deconv}(\boldsymbol{\Gamma}_1, \mathbf{D}_1)$

Uniqueness Theorem

$$\|\mathbf{\Gamma}_l\|_{0,\infty} \leq \lambda_l < \frac{1}{2}\left(1 + \frac{1}{\mu(\mathbf{D}_l)}\right)$$

$\{\mathbf{\Gamma}_l\}_{l=1}^{L}$ are the unique feature maps of $\mathbf{X}$

Success of Forward Pass

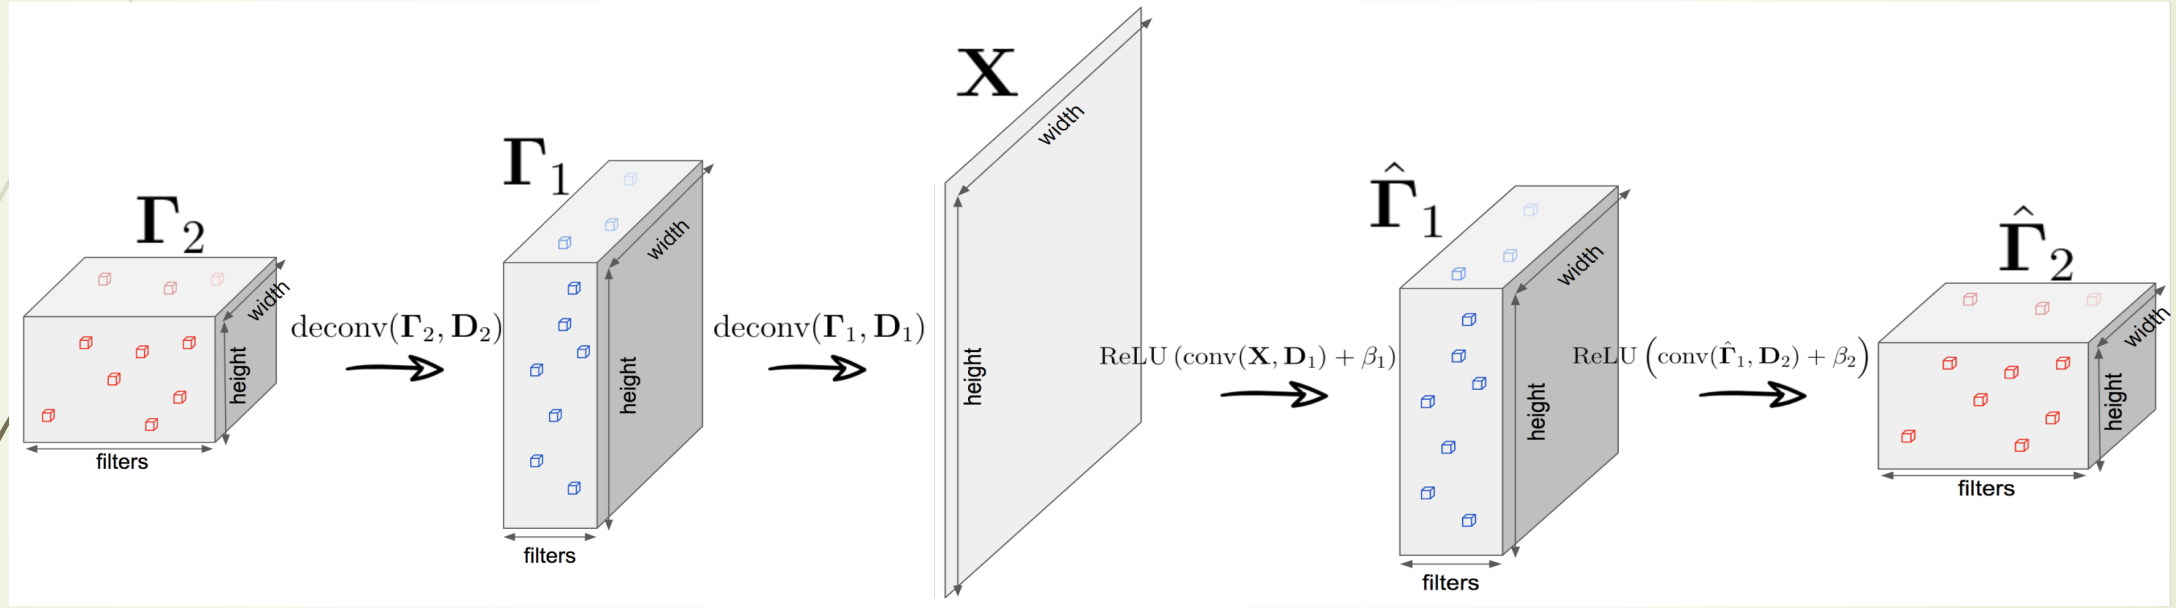

Success of Forward Pass Theorem

$$\|\mathbf{\Gamma}_l\|_{0,\infty} < \frac{1}{2}\left(1 + \frac{1}{\mu(\mathbf{D}_l)}\frac{|\Gamma_l^{\min}|}{|\Gamma_l^{\max}|}\right) - \frac{1}{\mu(\mathbf{D}_l)}\frac{\epsilon_{l-1}}{|\Gamma_l^{\max}|}$$

Layered thresholding guaranteed:

1. Find correct places of nonzeros

$$\|\hat{\mathbf{\Gamma}}_l - \mathbf{\Gamma}_l\|_{2,\infty} \leq \epsilon_l$$

Forward pass always fails at recovering representations exactly

Success depends on ratio

Distance increases with layer

Generative Model and Crude Inference

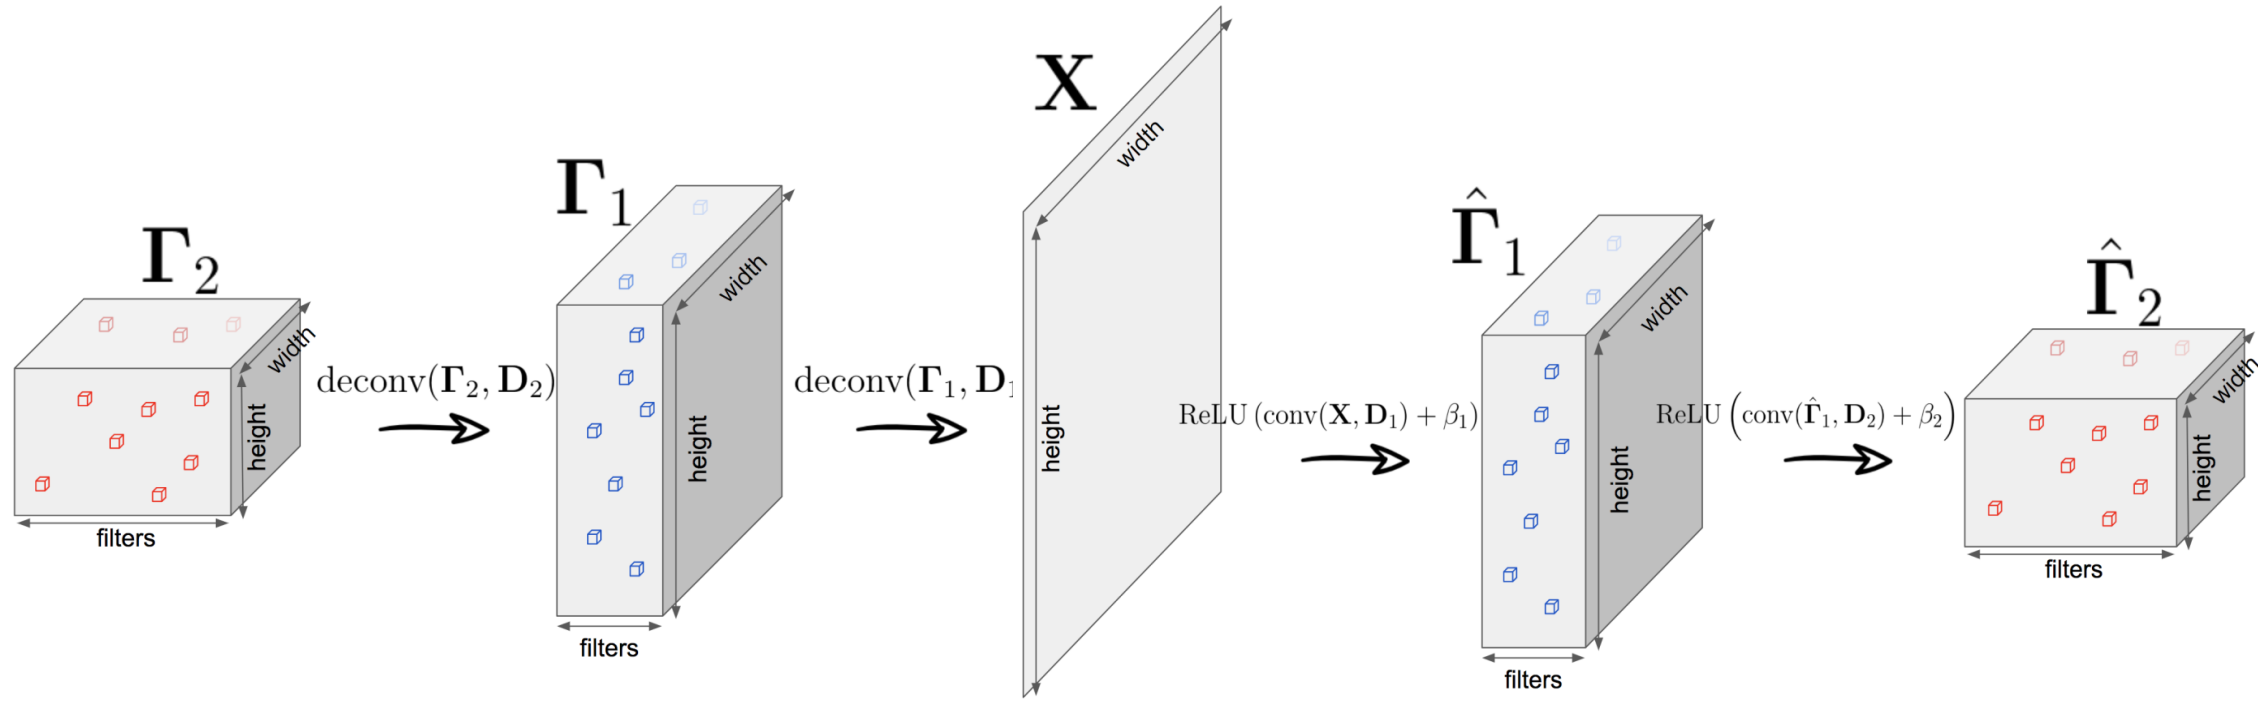

Layered Lasso

#StatsDepartment

$$\hat{\mathbf{\Gamma}}_1 = \arg\min_{\mathbf{\Gamma}_1} \frac{1}{2}\|\mathbf{Y} - \mathbf{D}_1\mathbf{\Gamma}_1\|_2^2 + \alpha_1\|\mathbf{\Gamma}_1\|_1$$

$$\hat{\mathbf{\Gamma}}_2 = \arg\min_{\mathbf{\Gamma}_2} \frac{1}{2}\|\hat{\mathbf{\Gamma}}_1 - \mathbf{D}_2\mathbf{\Gamma}_2\|_2^2 + \alpha_2\|\mathbf{\Gamma}_2\|_1$$

Success of Layered Lasso

$$\|\mathbf{\Gamma}_l\|_{0,\infty} < \frac{1}{3}\left(1 + \frac{1}{\mu(\mathbf{D}_L)}\right)$$

Layered Lasso guaranteed:

1. Find only correct places of nonzeros
2. Find all coefficients that are big enough

$$\|\hat{\mathbf{\Gamma}}_l - \mathbf{\Gamma}_l\|_{2,\infty} \leq \epsilon_l$$

~~Forward pass always fails at recovering representations exactly~~

~~Success depends on ratio~~

Distance increases with layer

Layered Iterative Thresholding

$$\mathbf{\Gamma}_1^t = \mathcal{S}_{\alpha_1}\left(\mathbf{D}_1^T\mathbf{Y} + \left(\mathbf{I} - \mathbf{D}_1^T\mathbf{D}_1\right)\mathbf{\Gamma}_1^{t-1}\right)$$

$$\mathbf{\Gamma}_2^t = \mathcal{S}_{\alpha_2}\left(\mathbf{D}_2^T\hat{\mathbf{\Gamma}}_1 + \left(\mathbf{I} - \mathbf{D}_2^T\mathbf{D}_2\right)\mathbf{\Gamma}_2^{t-1}\right)$$

Multi-layered Convolutional Sparse Modeling

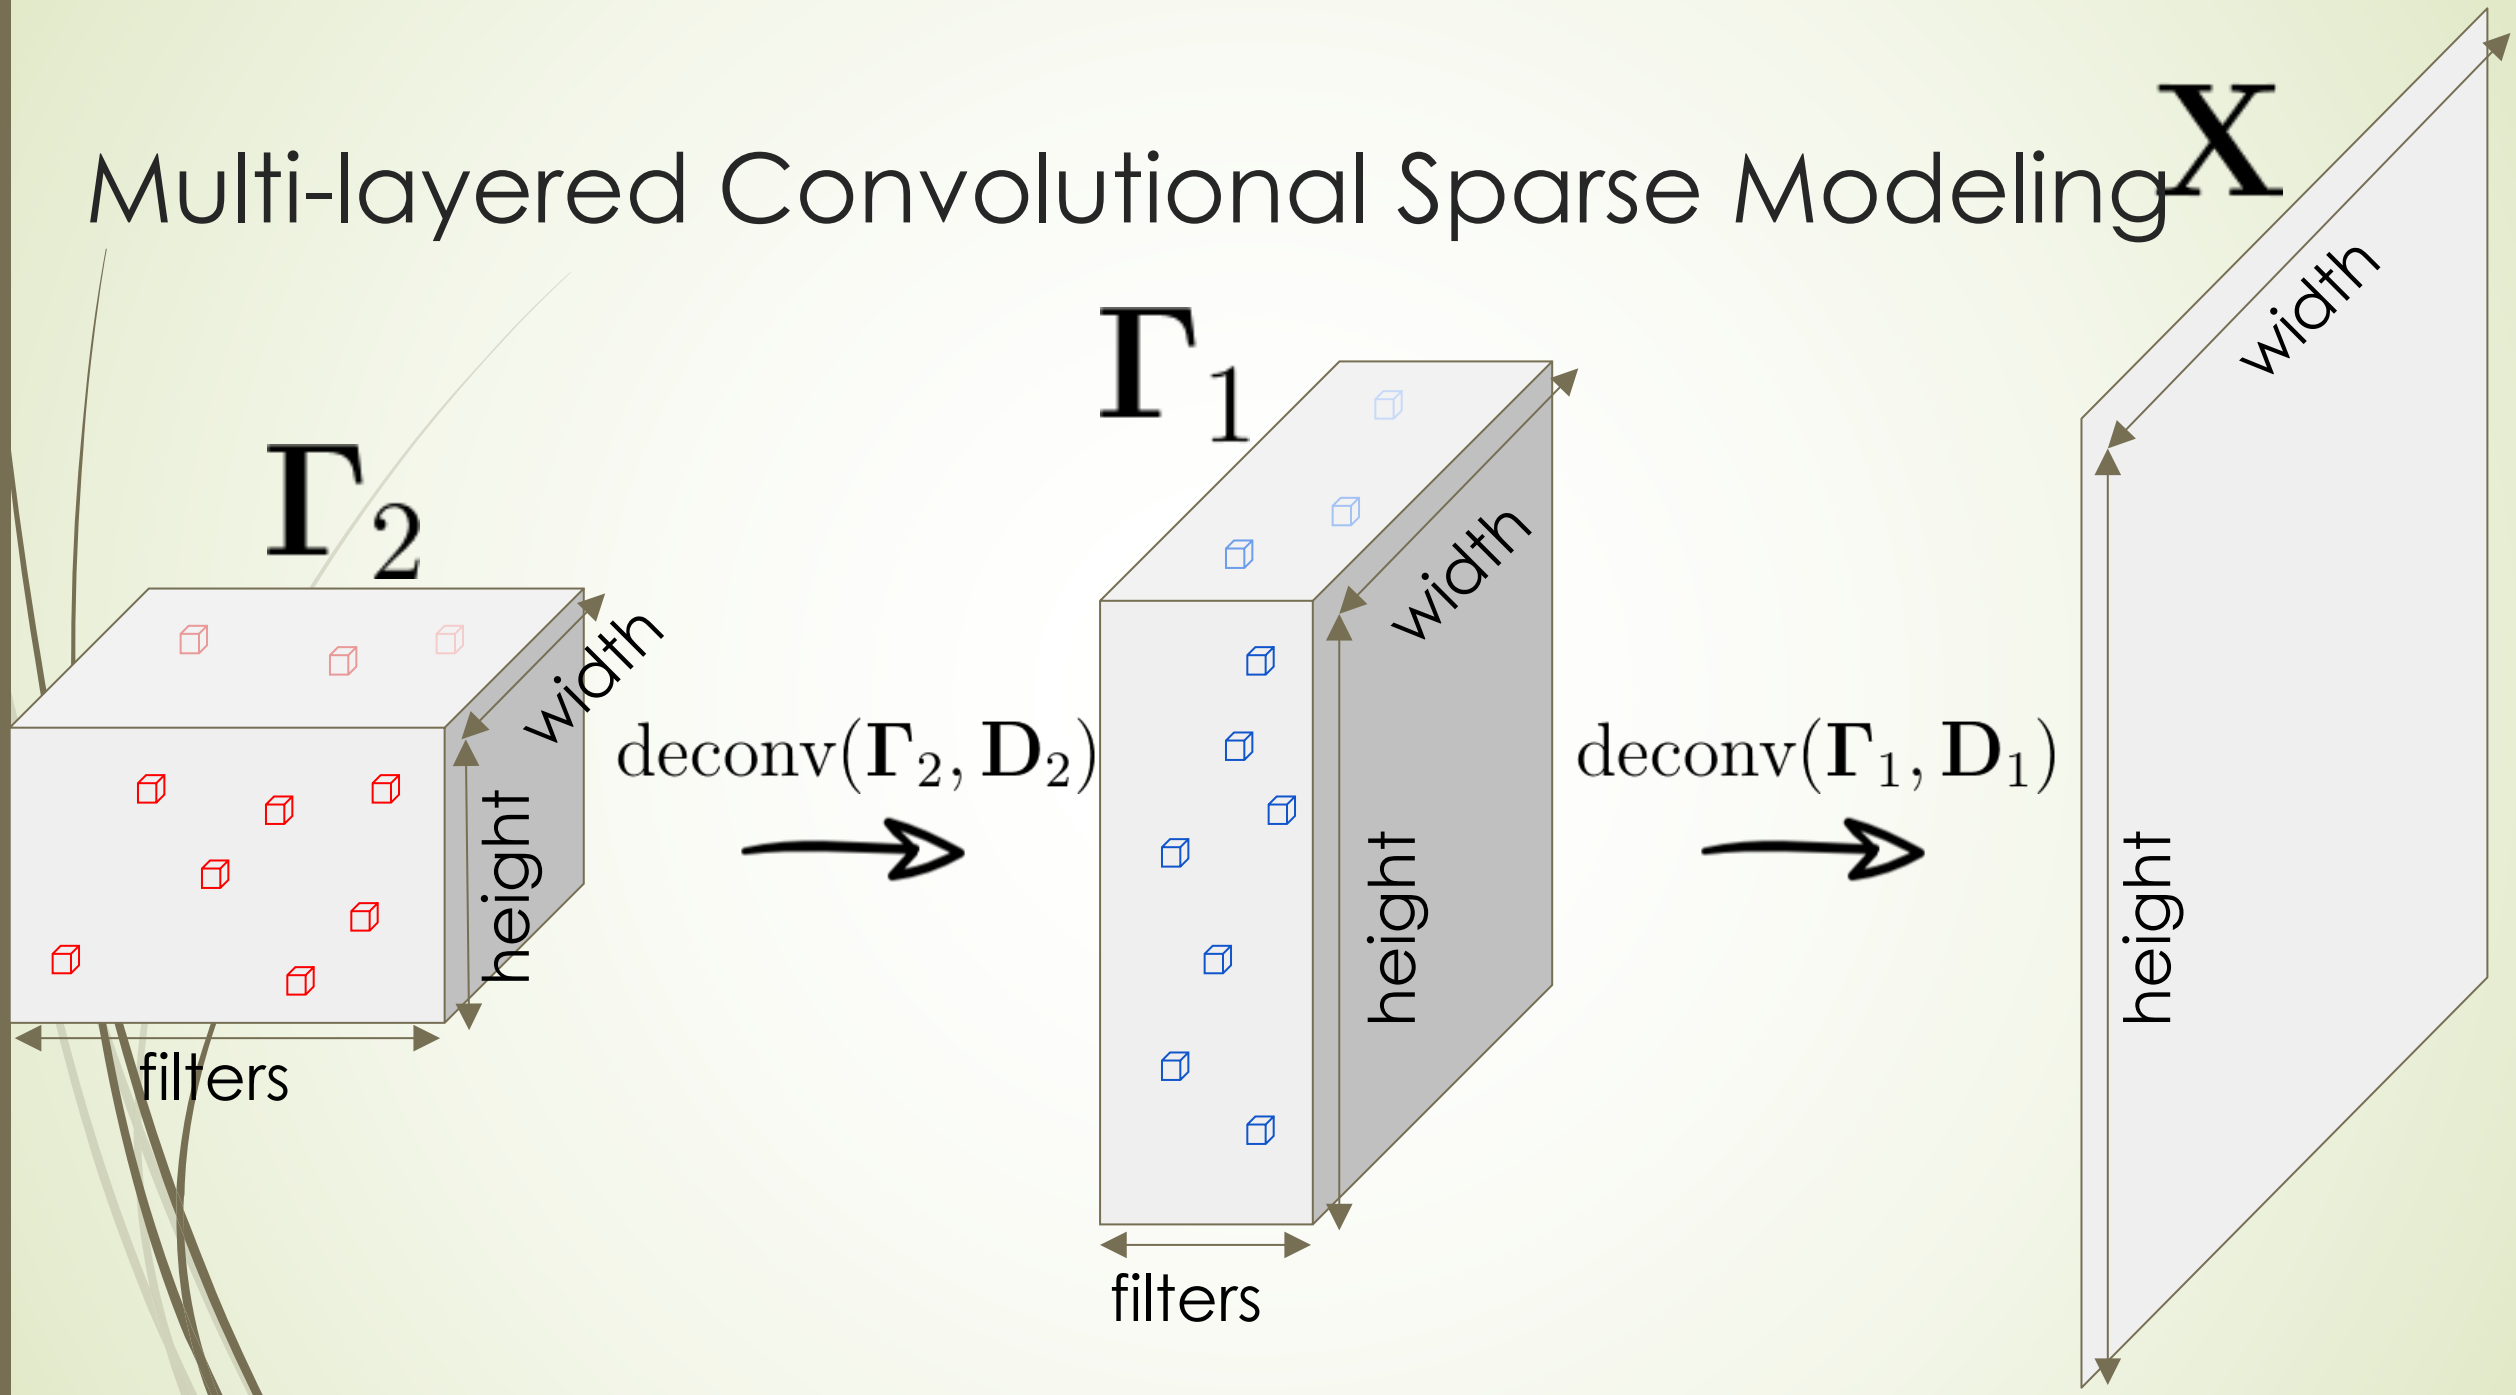

Summary

1. 

   Sparsity well established theoretically

2. 

   Sparsity is covertly exploited in practice: ReLU, dropout, stride, dilation, ...

3. 

   Sparsity is the secret sauce behind CNN

4. 

   Need to bring sparsity to the surface to better understand CNNs

5. 

   Andrej Karpathy agrees

Thank you!