

# Exponentially Weighted Imitation Learning for Batched Historical Data

Qing Wang<sup>1</sup>, Jiechao Xiong<sup>1</sup>, Lei Han<sup>1</sup>, Peng Sun<sup>1</sup>, Han Liu<sup>12</sup>, Tong Zhang<sup>1</sup>


<sup>1</sup>Tencent AI Lab <sup>2</sup>Northwestern University

# Outline

- Background:
  - Reinforcement Learning: from Simulation
  - Our Problem: Deep Policy Learning from Data
- Method:
  - Imitation Learning
  - Off-policy Reinforcement Learning
  - Monotonic Advantage Re-Weighted IL
- Experiments:
  - HFO, TORCS, King of Glory
- Discussion

# Reinforcement Learning

1. Environment provide initial **state** to the agent.
2. Agent observes the state and takes an **action**.
3. Environment receives the action, outputs a **reward**, and evolves to the next **state**.
4. Agent observes the next **state** and takes another **action**.



$S$  : state in state space  $\mathcal{S}$

$a$  : action in action space  $\mathcal{A}$

$r$  : reward  $\mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

# Example: Breakout (Atari game)

**state:** (stacked) pixels on the screen

**action:** left, right, ...

**reward:** +1 for hit a brick

**terminal:** drop the ball

or clear two screen of bricks




# Reinforcement Learning

- **Initial State:** the distribution of initial state  $s_0$  is denoted by  $d_0 \in \Delta(\mathcal{S})$
- **Stochastic Policy:** a map from state to a probabilistic distribution in action space

$$\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$$

- **Objective:** the performance of a policy is measured by its expected sum of discounted reward

$$\eta(\pi) = \mathbb{E}_{d_0, \pi} \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$



$s$  : state in state space  $\mathcal{S}$

$a$  : action in action space  $\mathcal{A}$

$r$  : reward  $\mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

# A Mathematical Formulation of RL

- Markov Decision Process (MDP)

$$(\mathcal{S}, \mathcal{A}, P, r, d_0, \gamma)$$

$\mathcal{S}$  is the state space

$\mathcal{A}$  is the action space


$P$  is the transition probability  $\mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$

$r$  is the reward function  $\mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

$d_0$  is the distribution of initial state

$\gamma \in (0, 1)$  is the discount factor

- Markovian: *given current state, future is independent of past*
  - or else, a POMDP



$s$  : state in state space  $\mathcal{S}$

$a$  : action in action space  $\mathcal{A}$

$r$  : reward  $\mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

Different definitions exist, e.g. finite horizon vs infinite horizon

# State Value, Action Value, and Advantage

- We define the value of a state as

$$V^{\pi}(s_t) = \mathbb{E}_{\pi|s_t} \sum_{l=0}^{\infty} \gamma^l r(s_{t+l}, a_{t+l})$$

- And the value of an action (under current state) as

$$Q^{\pi}(s_t, a_t) = \mathbb{E}_{\pi|s_t, a_t} \sum_{l=0}^{\infty} \gamma^l r(s_{t+l}, a_{t+l})$$


- The advantage is defined as the difference of Q and V

$$A^{\pi}(s_t, a_t) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)$$

- They could be estimated with Monte-Carlo method or Bootstrap

$$V(S_t) \leftarrow V(S_t) + \alpha \left[ G_t - V(S_t) \right] \quad \text{or} \quad V(S_t) \leftarrow V(S_t) + \alpha \left[ R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \right]$$

# State Value, Action Value, and Advantage





# DQN, DDPG, TRPO

**DQN** 
$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

**DDPG** 
$$\begin{aligned} \nabla_{\theta^\mu} J &\approx \mathbb{E}_{s_t \sim \rho^\beta} \left[ \nabla_{\theta^\mu} Q(s, a | \theta^Q) \big|_{s=s_t, a=\mu(s_t | \theta^\mu)} \right] \\ &= \mathbb{E}_{s_t \sim \rho^\beta} \left[ \nabla_a Q(s, a | \theta^Q) \big|_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) \big|_{s=s_t} \right] \end{aligned}$$

**TRPO** 
$$\begin{aligned} &\underset{\theta}{\text{maximize}} \sum_s \rho_{\theta_{\text{old}}}(s) \sum_a \pi_\theta(a|s) A_{\theta_{\text{old}}}(s, a) \\ &\text{subject to } \overline{D}_{\text{KL}}^{\rho_{\theta_{\text{old}}}}(\theta_{\text{old}}, \theta) \leq \delta. \end{aligned}$$

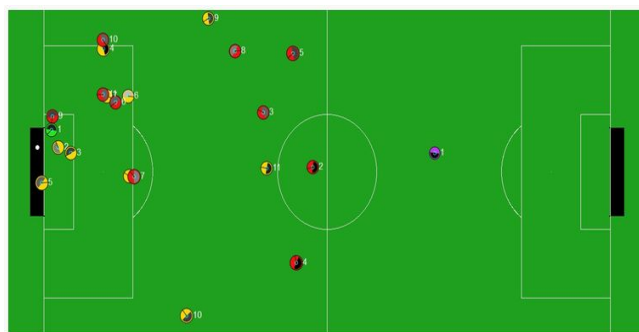
See DQN (Mnih et al., 2013), DDPG (Lillicrap et al., 2016), TRPO (John Schulman et al., 2015)

# Our Problem

- Deep **Policy Learning** from **Data** (without a **Simulator** !)
- $\pi : s \rightarrow a$  for complex environment
- Examples:
  - Autonomous Driving, Sports, Game bot, ...



TORCS



Robocup 2D



王者荣耀

# Imitation Learning (behavior cloning)

- Learn the map from  $s$  to  $a$
- Let  $\pi_\theta(a|s)$  close to  $\pi(a|s)$   
target policy behavior policy
- e.g. minimize KL-divergence

$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$

$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$

$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$

$$L_p = D_{\text{KL}}^{d_\pi}(\pi || \pi_\theta) = -\mathbb{E}_{s \sim d_\pi(s), a \sim \pi(a|s)} \log \pi_\theta(a|s) + C$$

- where

$$D_{\text{KL}}^{d_\pi}(\pi' || \pi) = \sum_s d(s) \sum_a \pi'(a|s) \log \frac{\pi'(a|s)}{\pi(a|s)}$$

$$d_\pi(s) = (1 - \gamma) \mathbb{E}_{d_0, \pi} \sum_{t=0}^{\infty} \gamma^t \mathbf{1}(s_t = s)$$

use of reward information?

quality of expert trajectories?

# Off-policy Deep Reinforcement Learning

- Value iteration:
  - DQN, DDPG, Hybrid action?
  - How to make sure the learned policy performs similar to (or better than) the behavior policy?
- Policy iteration:
  - TRPO?, Retrace?
  - What if we do not have probability of behavior policy (common in practice)?

$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$

$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$

$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$

See DQN (Mnih et al., 2013), DDPG (Lillicrap et al., 2016), TRPO (John Schulman et al., 2015), and Retrace (Remi Munos et al., 2016)

# Challenges

- Large state space
  - tabular algorithm? function approximation?

$$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$$

- Hybrid action space
  - discrete action: 0, 1, 2, 3, ..
  - continuous action: (x, y), r,  $\theta$

$$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$$

$$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$$

- Unknown behavior policy
  - hard for off-policy policy gradient method (Schulman et al., 2015)


$$\frac{\pi_{\theta}(a|s)}{\pi(a|s)} \beta A^{\pi}(s, a)$$

explicit value needed!

# Can we do better imitation learning?

- action: Good 😊 vs Bad 😞
- look ahead in the historic data to determine what action results in good consequence.
- What if we imitate good action only?
- What if we put larger sample weight on good actions?


$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$



$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$



$s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots$



# Imitate a better policy?

- Is there a better policy?
- Consider two policy  $\pi$  and  $\tilde{\pi}$  satisfying:

$$\tilde{\pi}(a|s) \geq \pi(a|s) \quad A^{\pi}(s, a) \geq 0$$

$$\tilde{\pi}(a|s) \leq \pi(a|s) \quad A^{\pi}(s, a) \leq 0$$

- then  $\tilde{\pi}$  is uniformly as good as, or better than  $\pi$  , i.e.

$$V^{\tilde{\pi}}(s) \geq V^{\pi}(s), \forall s \in S.$$

# Imitate a better policy?

- We can imitate  $\tilde{\pi}$  instead of  $\pi$
- For example, we can choose<sup>1</sup>:

$$\tilde{\pi} \propto \pi \exp(\beta A^\pi)$$

- then  $\tilde{\pi}$  is a better policy, we can imitate this one by

$$\begin{aligned} \arg \min_{\theta} D_{\text{KL}}^d(\tilde{\pi} || \pi_{\theta}) &= \arg \max_{\theta} \sum_s d(s) \sum_a \tilde{\pi}(a|s) \log \pi_{\theta}(a|s) \\ &= \arg \max_{\theta} \sum_s d(s) \exp(C(s)) \sum_a \pi(a|s) \boxed{\exp(\beta A^{\pi}(s, a))} \log \pi_{\theta}(a|s) \\ &\quad \text{better action, larger weight} \end{aligned}$$

<sup>1</sup> The derivation of  $\exp()$  is related to a few previous works e.g. (Peters et al., 2010), (Azar et al., 2012)



# Monotonic Advantage Re-Weighted IL

---

**Algorithm 1** Monotonic Advantage Re-Weighted Imitation Learning (MARWIL)

---

**Input:** Historical data  $\mathcal{D}$  generated by  $\pi$ , hyper-parameter  $\beta$ .

For each trajectory  $\tau$  in  $\mathcal{D}$ , estimate advantages  $\hat{A}^\pi(s_t, a_t)$  for time  $t = 1, \dots, T$ .

Maximize  $\sum_{\tau \in \mathcal{D}} \sum_{(s_t, a_t) \in \tau} \exp(\beta \hat{A}^\pi(s_t, a_t)) \log \pi_\theta(a_t | s_t)$  with respect to  $\theta$ .

---

monotonic advantage reweighting

also works with other formulation, e.g. ReLU instead of Exp

$$\sum_a \pi(a|s) ((\beta A^\pi(s, a))_+ + \epsilon) \log \pi_\theta(a|s)$$

# A lower bound on policy improvement

The performance of a policy is measured by its **expected sum of discounted reward**:

$$\eta(\pi) = \mathbb{E}_{d_0, \pi} \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

For practical usage, we usually seek a parametric approximation of  $\tilde{\pi}$ . The following proposition gives a lower bound of policy improvement for the parametric policy  $\pi_\theta$ .

**Proposition 2.** *Suppose we use parametric policy  $\pi_\theta$  to approximate the improved policy  $\tilde{\pi}$  defined in Formula 3, we have the following lower bound on the policy improvement*

$$\eta(\pi_\theta) - \eta(\pi) \geq -\frac{\sqrt{2}}{1-\gamma} \delta_1^{\frac{1}{2}} M^{\pi_\theta} + \frac{1}{(1-\gamma)\beta} \delta_2 - \frac{\sqrt{2}\gamma\epsilon_\pi^{\tilde{\pi}}}{(1-\gamma)^2} \delta_2^{\frac{1}{2}} \quad (8)$$

where  $\delta_1 = \min(D_{KL}^{d_{\tilde{\pi}}}(\pi_\theta||\tilde{\pi}), D_{KL}^{d_{\tilde{\pi}}}(\tilde{\pi}||\pi_\theta))$ ,  $\delta_2 = D_{KL}^{d_\pi}(\tilde{\pi}||\pi)$ ,  $\epsilon_\pi^{\pi'} = \max_s |\mathbb{E}_{a \sim \pi'} A^\pi(s, a)|$ , and  $M^\pi = \max_{s,a} |A^\pi(s, a)| \leq \max_{s,a} |r(s, a)|/(1-\gamma)$ .

# Experiments

- We consider 4 policy loss in our experiments

IL  $L_p = D_{\text{KL}}^{d_\pi}(\pi || \pi_\theta) = -\mathbb{E}_{s \sim d_\pi(s), a \sim \pi(a|s)} \log \pi_\theta(a|s) + C$

PG  $L_p = -\mathbb{E}_{s \sim d_\pi(s), a \sim \pi(a|s)} (\beta A^\pi(s, a) + 1) \log \pi_\theta(a|s) + C$

PGIS  $L_p = D_{\text{KL}}^{d_\pi}(\pi || \pi_\theta) - (1 - \gamma) \beta L^{d_\pi, \pi}(\pi_\theta)$   
 $= -\mathbb{E}_{s \sim d_\pi(s), a \sim \pi(a|s)} \left( \frac{\pi_\theta(a|s)}{\pi(a|s)} \beta A^\pi(s, a) + \log \pi_\theta(s, a) \right) + C$

MARWIL  $L_p = D_{\text{KL}}^d(\tilde{\pi} || \pi_\theta) = -\mathbb{E}_{s \sim d_\pi(s), a \sim \pi(a|s)} \log(\pi_\theta(a|s)) \exp(\beta A^\pi(s, a)) + C$

# Experiments

---

**Algorithm 2** Stochastic Gradient Algorithm for MARWIL

---

**Input:** Policy loss  $L_p$  being one of 9 to 12. base policy  $\pi$ , parameter  $m, c_v$ .

Randomly initialize  $\pi_\theta$ . Empty replay memory  $D$ .

Fill  $D$  with trajectories from  $\pi$  and calculate  $R_t$  for each  $(s_t, a_t)$  in  $D$ .

**for**  $i = 1$  **to**  $N$  **do**

    Sample a batch  $B = \{(s_k, a_k, R_k)\}_m$  from  $D$ .

    Compute mini-batch gradient  $\nabla_\theta \hat{L}_p, \nabla_\theta \hat{L}_v$  of  $B$ .



    Update  $\theta$ :  $-\Delta\theta \propto \nabla_\theta \hat{L}_p + c_v \nabla_\theta \hat{L}_v$

**end for**

---

# Experiments: HFO

- Half Field Offense:
- state: 59 floats
- action: Dash( $r, \theta$ ), Turn( $\theta$ ), Kick( $r, \theta$ )



# Experiments: HFO


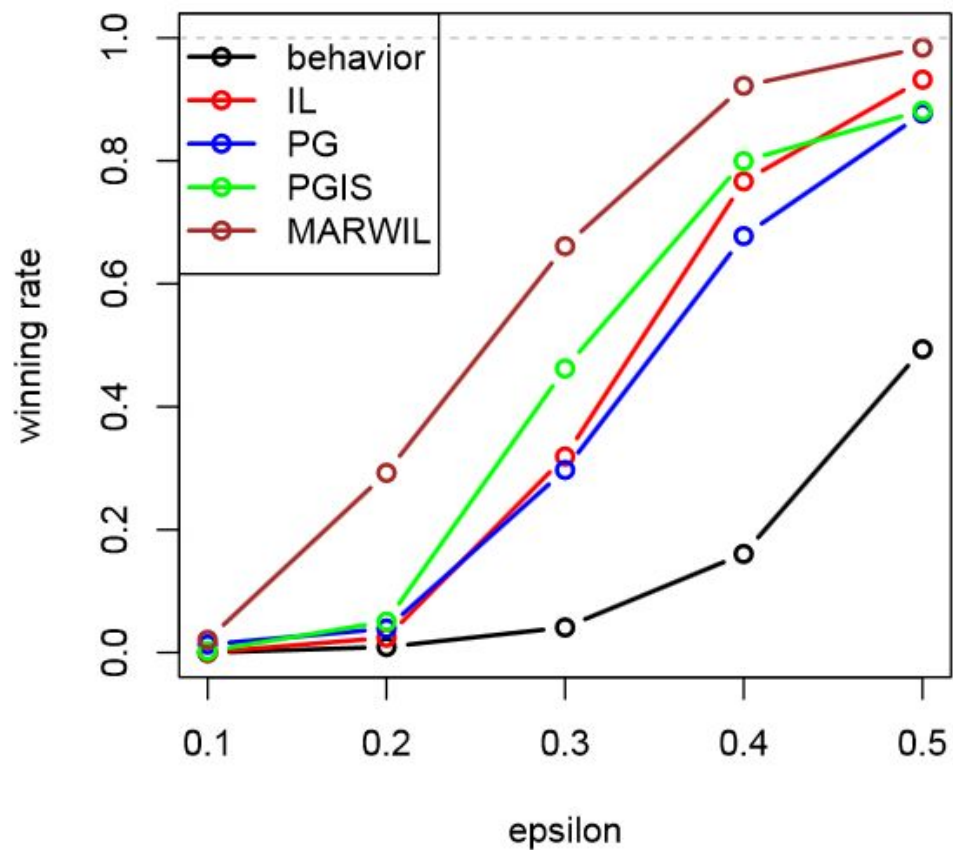
- We add different level of noise in the historic data:

$$a_t \sim \begin{cases} \pi_{\text{perfect}}(\cdot|s_t) + N(0, \sigma) & \text{w.p. } \epsilon \\ \pi_{\text{random}}(\cdot|s_t) + N(0, \sigma) & \text{w.p. } 1 - \epsilon \end{cases}$$

- the policy is modeled as

$$\pi_{\theta}((k, x_k)|s) = p_{\theta}(k|s)N(x_k|\mu_{\theta,k}, \sigma), \quad k \in \{1, 2, 3\}, x_k \in \mathbb{R}^2$$

# Experiments: HFO



# Experiments: TORCS

- Raw image input to steering angle



64x64x3



$[-\pi, \pi]$

- We vary the parameter beta:  $-\mathbb{E}_{s \sim d_{\pi}(s), a \sim \pi(a|s)} \log(\pi_{\theta}(a|s)) \exp(\beta A^{\pi}(s, a))$


Table 1: Performance of PG and MARWIL in TORCS, where  $\beta = 0$  is the case of IL. Different  $\beta$  are tested in the experiments. The performance is evaluated on the sum of rewards per episode.

$\beta$	0.0	0.25	0.5	0.75	1.0
PG	2710	6396	6735	6758	7152
MARWIL	(2710)	5583	6832	7670	9492




# Experiments: 王者荣耀

- state:



- action: 貂蝉: 6 discrete type



# Experiments: 王者荣耀

- Network structure




Figure 2: Network structure for our AI agent in King of Glory

# Discussion & Future work

- Imitating a **better** policy, by **monotonic advantage reweighting**
  - suitable for large state space and hybrid action space
  - do not require complete knowledge of behavior policy
  - robust under complex function approximation (with lower bound)
- The method also works in full reinforcement learning
  - improve over self-generated trajectories

Questions?

Thanks ~ 

# Backup slides

The performance of a policy  $\pi$  is measured by its expected discounted reward:

$$\eta(\pi) = \mathbb{E}_{d_0, \pi} \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

where  $\mathbb{E}_{d_0, \pi}$  means  $s_0 \sim d_0$ ,  $a_t \sim \pi(a_t|s_t)$ , and  $s_{t+1} \sim P(s_{t+1}|s_t, a_t)$ . We omit the subscript  $d_0$  when there is no ambiguity. In [Kakade and Langford, 2002], a useful equation has been proved that

$$\eta(\pi') - \eta(\pi) = \frac{1}{1 - \gamma} \sum_s d_{\pi'}(s) \sum_a \pi'(a|s) A^\pi(s, a)$$

where  $d_\pi$  is the discounted visiting frequencies defined as  $d_\pi(s) = (1 - \gamma) \mathbb{E}_{d_0, \pi} \sum_{t=0}^{\infty} \gamma^t \mathbf{1}(s_t = s)$  and  $\mathbf{1}(\cdot)$  is an indicator function. In addition, define  $L^{d, \pi}(\pi')$  as

$$L^{d, \pi}(\pi') = \frac{1}{1 - \gamma} \sum_s d(s) \sum_a \pi'(a|s) A^\pi(s, a)$$

then from [Schulman et al., 2015, Theorem 1], the difference of  $\eta(\pi')$  and  $\eta(\pi)$  can be approximated by  $L^{d_\pi, \pi}(\pi')$ , where the approximation error is bounded by total variance  $D_{\text{TV}}^{d_\pi}(\pi', \pi)$ , which can be further bounded by  $D_{\text{KL}}^{d_\pi}(\pi' || \pi)$  or  $D_{\text{KL}}^{d_\pi}(\pi || \pi')$ .

# Backup slides

Consider the problem

$$\tilde{\pi} = \arg \max_{\pi' \in \Pi} ((1 - \gamma)\beta L^{d_{\pi}, \pi}(\pi') - D_{\text{KL}}^{d_{\pi}}(\pi' || \pi)) \quad (3)$$

which has an analytical solution in the policy space  $\Pi$  [Azar et al., 2012, Appendix A, Proposition 1]


$$\tilde{\pi}(a|s) = \pi(a|s) \exp(\beta A^{\pi}(s, a) + C(s)) \quad (4)$$

where  $C(s)$  is a normalizing factor to ensure that  $\sum_{a \in \mathcal{A}} \tilde{\pi}(a|s) = 1$  for each state  $s$ . Then

$$\begin{aligned} \arg \min_{\theta} D_{\text{KL}}^d(\tilde{\pi} || \pi_{\theta}) &= \arg \max_{\theta} \sum_s d(s) \sum_a \tilde{\pi}(a|s) \log \pi_{\theta}(a|s) \\ &= \arg \max_{\theta} \sum_s d(s) \exp(C(s)) \sum_a \pi(a|s) \exp(\beta A^{\pi}(s, a)) \log \pi_{\theta}(a|s) \end{aligned} \quad (5)$$

# Remark: MDP vs SG

Markov Decision Process  
(MDP)



Example:

- Most RL env in OpenAI Gym.
- Game against fixed opponent.

Stochastic Game  
(SG)



Example:

- Chess, Go
- VizDoom, MOBA, Starcraft

- We have the simulator for the game, but we do not have the ``simulator" of opponent human player (as part of the environment)
- (Recall AlphaGo, this work is similar to learning from human replay (AlphaGo Master). It is also possible to learn with pure self-play (AlphaGo Zero))

# AlphaGo as Imitating MCTS policy

- Self-play with MCTS:

$$(\mathbf{p}, v) = f_{\theta}(s) \text{ and } l = (\boxed{z} - v)^2 - \boxed{\pi}^T \log \mathbf{p} + c \|\theta\|^2$$

Step with self-play

MCTS search

- MCTS as a **policy improvement operator**.
  - In our case: **monotonic advantage reweighting**, instead of MCTS
- Self-play with search as a **policy evaluation operator**.



# Reference

- R. Sutton and A. Barto, Reinforcement Learning: An Introduction, 2017
- J. Schulman et al., Trust Region Policy Optimization, 2015
- V. Minh et al., Playing Atari with Deep Reinforcement Learning, 2013
- D. Silver et al., Mastering the game of go with deep neural networks and tree search, 2016
- T. P. Lillicrap et al., Continuous Control with Deep Reinforcement Learning, 2016
- R. Munos et al., Safe and efficient off-policy reinforcement learning, 2016
- M. Bain et al., A Framework for Behavioural cloning, 1999
- S. Ross et al., A Reduction of Imitation Learning and Structured Prediction to no-regret Online Learning, 2011
- J. Peters et al., Relative Entropy Policy Search, 2010
- M. G. Azar et al., Dynamic Policy Programming, 2012
- D. Jiang et al., Feedback-based tree search for reinforcement learning, 2018